HIERARCHICAL

==========

RUN 1: Using Split A for validation, Split B for testing

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

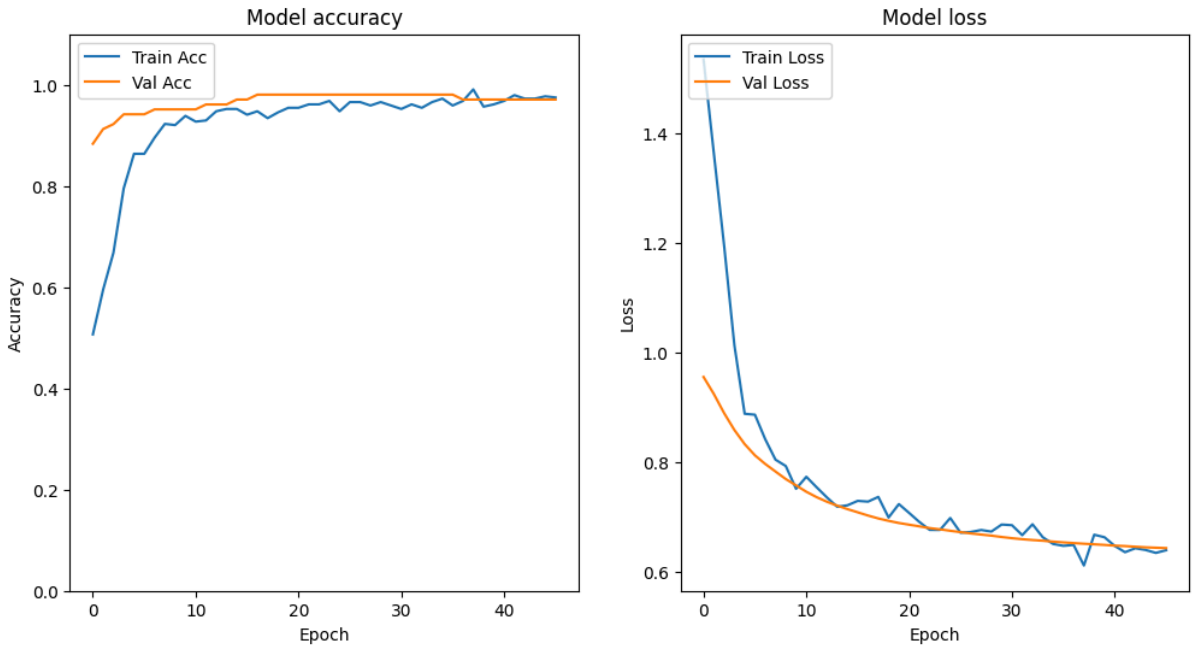

[Threshold Optimization] Best balanced_accuracy: 0.8889 at threshold=0.030

Optimal threshold (Stage 1 (I vs Rest)): 0.030

	precision	recall	f1-score	support
0	1.0000	0.7900	0.8827	100
1	0.1250	1.0000	0.2222	3
accuracy			0.7961	103
macro avg	0.5625	0.8950	0.5525	103
weighted avg	0.9745	0.7961	0.8634	103

Balanced Accuracy: 0.895

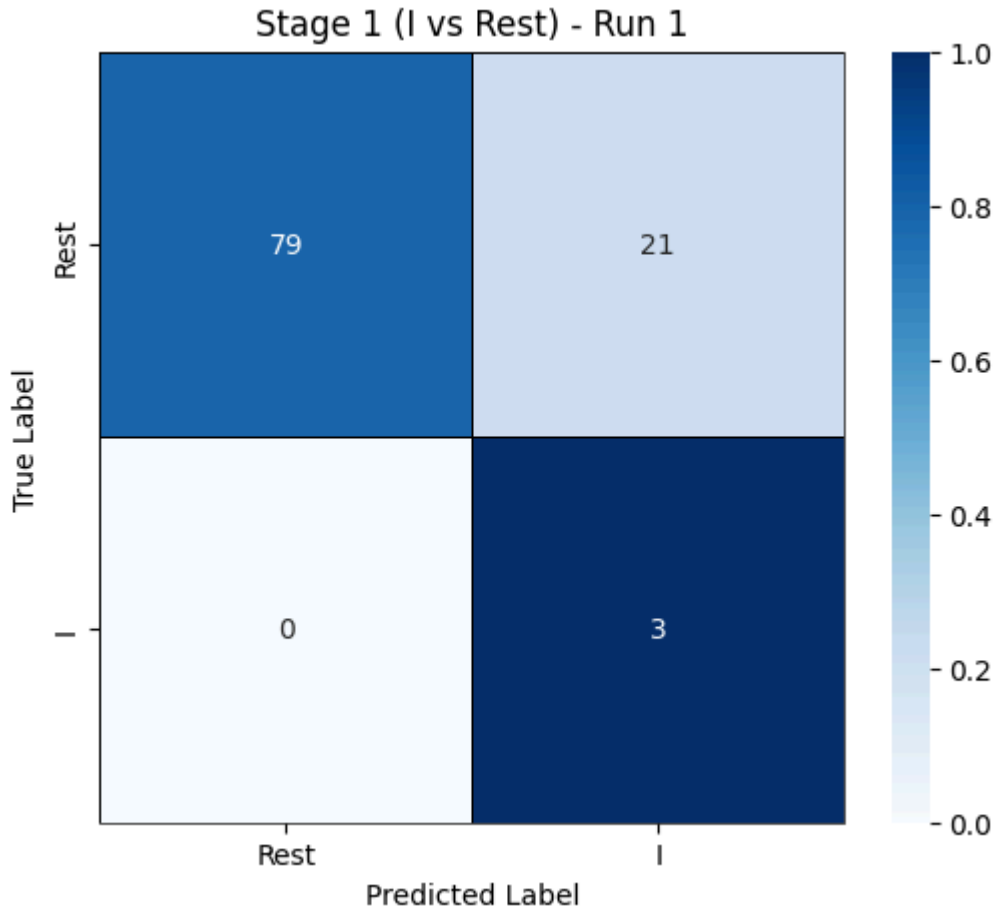

-- Stage 2 (DC vs Rest) --

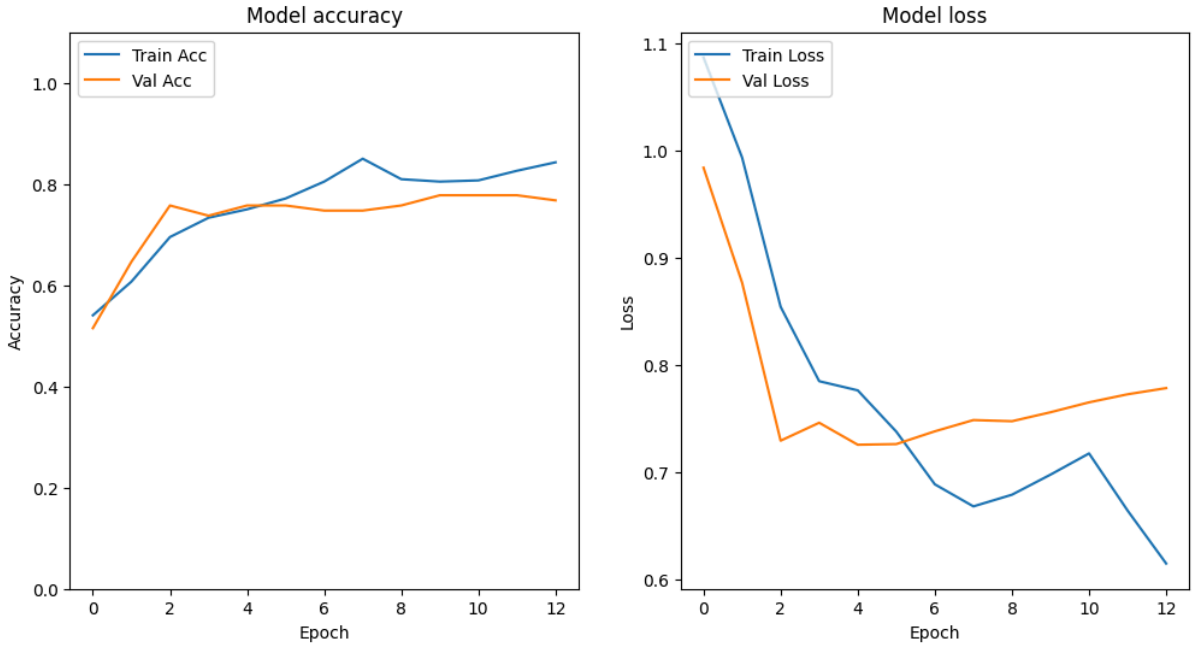

[Threshold Optimization] Best balanced_accuracy: 0.6451 at threshold=0.126
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.9032	0.6437	0.7517	87
1	0.1842	0.5385	0.2745	13
accuracy			0.6300	100
macro avg	0.5437	0.5911	0.5131	100
weighted avg	0.8098	0.6300	0.6896	100

Balanced Accuracy: 0.5910698496905393

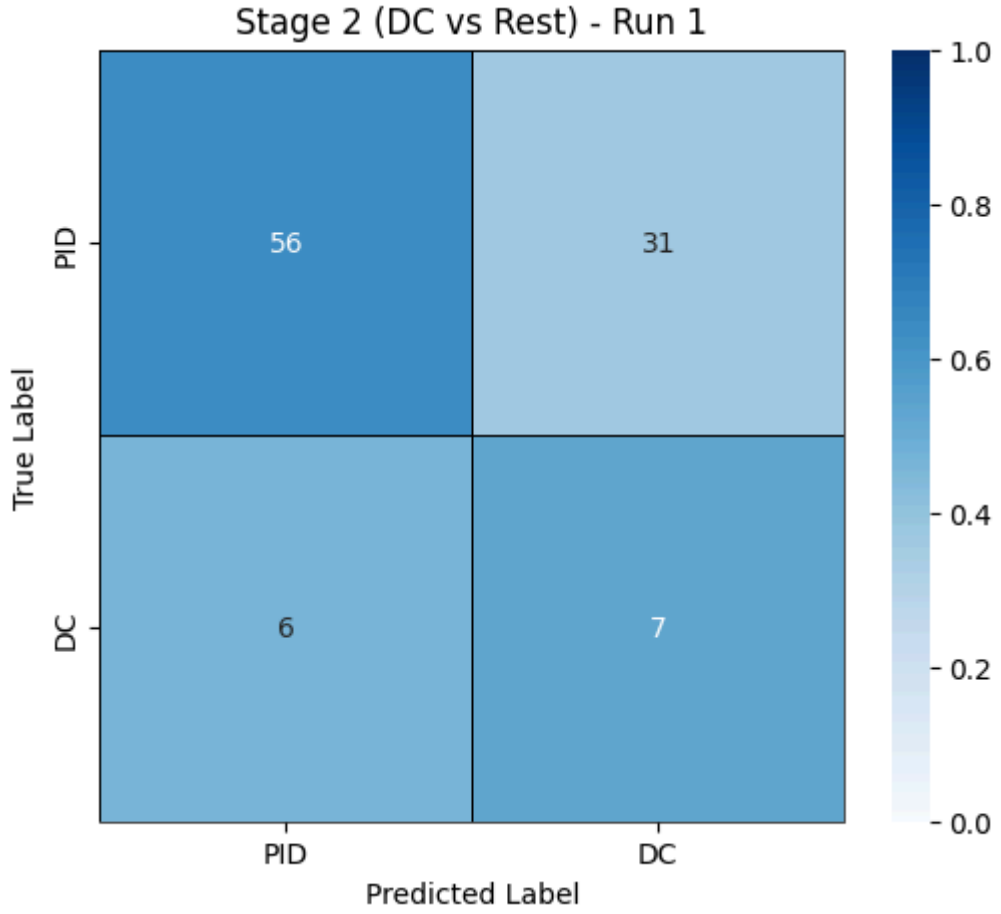

-- Stage 3 (Multiclass) --

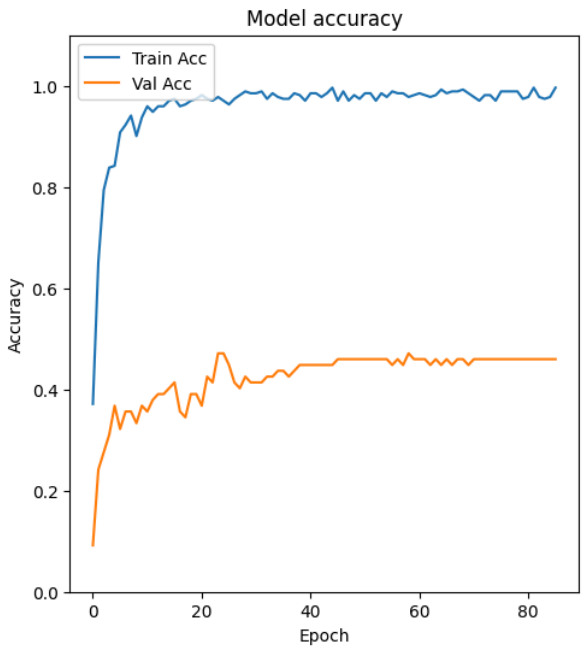

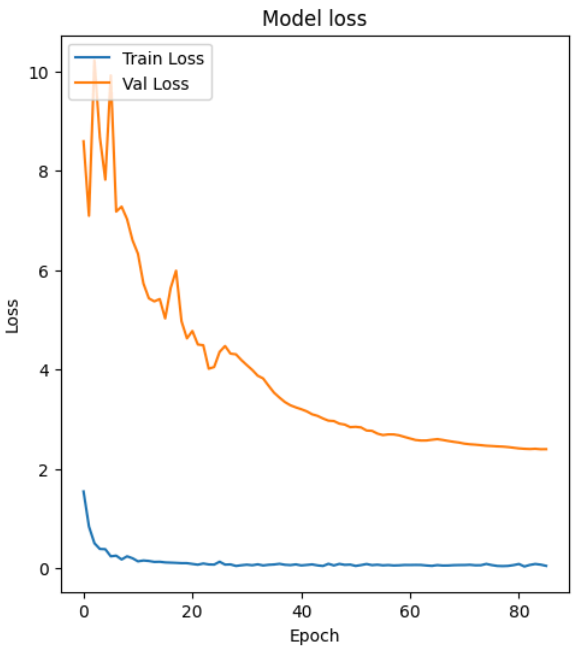

	precision	recall	f1-score	support
0	0.5000	0.2500	0.3333	4
1	0.2667	0.5000	0.3478	24
2	0.6250	0.4386	0.5155	57
3	0.0000	0.0000	0.0000	2
accuracy			0.4368	87
macro avg	0.3479	0.2971	0.2992	87
weighted avg	0.5060	0.4368	0.4490	87

Balanced Accuracy: 0.29714912280701755

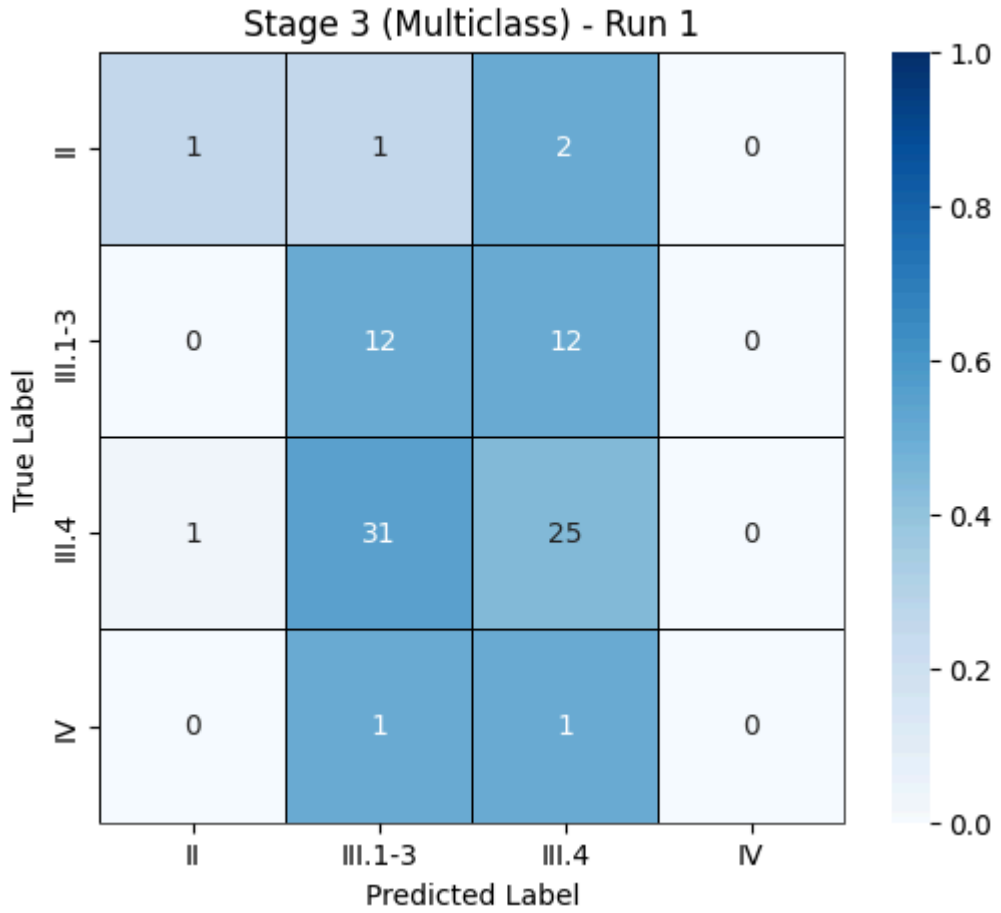

== Soft-Gated Overall (Test Set) - Run 1 ==

	precision	recall	f1-score	support
DC	0.1304	0.2308	0.1667	13
I	0.0423	1.0000	0.0811	3
II	0.0000	0.0000	0.0000	4
III.1-3	0.1667	0.0417	0.0667	24
III.4	0.6667	0.0351	0.0667	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.0874	103
macro avg	0.1677	0.2179	0.0635	103
weighted avg	0.4255	0.0874	0.0758	103

Balanced Accuracy: 0.21792060278902384

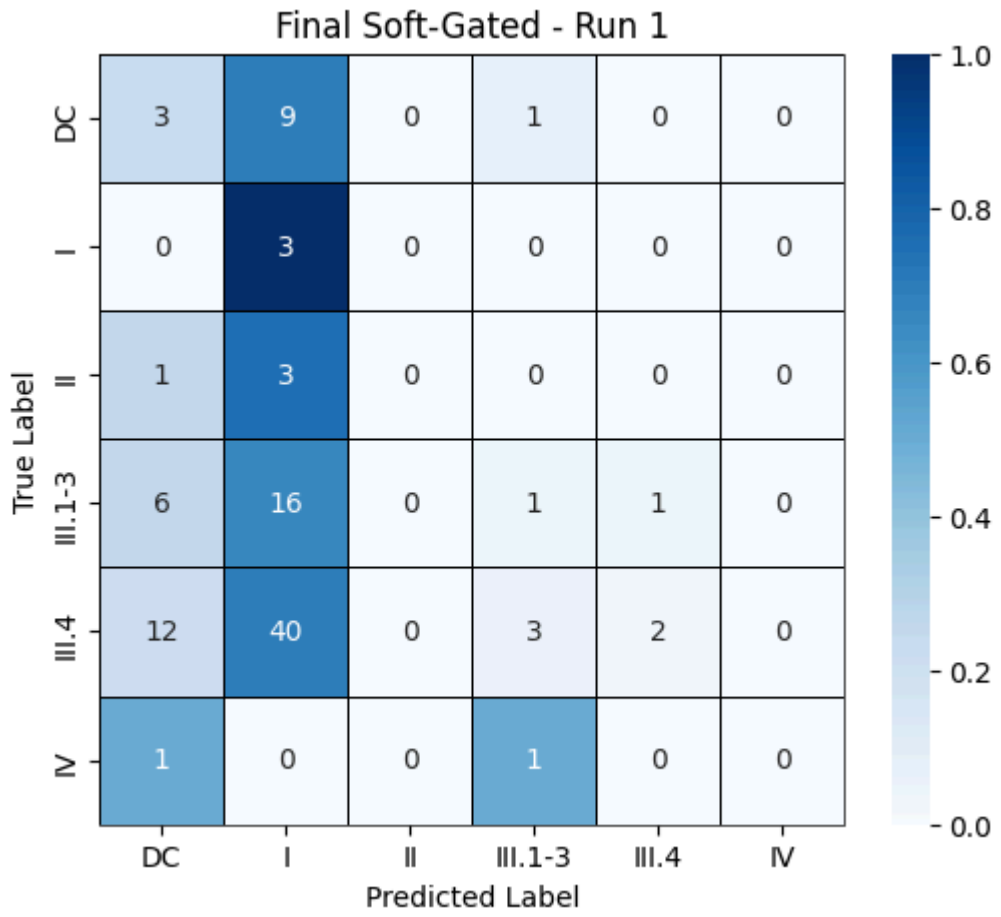

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

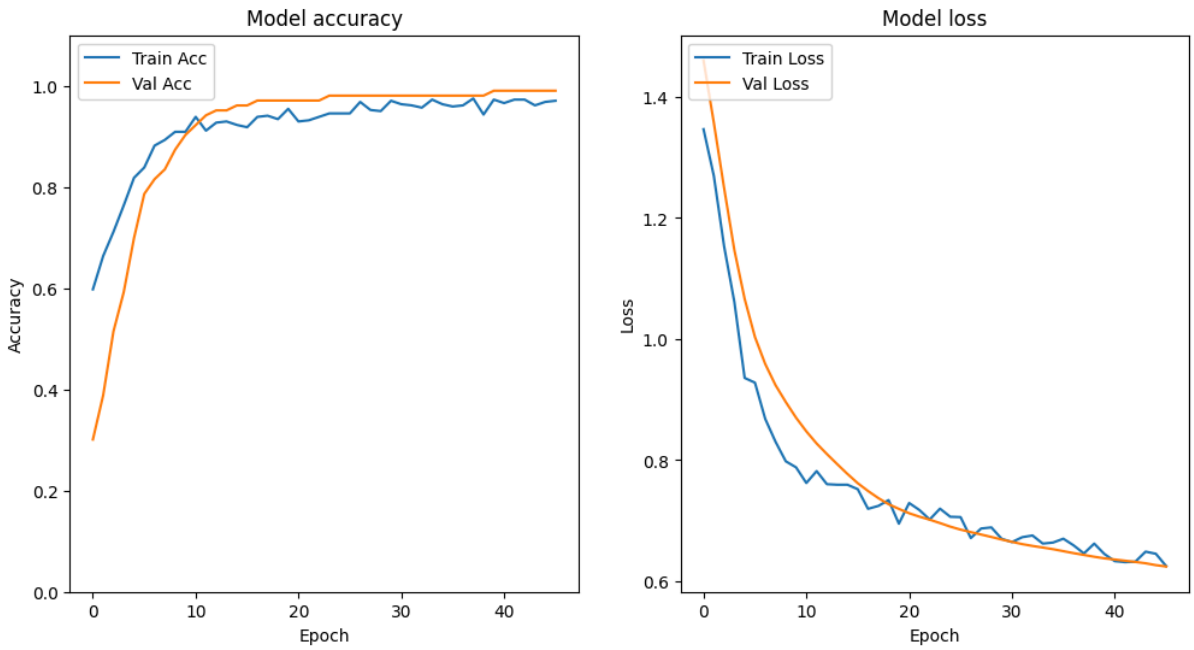

[Threshold Optimization] Best balanced_accuracy: 0.9950 at threshold=0.392
Optimal threshold (Stage 1 (I vs Rest)): 0.392

	precision	recall	f1-score	support
0	0.9892	0.9293	0.9583	99
1	0.3000	0.7500	0.4286	4
accuracy			0.9223	103
macro avg	0.6446	0.8396	0.6935	103
weighted avg	0.9625	0.9223	0.9378	103

Balanced Accuracy: 0.8396464646464646

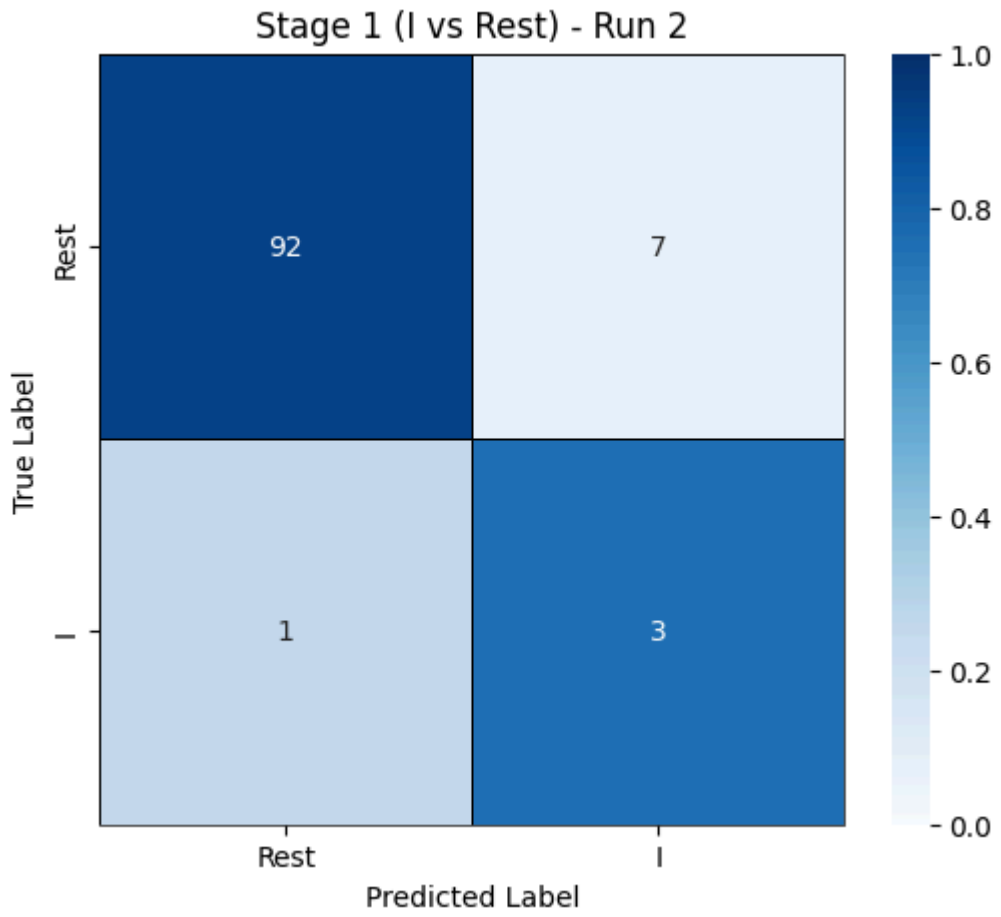

-- Stage 2 (DC vs Rest) --

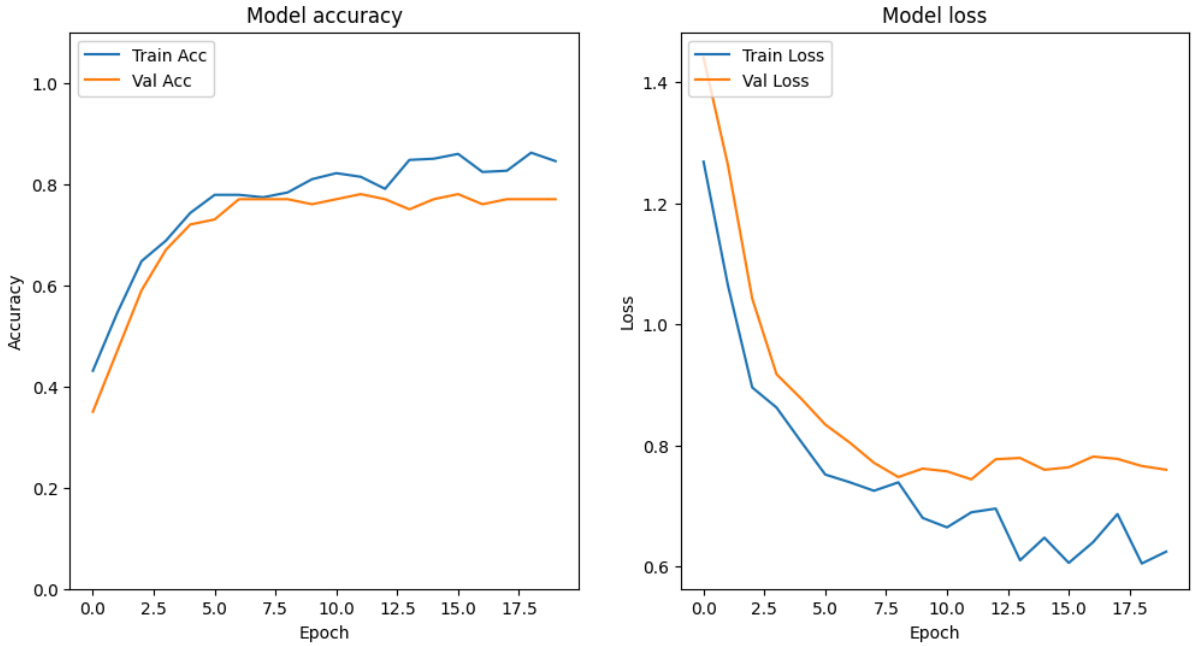

[Threshold Optimization] Best balanced_accuracy: 0.7679 at threshold=0.281
Optimal threshold (Stage 2 (DC vs Rest)): 0.281
precision recall f1-score support

0	0.8889	0.7356	0.8050	87
1	0.1481	0.3333	0.2051	12
accuracy			0.6869	99
macro avg	0.5185	0.5345	0.5051	99
weighted avg	0.7991	0.6869	0.7323	99

Balanced Accuracy: 0.5344827586206896

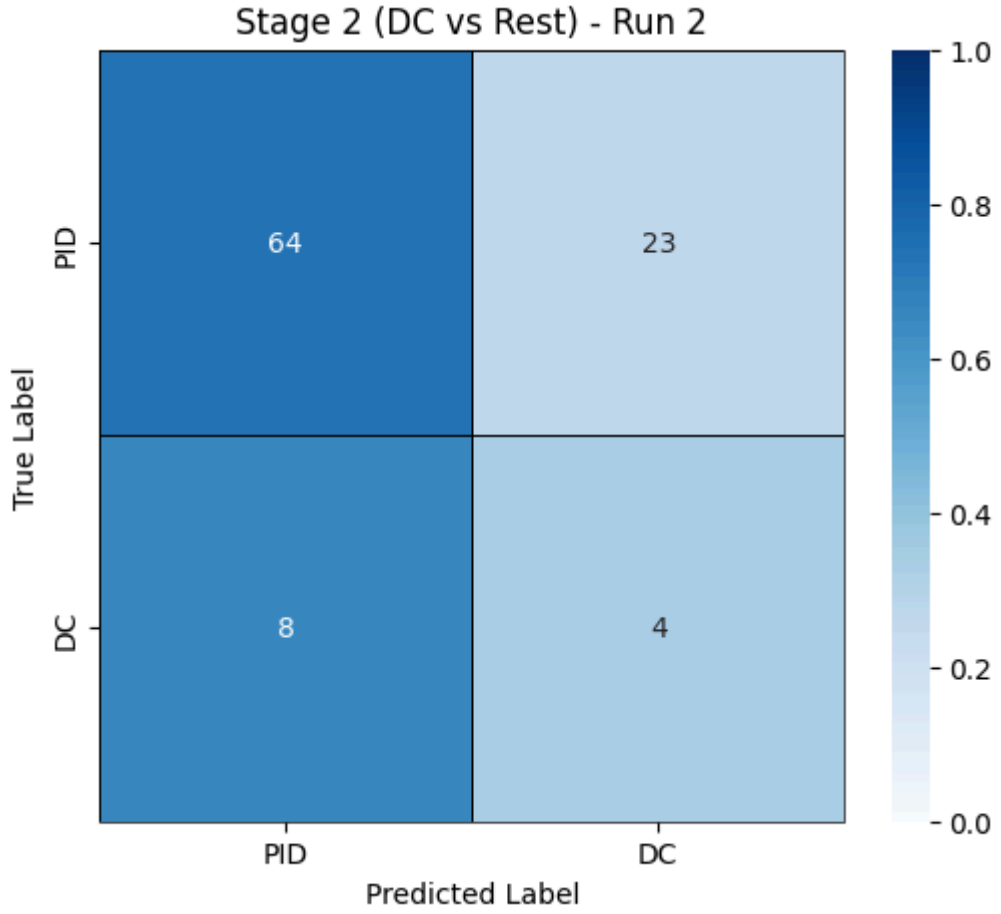

-- Stage 3 (Multiclass) --

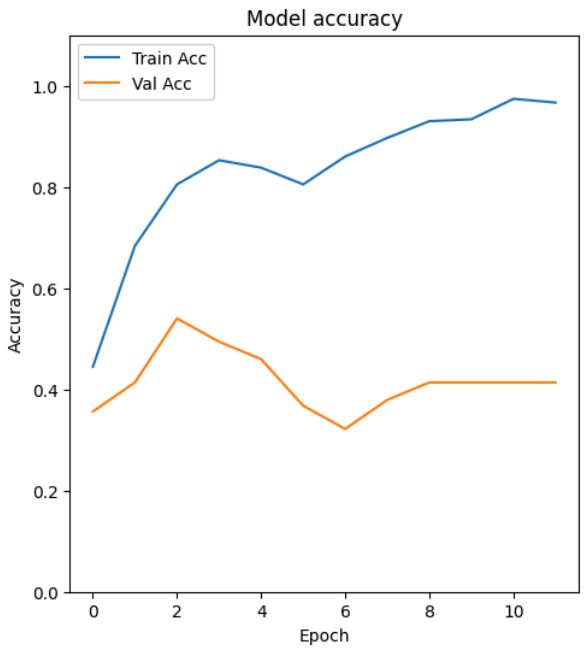

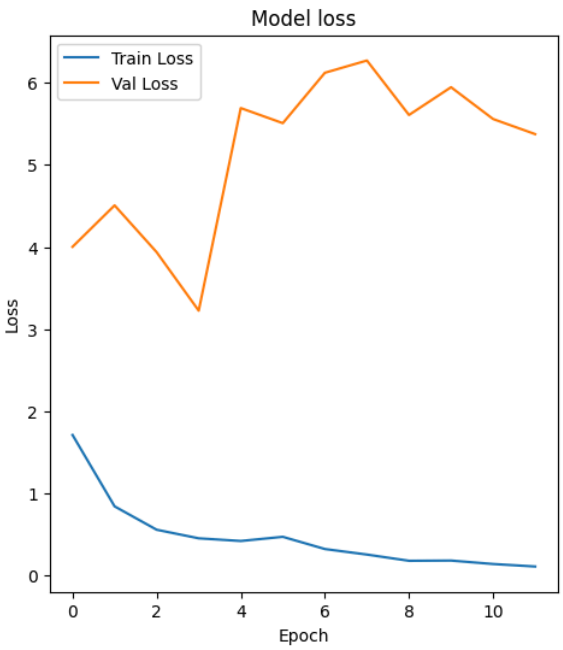

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	4
1	0.3529	0.4800	0.4068	25
2	0.7442	0.5714	0.6465	56
3	0.0000	0.0000	0.0000	2
accuracy			0.5057	87
macro avg	0.2743	0.2629	0.2633	87
weighted avg	0.5804	0.5057	0.5330	87

Balanced Accuracy: 0.26285714285714284

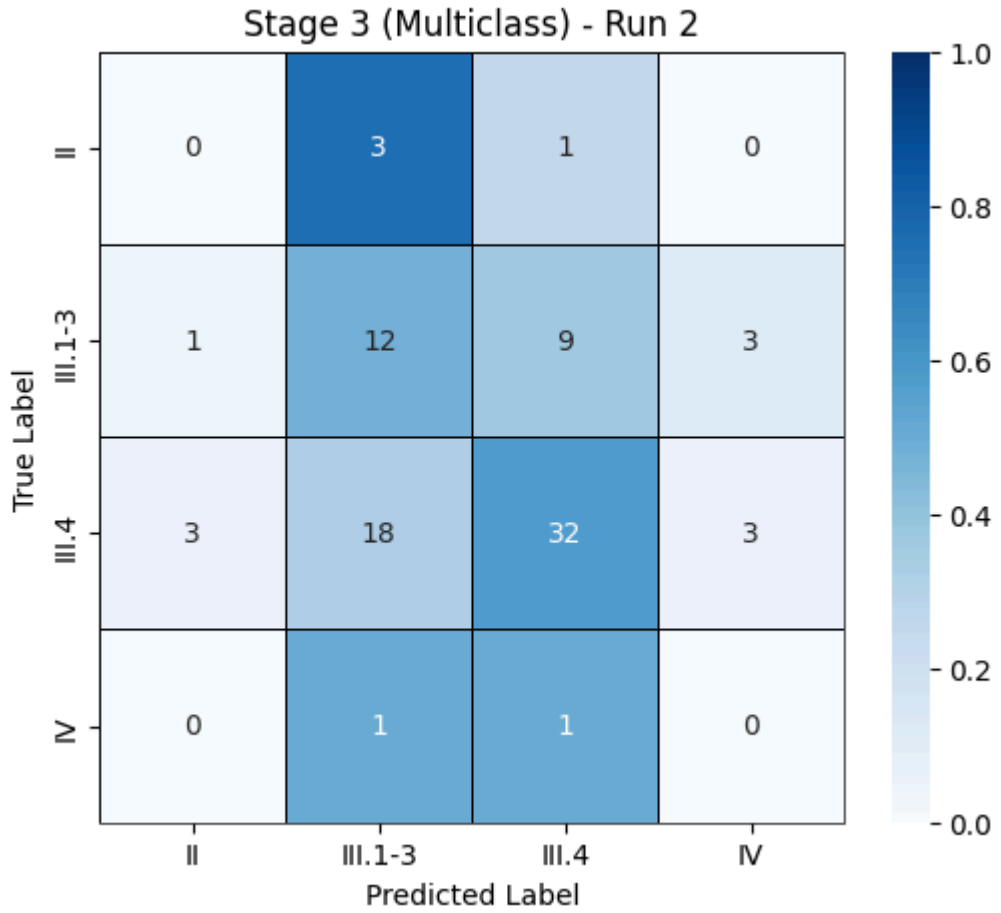

== Soft-Gated Overall (Test Set) - Run 2 ==

	precision	recall	f1-score	support
DC	0.1379	0.3333	0.1951	12
I	0.2727	0.7500	0.4000	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2308	0.2400	0.2353	25
III.4	0.7407	0.3571	0.4819	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.3204	103
macro avg	0.2304	0.2801	0.2187	103
weighted avg	0.4854	0.3204	0.3574	103

Balanced Accuracy: 0.2800793650793651

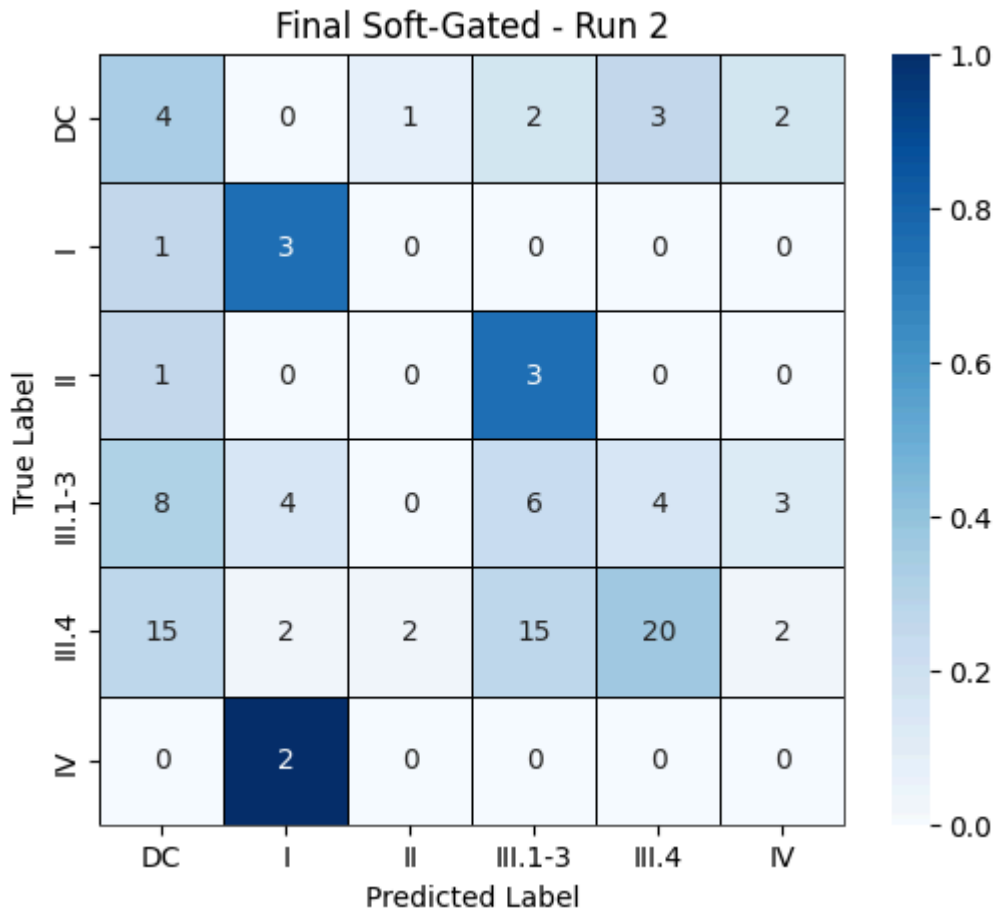

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.1346	0.2800	0.1818	25
I	0.0732	0.8571	0.1348	7
II	0.0000	0.0000	0.0000	8
III.1-3	0.2188	0.1429	0.1728	49
III.4	0.7333	0.1947	0.3077	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.2039	206
macro avg	0.1933	0.2458	0.1329	206
weighted avg	0.4731	0.2039	0.2365	206

Aggregated Balanced Accuracy: 0.2458

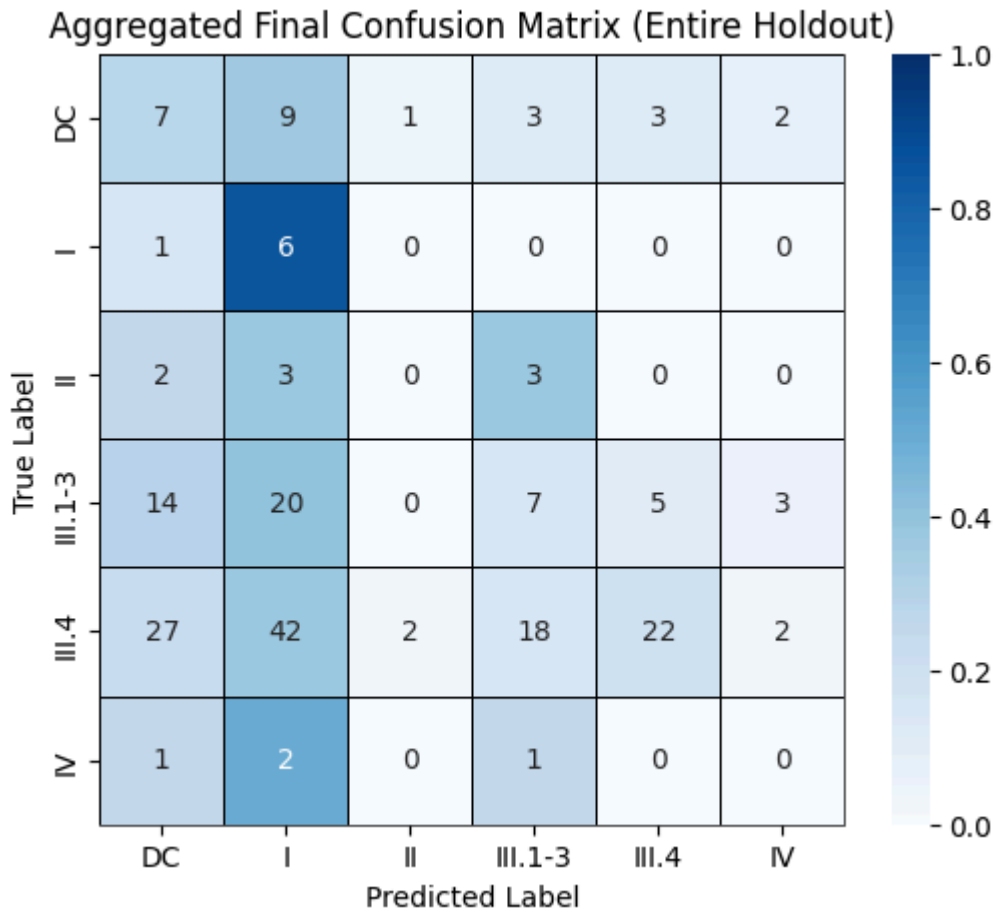

== Average Stage Balanced Accuracies ==
Stage 1 (I vs Rest): 0.8673
Stage 2 (DC vs Rest): 0.5628
Stage 3 (Multiclass): 0.2800
Final (Soft-Gated): 0.2490

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

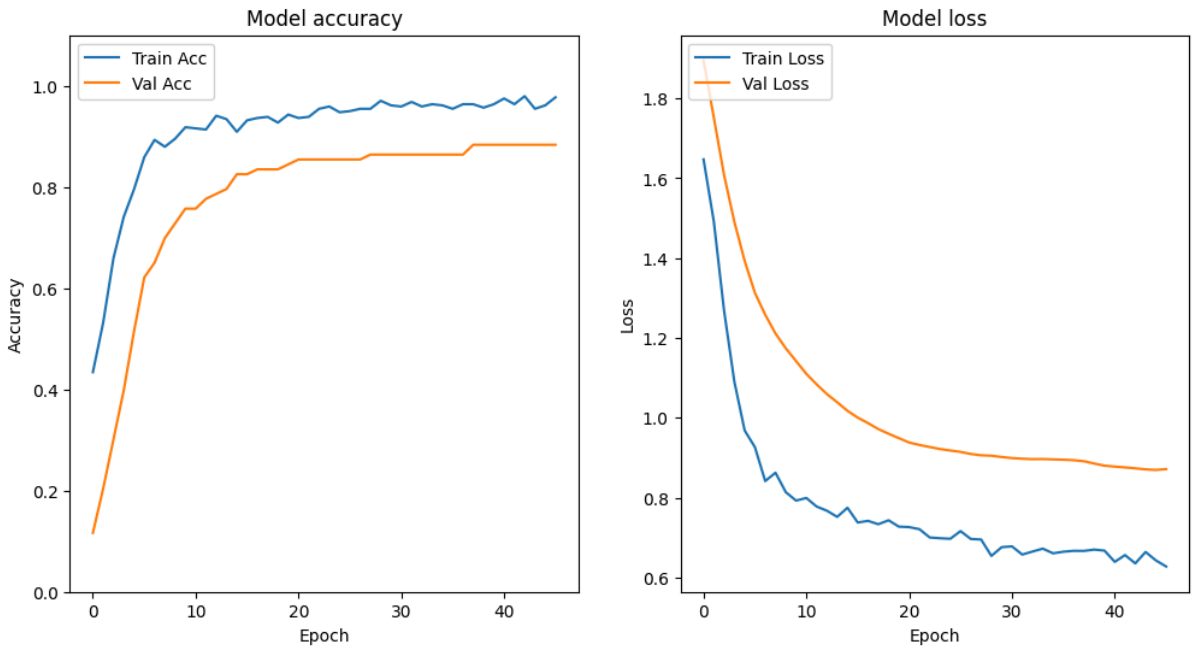

[Threshold Optimization] Best balanced_accuracy: 0.9444 at threshold=0.513
Optimal threshold (Stage 1 (I vs Rest)): 0.513

	precision	recall	f1-score	support
0	1.0000	0.9300	0.9637	100
1	0.3000	1.0000	0.4615	3
accuracy			0.9320	103
macro avg	0.6500	0.9650	0.7126	103
weighted avg	0.9796	0.9320	0.9491	103

Balanced Accuracy: 0.9650000000000001

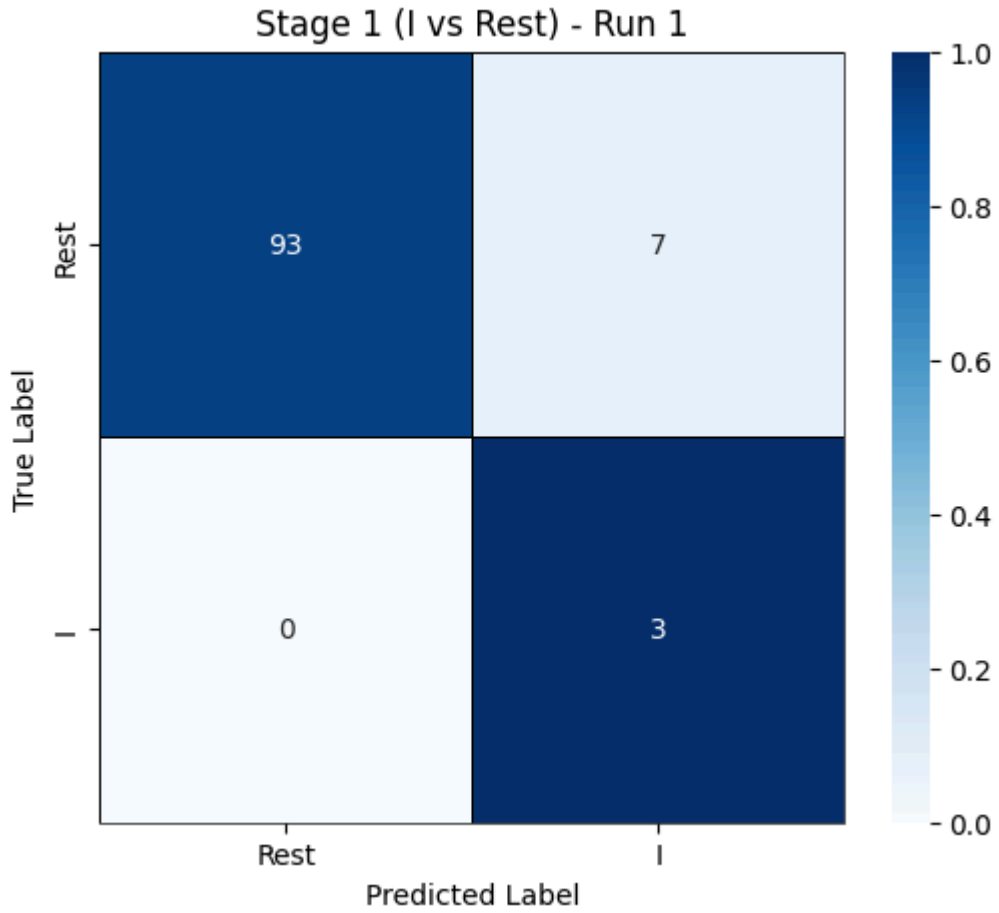

-- Stage 2 (DC vs Rest) --

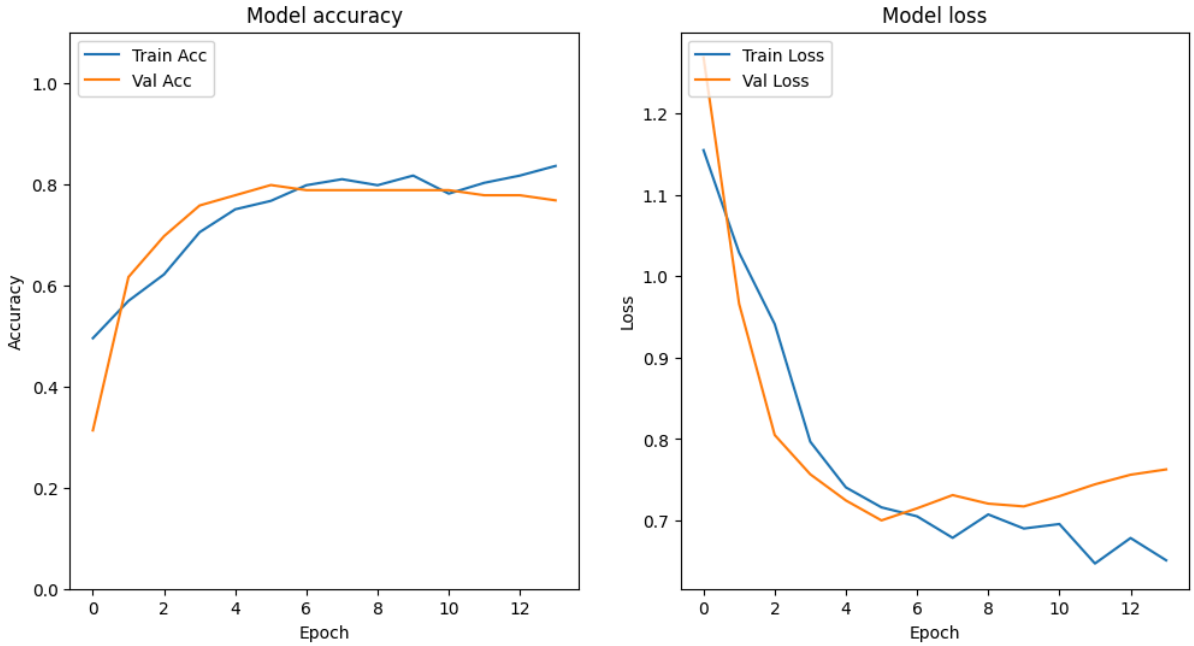

[Threshold Optimization] Best balanced_accuracy: 0.6422 at threshold=0.080
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.9216	0.5402	0.6812	87
1	0.1837	0.6923	0.2903	13
accuracy			0.5600	100
macro avg	0.5526	0.6163	0.4857	100
weighted avg	0.8256	0.5600	0.6304	100

Balanced Accuracy: 0.6162687886825817

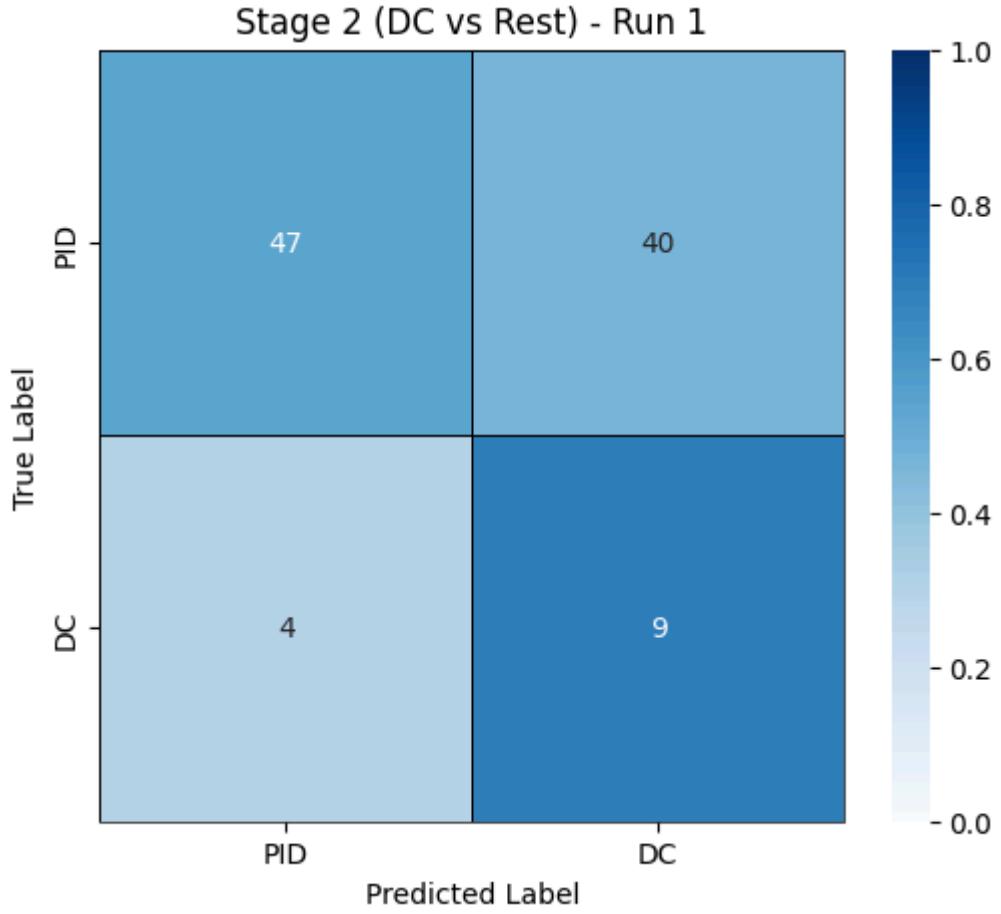

-- Stage 3 (Multiclass) --

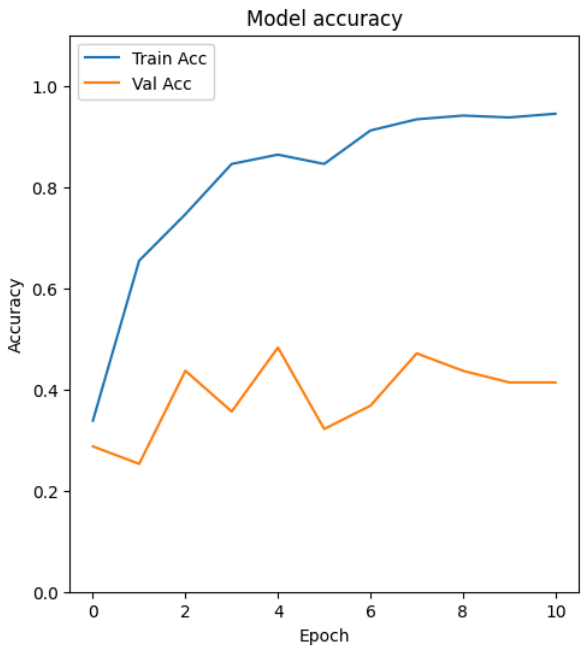

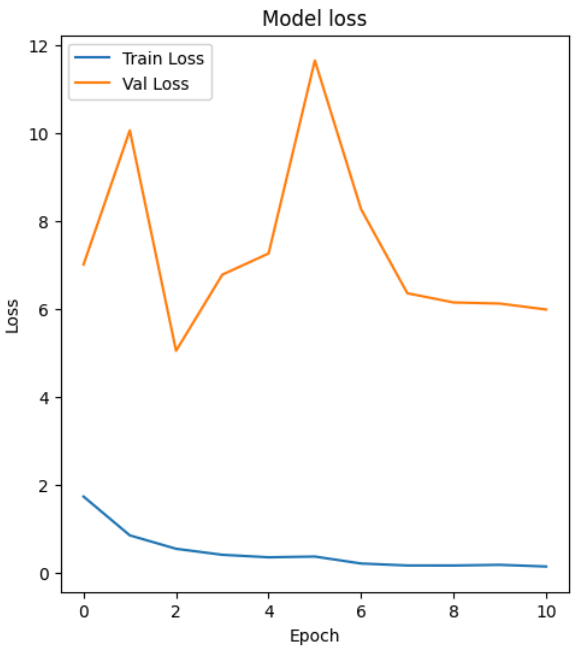

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	4
1	0.2857	0.2500	0.2667	24
2	0.6667	0.5614	0.6095	57
3	0.1333	1.0000	0.2353	2
accuracy			0.4598	87
macro avg	0.2714	0.4529	0.2779	87
weighted avg	0.5187	0.4598	0.4783	87

Balanced Accuracy: 0.45285087719298245

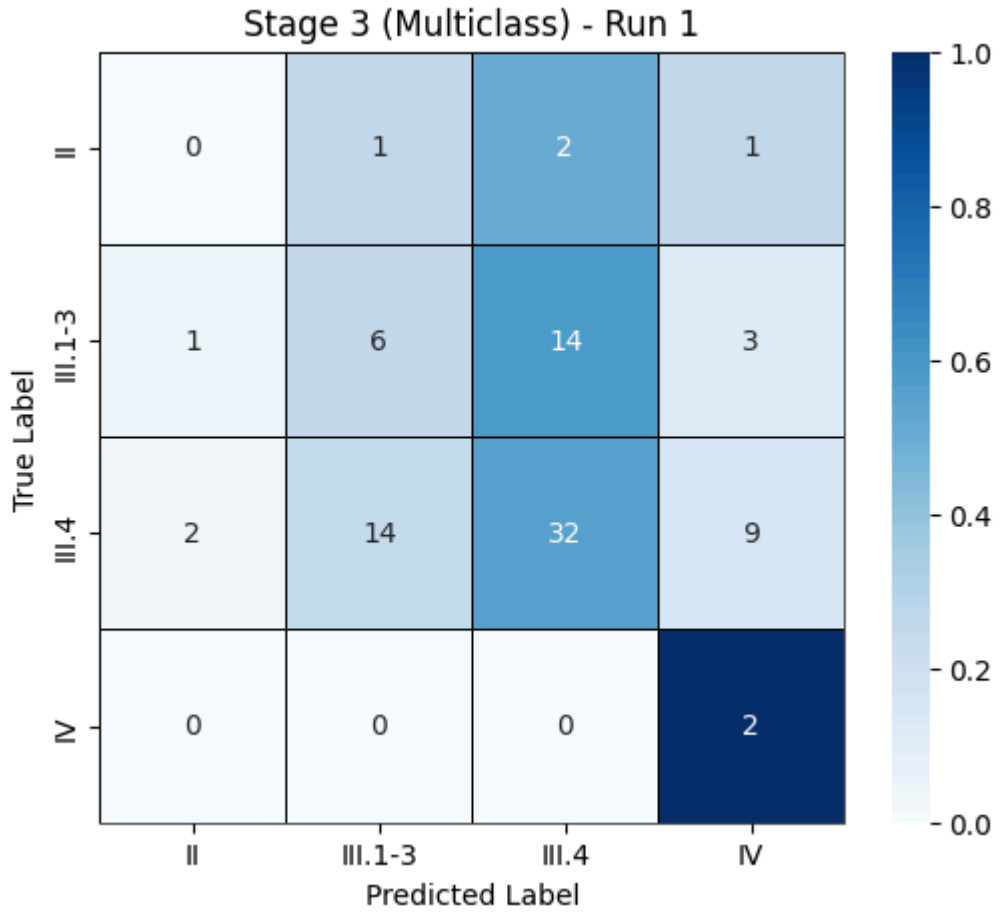

== Soft-Gated Overall (Test Set) - Run 1 ==

	precision	recall	f1-score	support
DC	0.1800	0.6923	0.2857	13
I	0.2500	1.0000	0.4000	3
II	0.0000	0.0000	0.0000	4
III.1-3	0.1000	0.0417	0.0588	24
III.4	0.7692	0.3509	0.4819	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.3204	103
macro avg	0.2165	0.3475	0.2044	103
weighted avg	0.4790	0.3204	0.3281	103

Balanced Accuracy: 0.34747525865946916

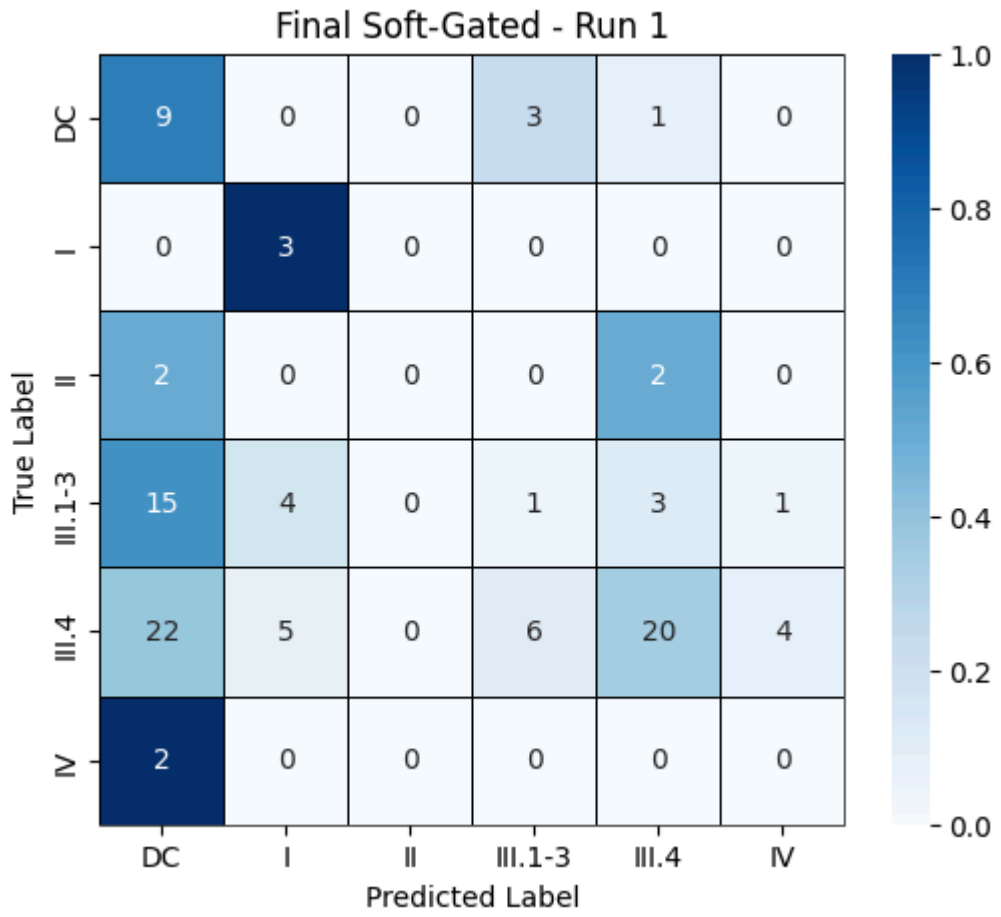

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

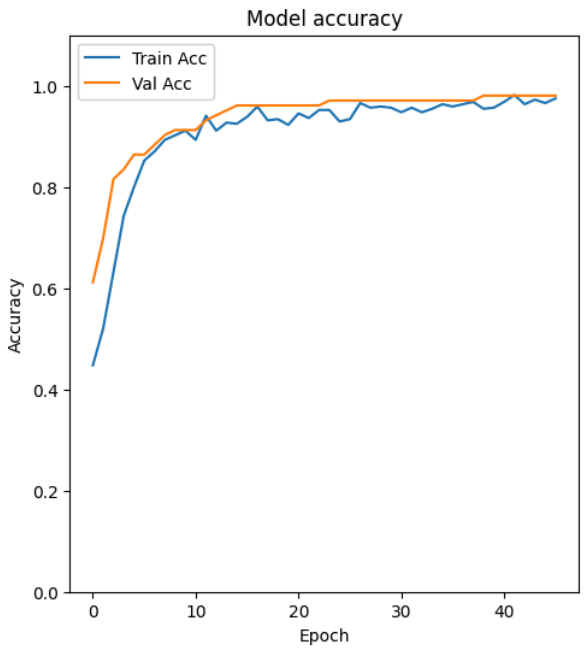

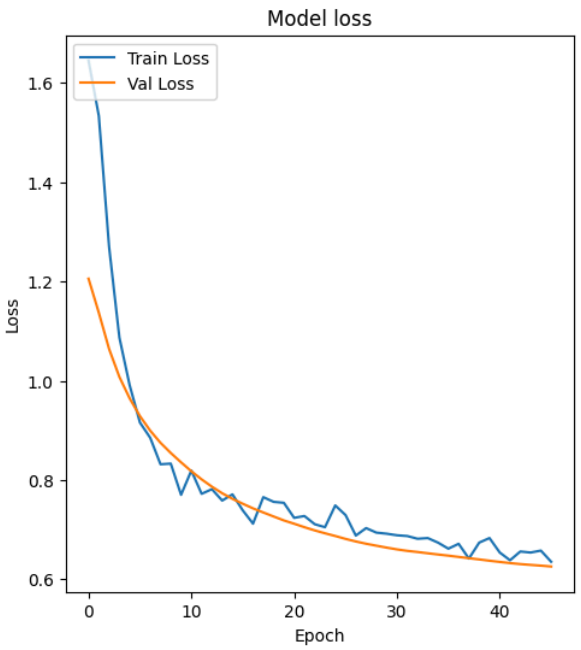

[Threshold Optimization] Best balanced_accuracy: 0.9050 at threshold=0.050
Optimal threshold (Stage 1 (I vs Rest)): 0.050

	precision	recall	f1-score	support
0	0.9870	0.7677	0.8636	99
1	0.1154	0.7500	0.2000	4
accuracy			0.7670	103
macro avg	0.5512	0.7588	0.5318	103
weighted avg	0.9532	0.7670	0.8379	103

Balanced Accuracy: 0.7588383838383839

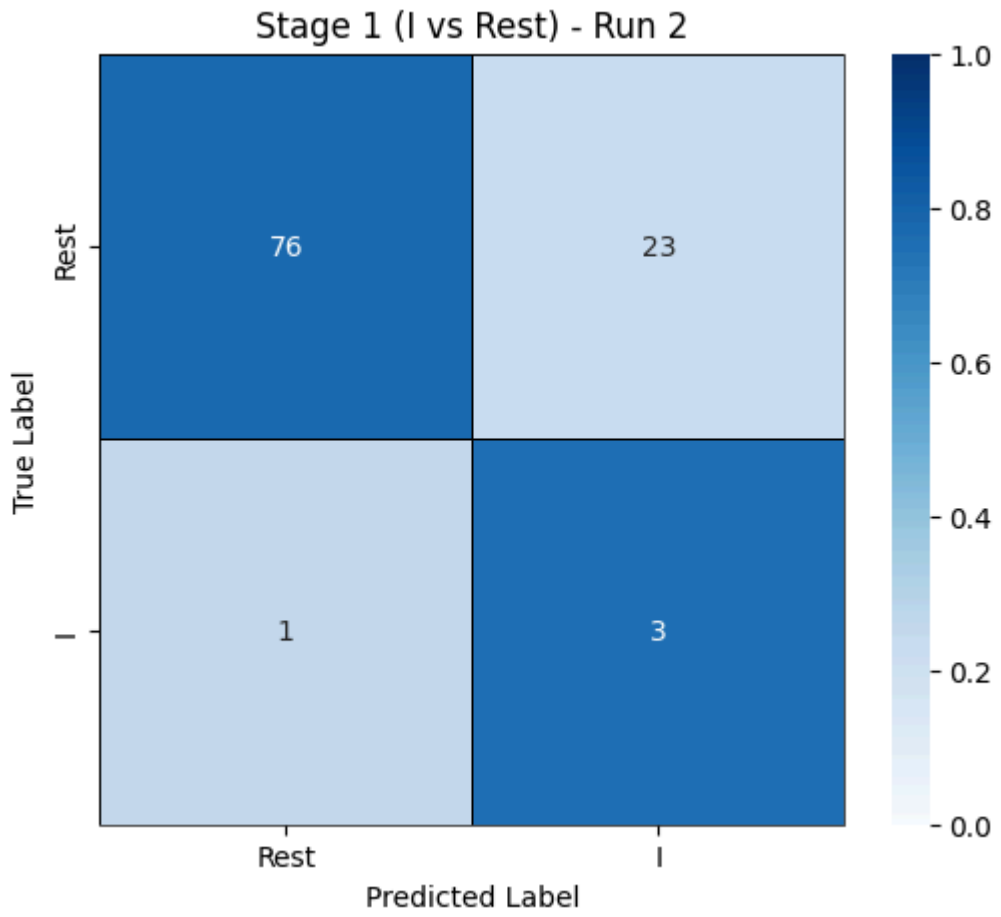

-- Stage 2 (DC vs Rest) --

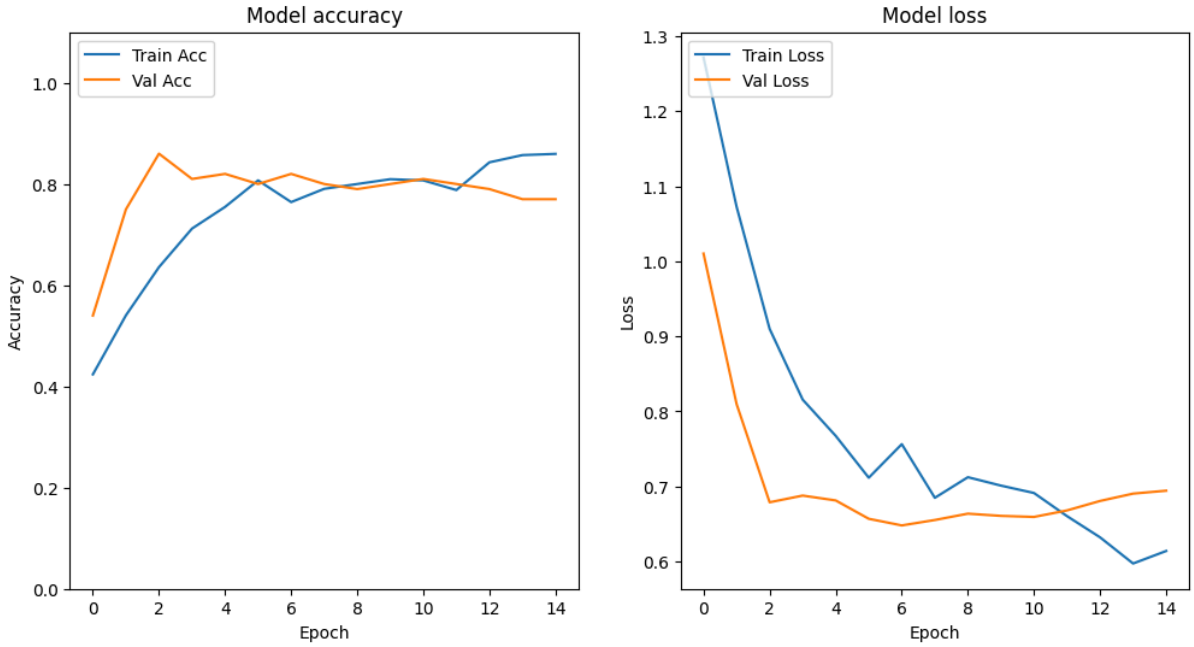

[Threshold Optimization] Best balanced_accuracy: 0.7104 at threshold=0.116
Optimal threshold (Stage 2 (DC vs Rest)): 0.200
precision recall f1-score support

0	0.8592	0.7011	0.7722	87
1	0.0714	0.1667	0.1000	12
accuracy			0.6364	99
macro avg	0.4653	0.4339	0.4361	99
weighted avg	0.7637	0.6364	0.6907	99

Balanced Accuracy: 0.4339080459770115

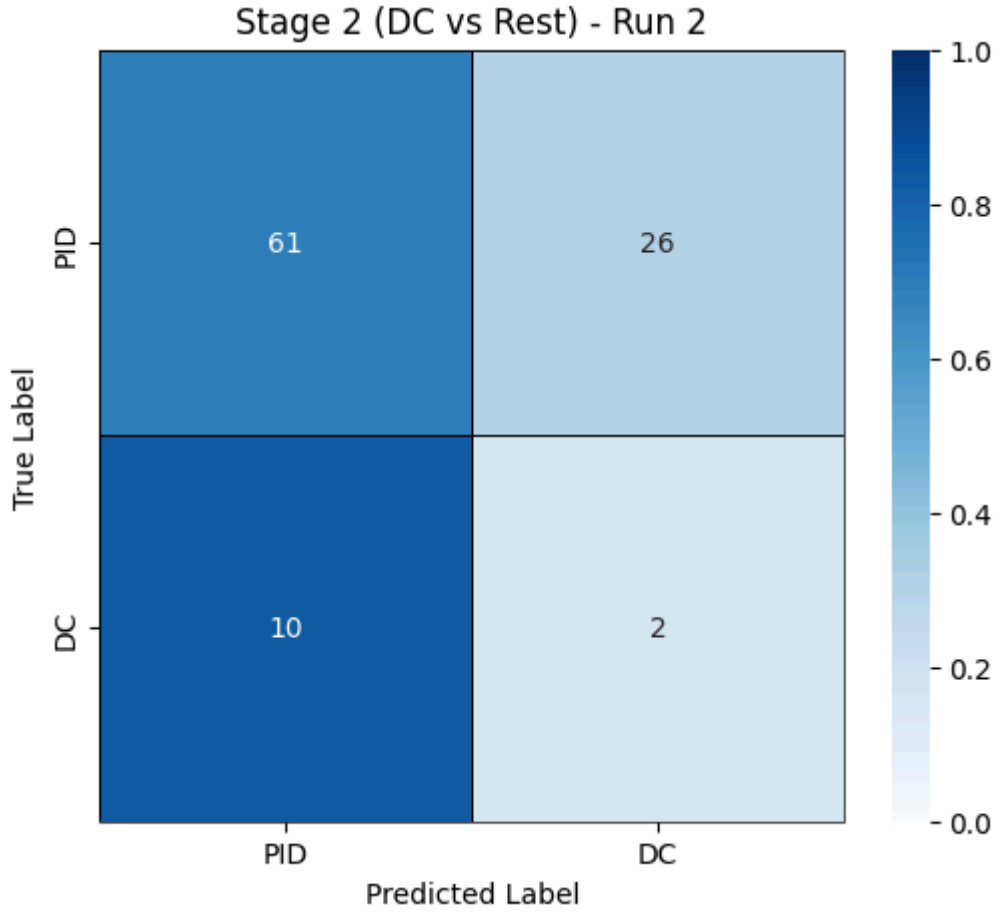

-- Stage 3 (Multiclass) --

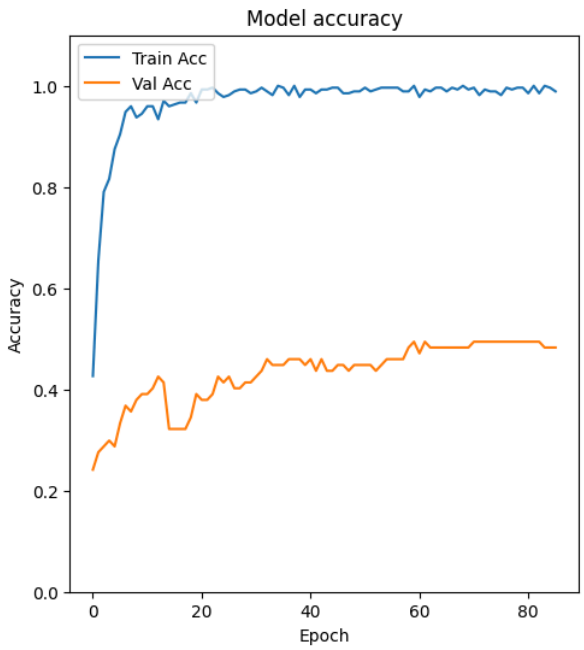

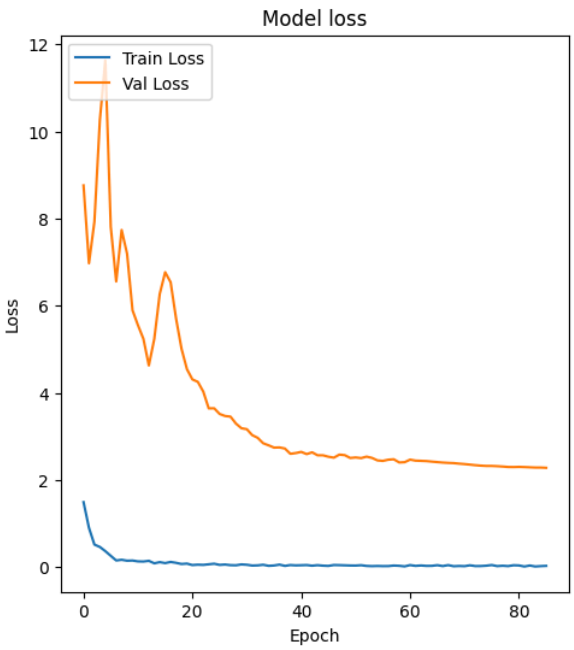

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	4
1	0.3235	0.4400	0.3729	25
2	0.6792	0.6429	0.6606	56
3	0.0000	0.0000	0.0000	2
accuracy			0.5402	87
macro avg	0.2507	0.2707	0.2584	87
weighted avg	0.5302	0.5402	0.5323	87

Balanced Accuracy: 0.27071428571428574

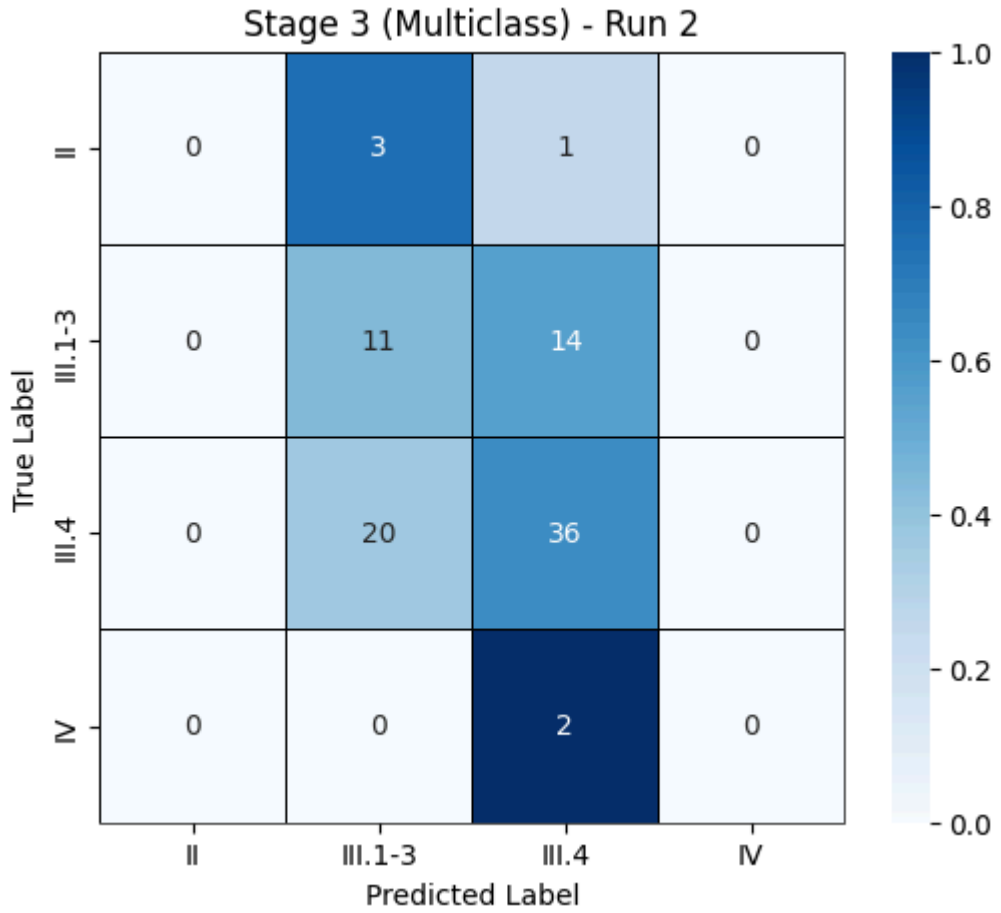

== Soft-Gated Overall (Test Set) - Run 2 ==

	precision	recall	f1-score	support
DC	0.0870	0.1667	0.1143	12
I	0.0690	1.0000	0.1290	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.1250	0.0400	0.0606	25
III.4	0.6429	0.1607	0.2571	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.1553	103
macro avg	0.1540	0.2279	0.0935	103
weighted avg	0.3927	0.1553	0.1728	103

Balanced Accuracy: 0.22789682539682543

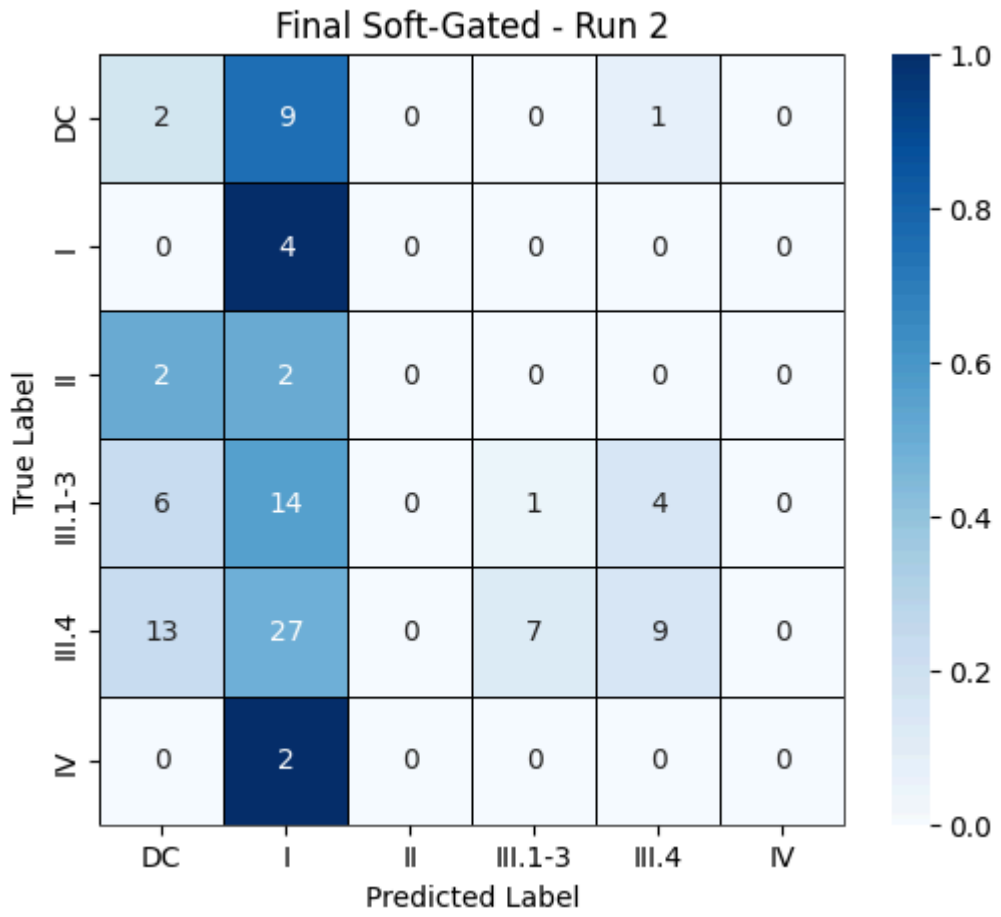

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.1507	0.4400	0.2245	25
I	0.1000	1.0000	0.1818	7
II	0.0000	0.0000	0.0000	8
III.1-3	0.1111	0.0408	0.0597	49
III.4	0.7250	0.2566	0.3791	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.2379	206
macro avg	0.1811	0.2896	0.1408	206
weighted avg	0.4458	0.2379	0.2556	206

Aggregated Balanced Accuracy: 0.2896

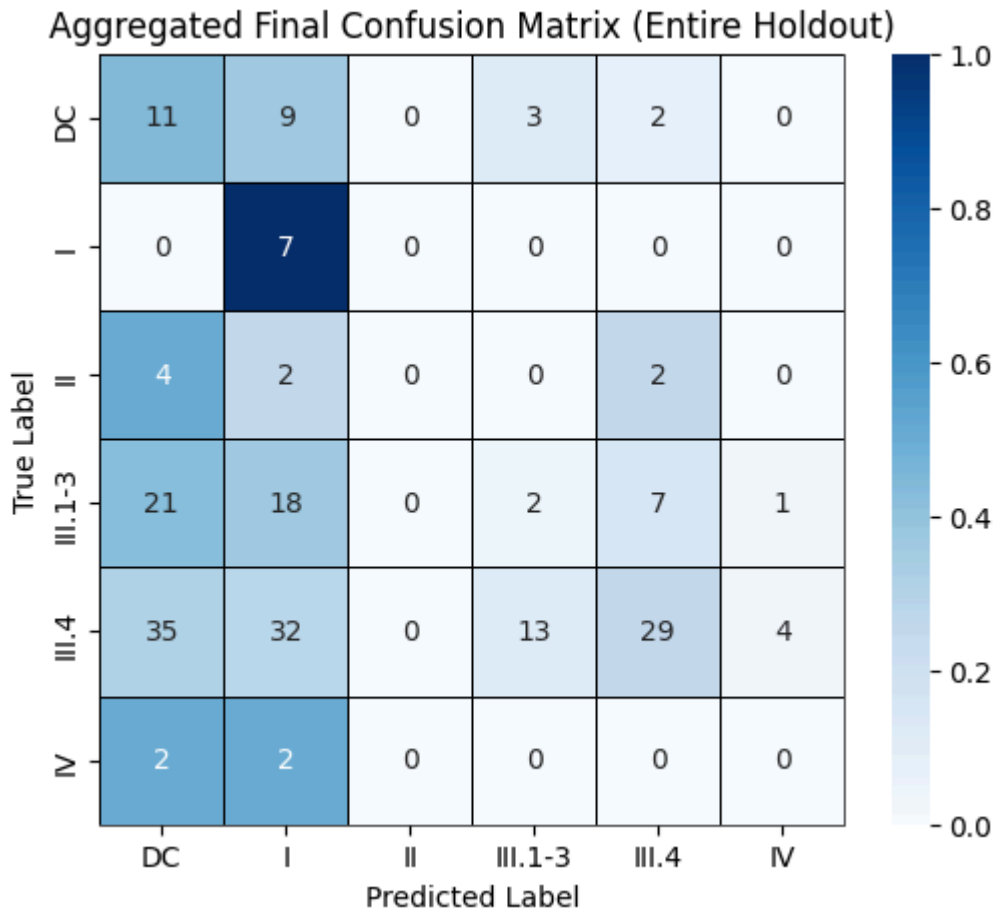

== Average Stage Balanced Accuracies ==
Stage 1 (I vs Rest): 0.8619
Stage 2 (DC vs Rest): 0.5251
Stage 3 (Multiclass): 0.3618
Final (Soft-Gated): 0.2877

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

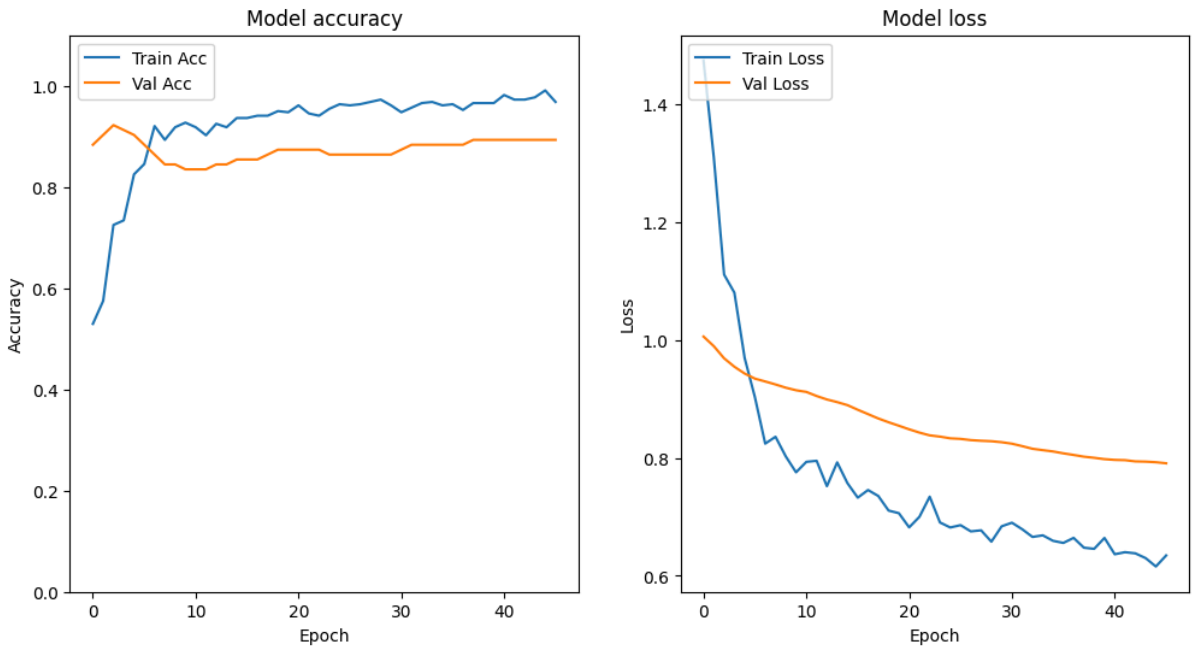

[Threshold Optimization] Best balanced_accuracy: 0.9091 at threshold=0.291
Optimal threshold (Stage 1 (I vs Rest)): 0.291

	precision	recall	f1-score	support
0	1.0000	0.8400	0.9130	100
1	0.1579	1.0000	0.2727	3
accuracy			0.8447	103
macro avg	0.5789	0.9200	0.5929	103
weighted avg	0.9755	0.8447	0.8944	103

Balanced Accuracy: 0.9199999999999999

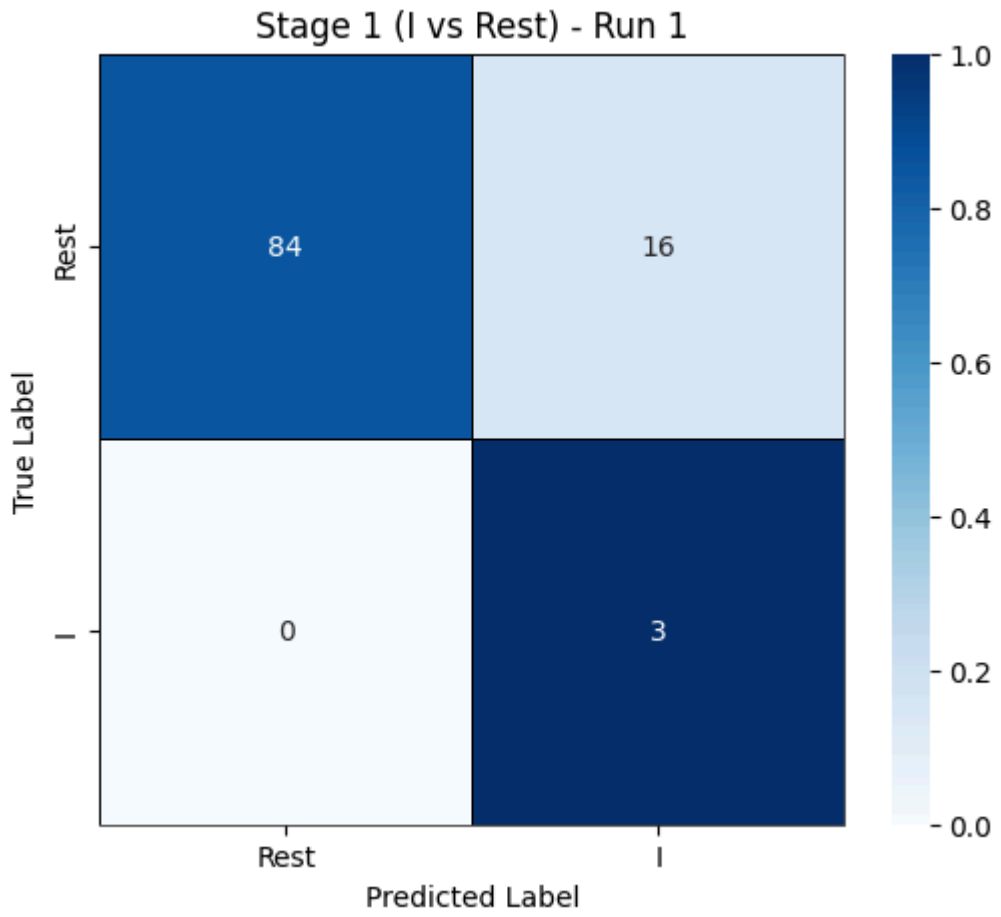

-- Stage 2 (DC vs Rest) --

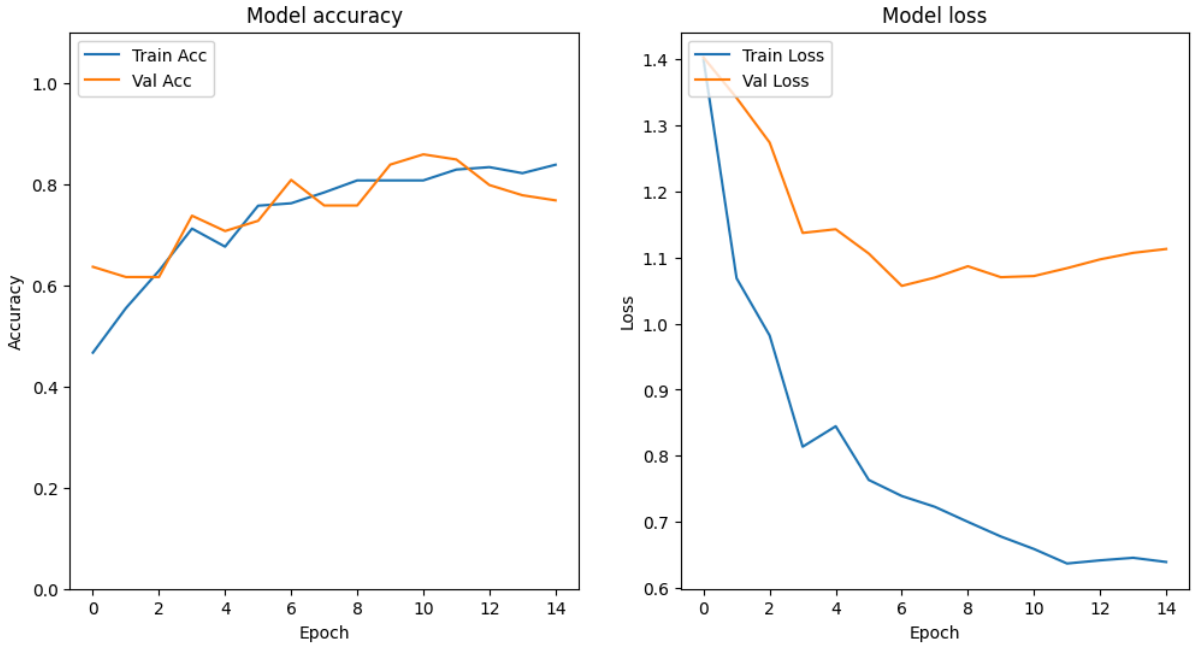

[Threshold Optimization] Best balanced_accuracy: 0.6279 at threshold=0.090
Optimal threshold (Stage 2 (DC vs Rest)): 0.200
precision recall f1-score support

0	0.9444	0.5862	0.7234	87
1	0.2174	0.7692	0.3390	13
accuracy			0.6100	100
macro avg	0.5809	0.6777	0.5312	100
weighted avg	0.8499	0.6100	0.6734	100

Balanced Accuracy: 0.6777188328912467

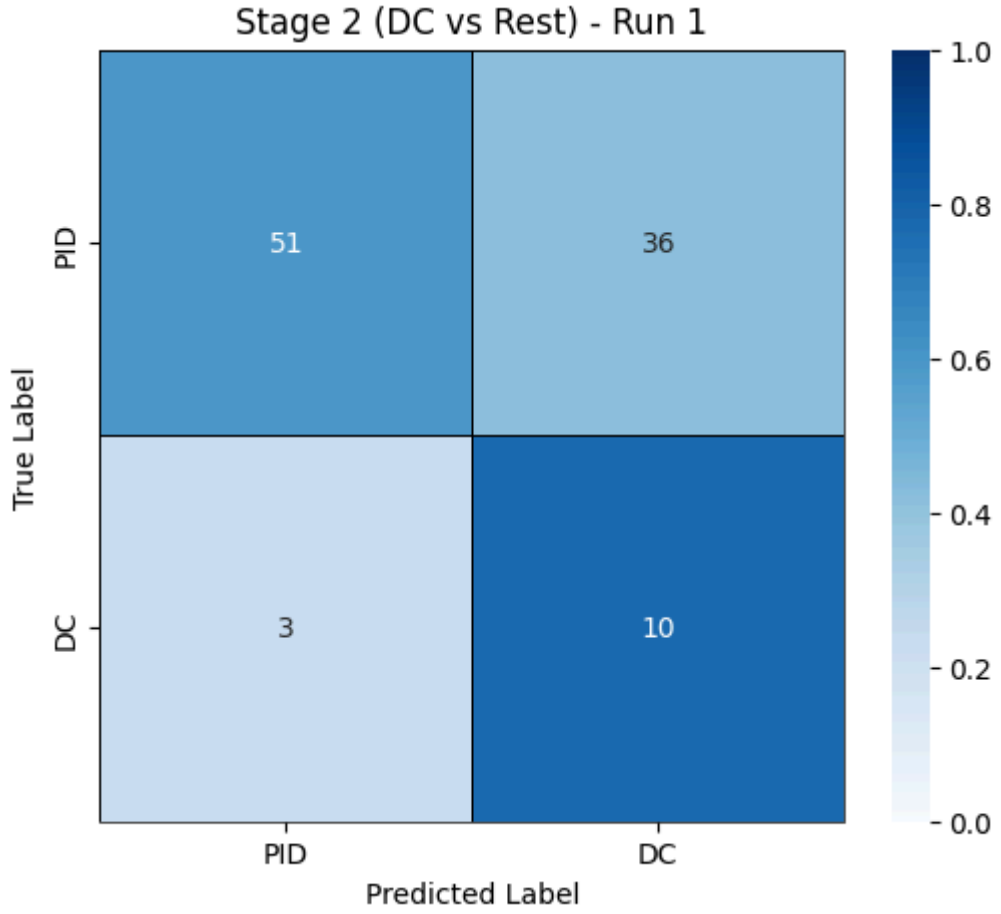

-- Stage 3 (Multiclass) --

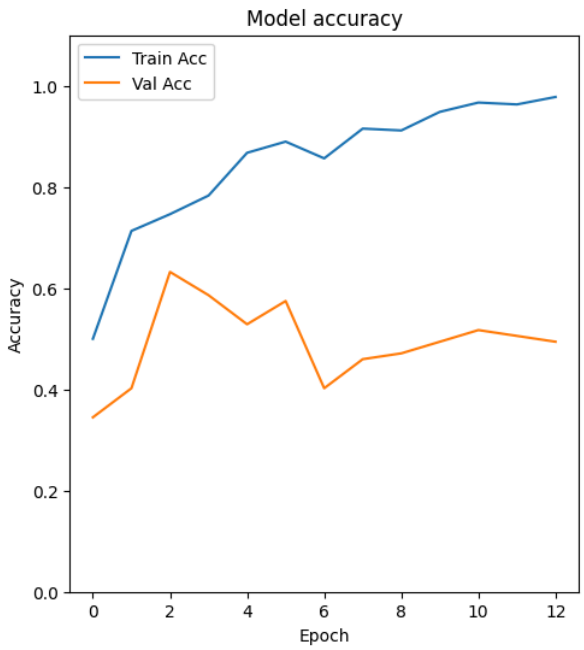

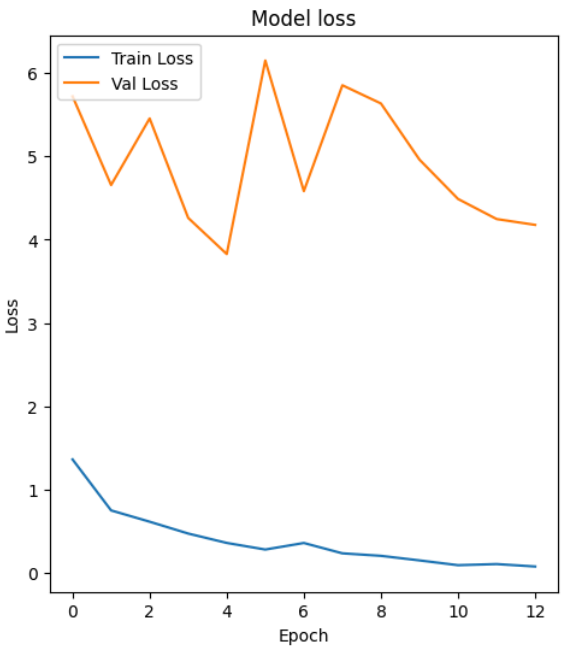

	precision	recall	f1-score	support
0	1.0000	0.2500	0.4000	4
1	0.3667	0.4583	0.4074	24
2	0.7273	0.7018	0.7143	57
3	0.0000	0.0000	0.0000	2
accuracy			0.5977	87
macro avg	0.5235	0.3525	0.3804	87
weighted avg	0.6236	0.5977	0.5988	87

Balanced Accuracy: 0.3525219298245614

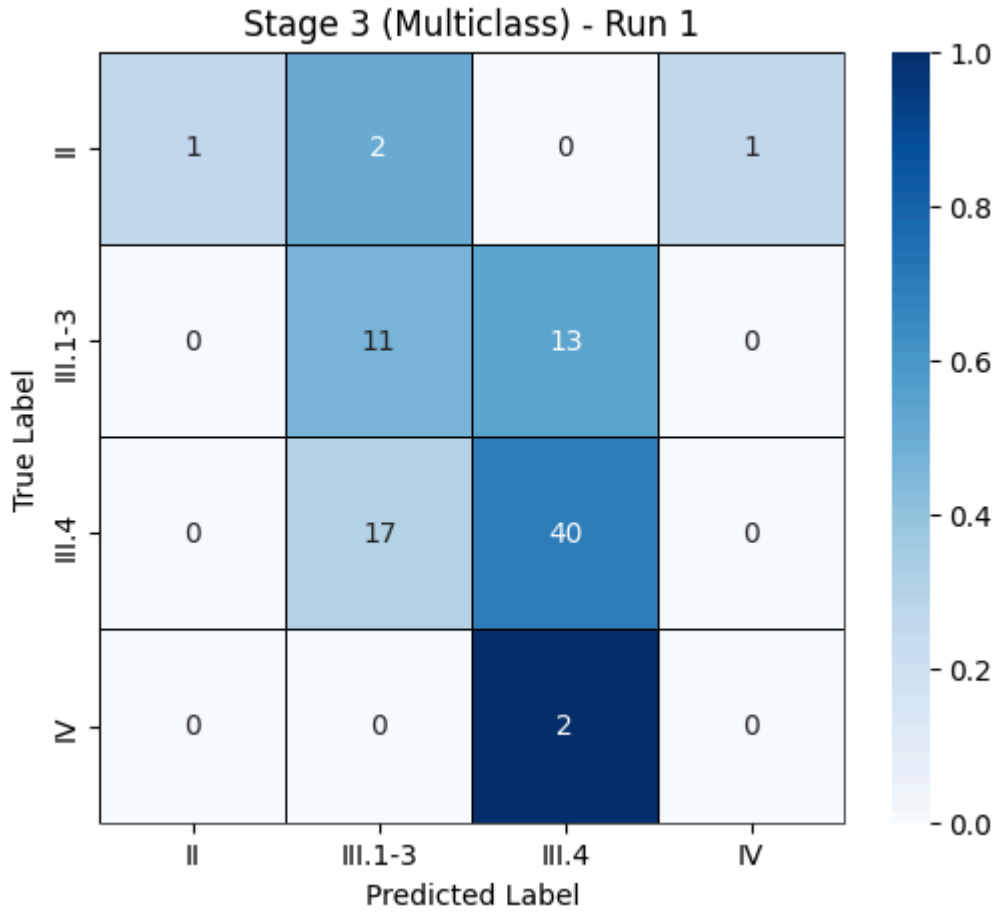

== Soft-Gated Overall (Test Set) - Run 1 ==

	precision	recall	f1-score	support
DC	0.2222	0.7692	0.3448	13
I	0.1429	1.0000	0.2500	3
II	0.0000	0.0000	0.0000	4
III.1-3	0.1667	0.0833	0.1111	24
III.4	0.7200	0.3158	0.4390	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.3204	103
macro avg	0.2086	0.3614	0.1908	103
weighted avg	0.4695	0.3204	0.3196	103

Balanced Accuracy: 0.36139226270805214

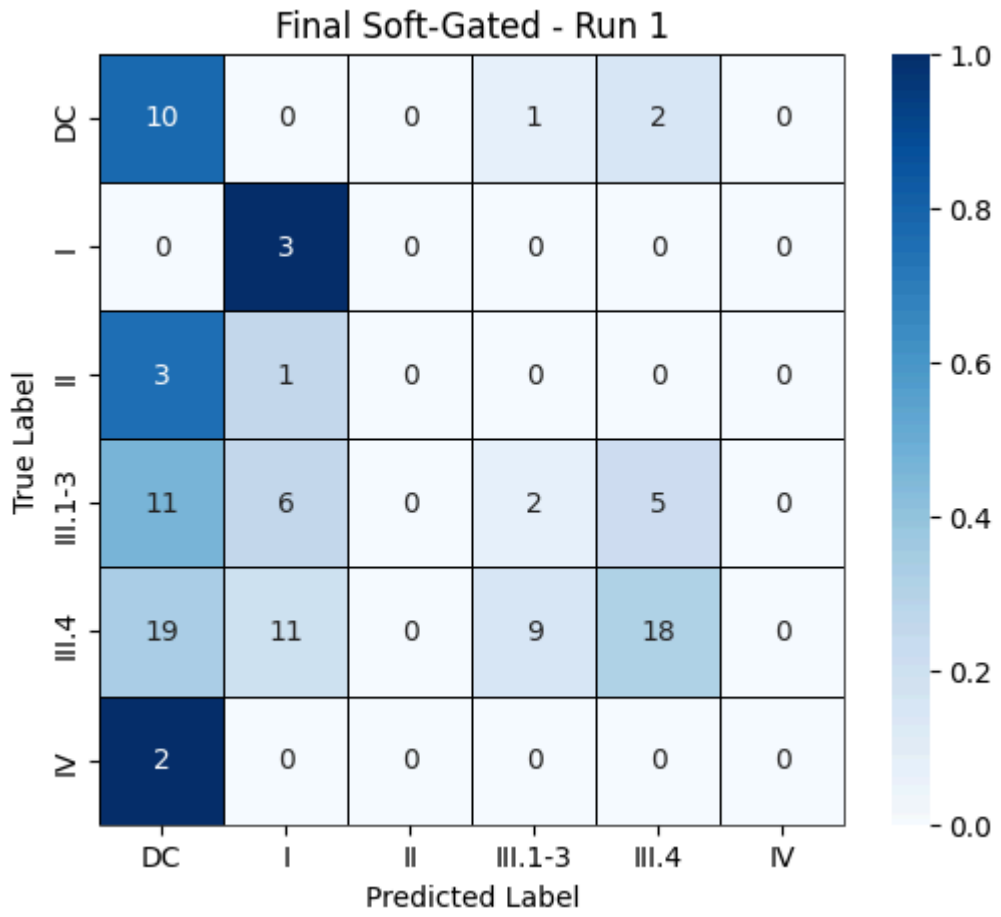

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

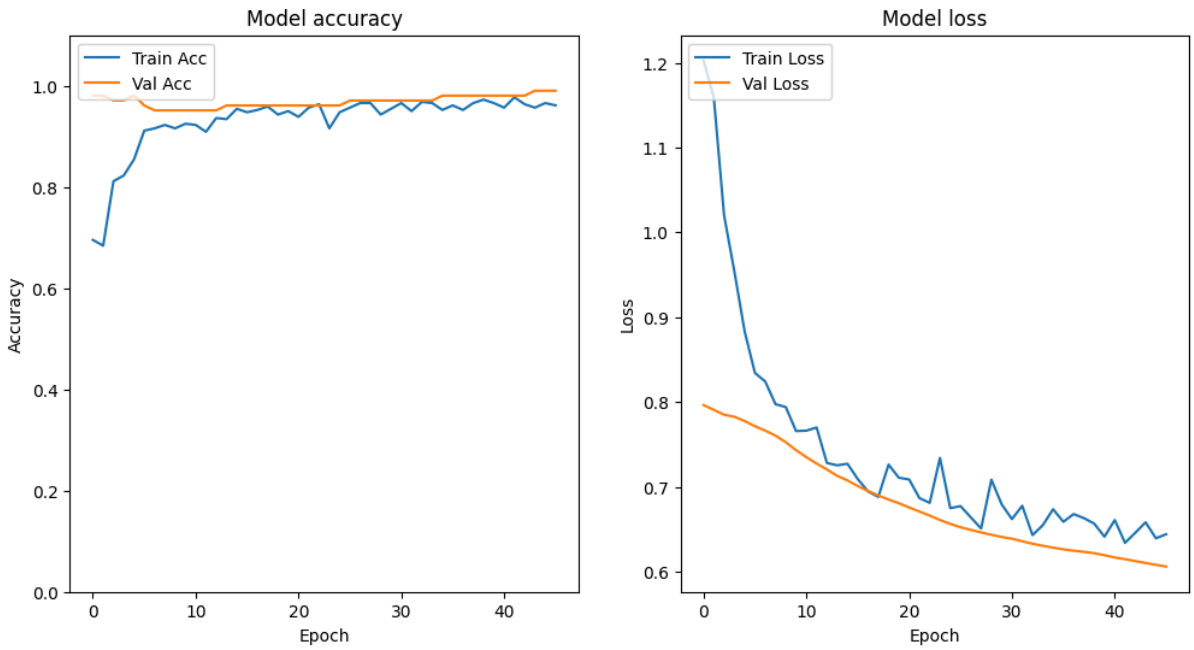

[Threshold Optimization] Best balanced_accuracy: 0.9850 at threshold=0.281
Optimal threshold (Stage 1 (I vs Rest)): 0.281

	precision	recall	f1-score	support
0	0.9889	0.8990	0.9418	99
1	0.2308	0.7500	0.3529	4
accuracy			0.8932	103
macro avg	0.6098	0.8245	0.6474	103
weighted avg	0.9594	0.8932	0.9189	103

Balanced Accuracy: 0.8244949494949495

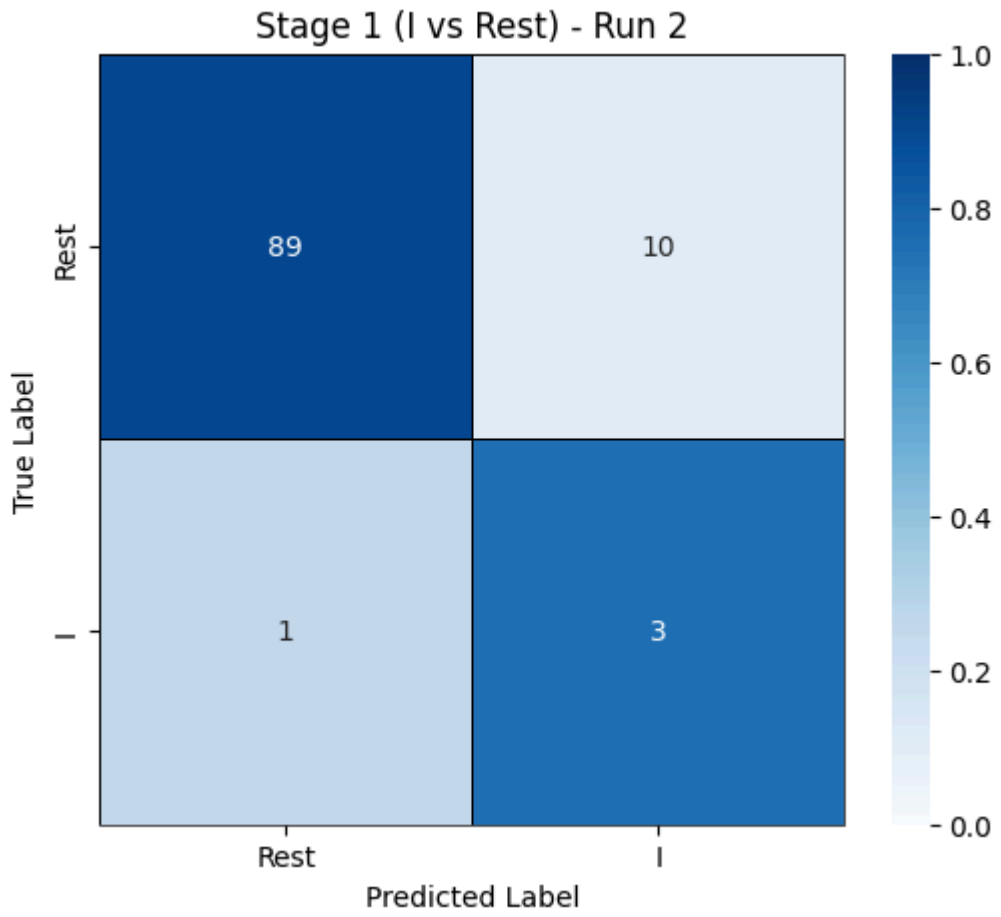

-- Stage 2 (DC vs Rest) --

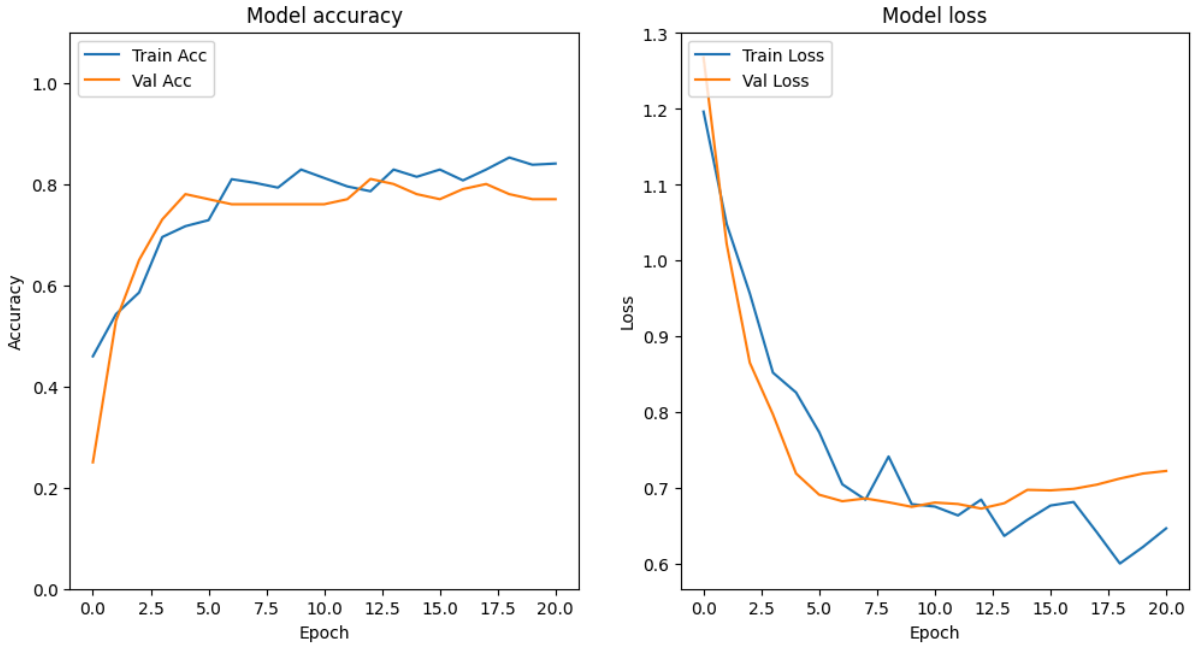

[Threshold Optimization] Best balanced_accuracy: 0.7259 at threshold=0.075
Optimal threshold (Stage 2 (DC vs Rest)): 0.200
precision recall f1-score support

0	0.8649	0.7356	0.7950	87
1	0.0800	0.1667	0.1081	12
accuracy			0.6667	99
macro avg	0.4724	0.4511	0.4516	99
weighted avg	0.7697	0.6667	0.7118	99

Balanced Accuracy: 0.4511494252873563

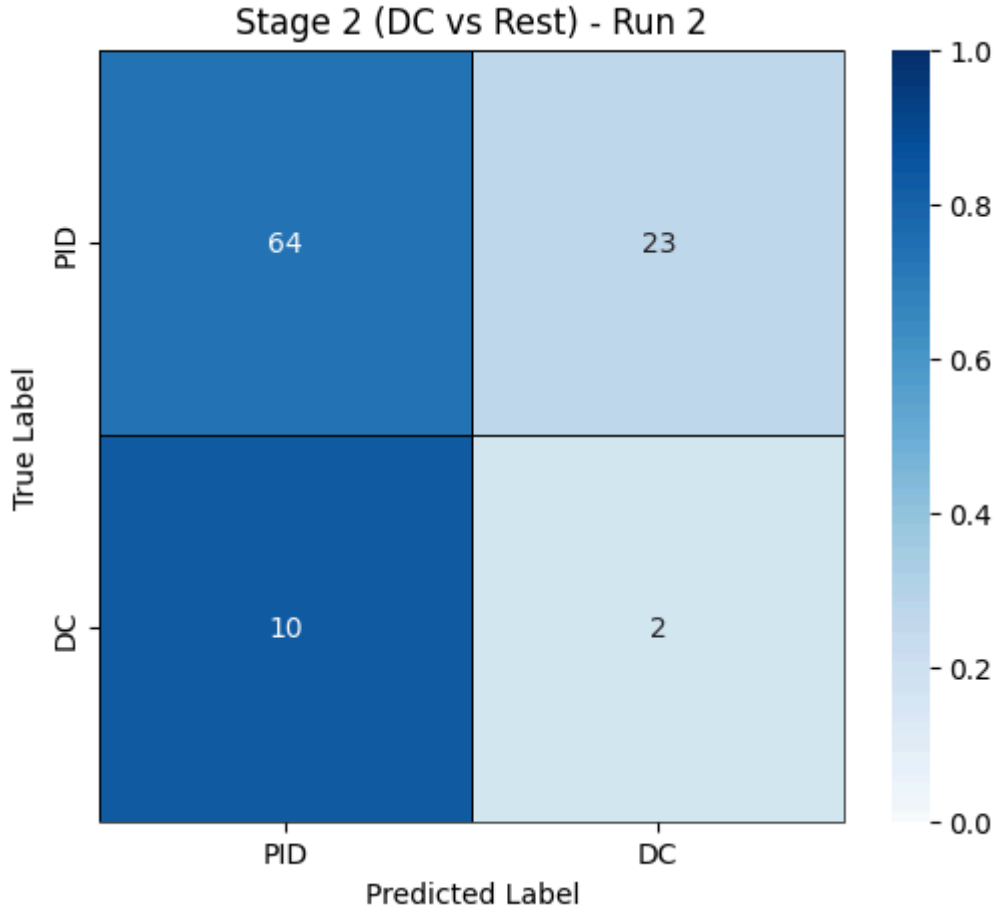

-- Stage 3 (Multiclass) --

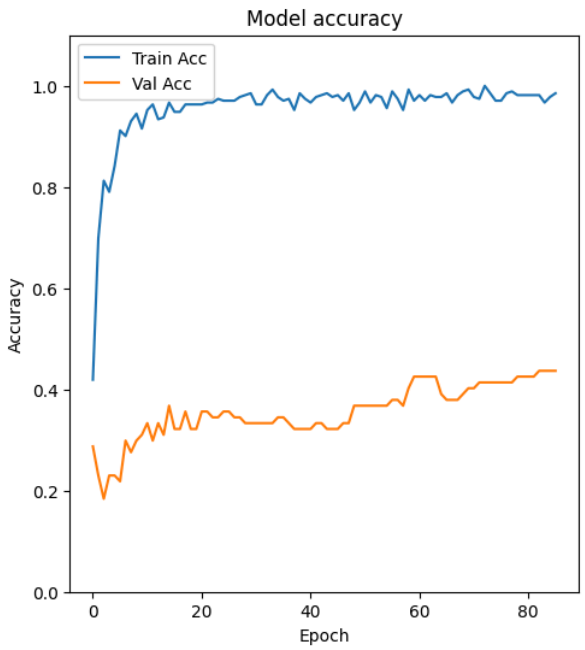

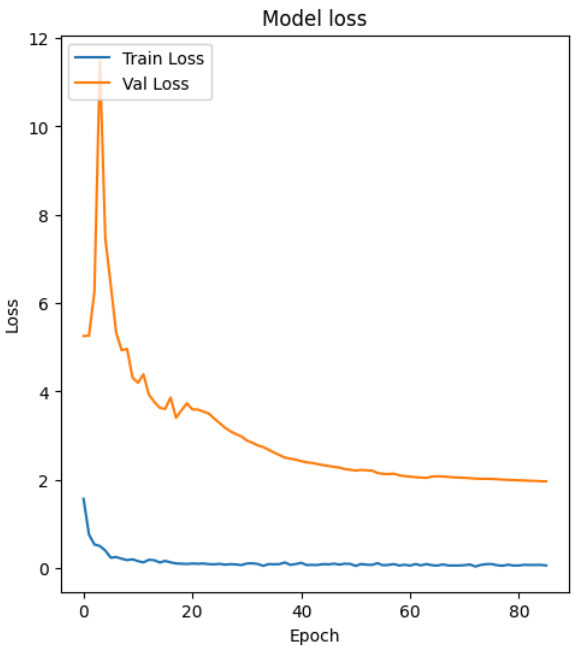

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	4
1	0.2889	0.5200	0.3714	25
2	0.6750	0.4821	0.5625	56
3	0.0000	0.0000	0.0000	2
accuracy			0.4598	87
macro avg	0.2410	0.2505	0.2335	87
weighted avg	0.5175	0.4598	0.4688	87

Balanced Accuracy: 0.2505357142857143

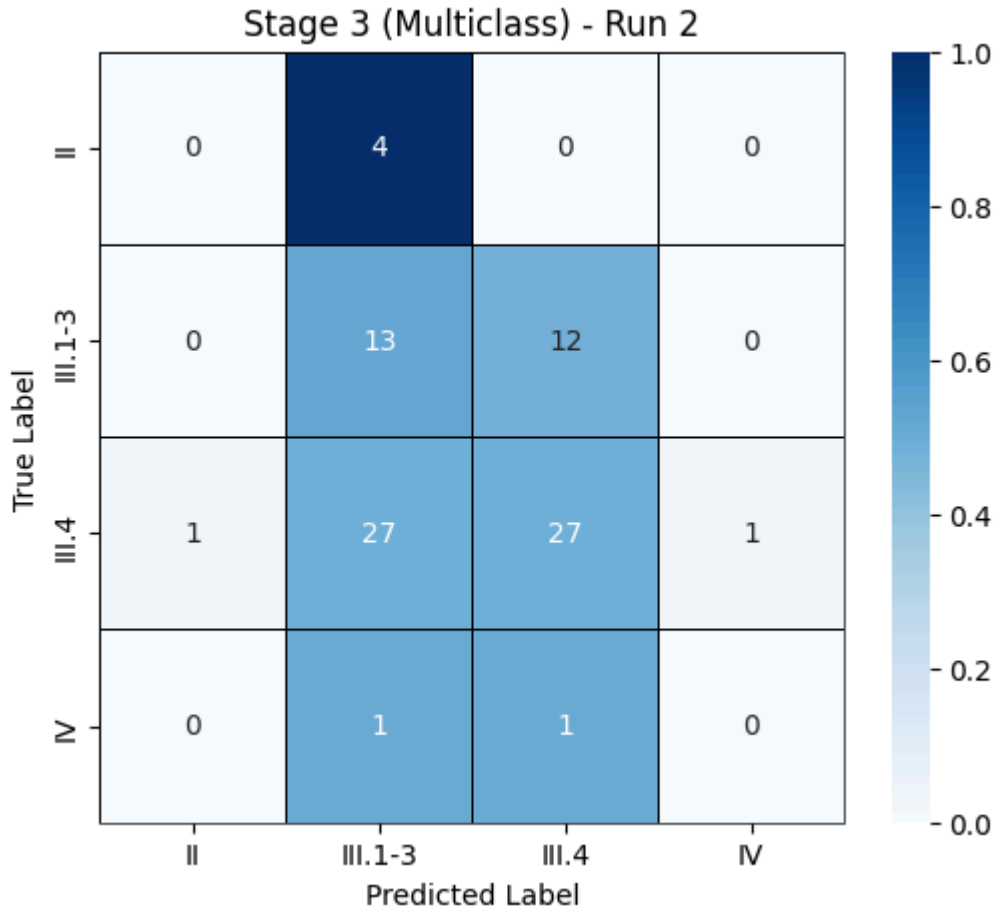

== Soft-Gated Overall (Test Set) - Run 2 ==

	precision	recall	f1-score	support
DC	0.0769	0.1667	0.1053	12
I	0.2000	0.7500	0.3158	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2353	0.3200	0.2712	25
III.4	0.5926	0.2857	0.3855	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.2816	103
macro avg	0.1841	0.2537	0.1796	103
weighted avg	0.3960	0.2816	0.3000	103

Balanced Accuracy: 0.25373015873015875

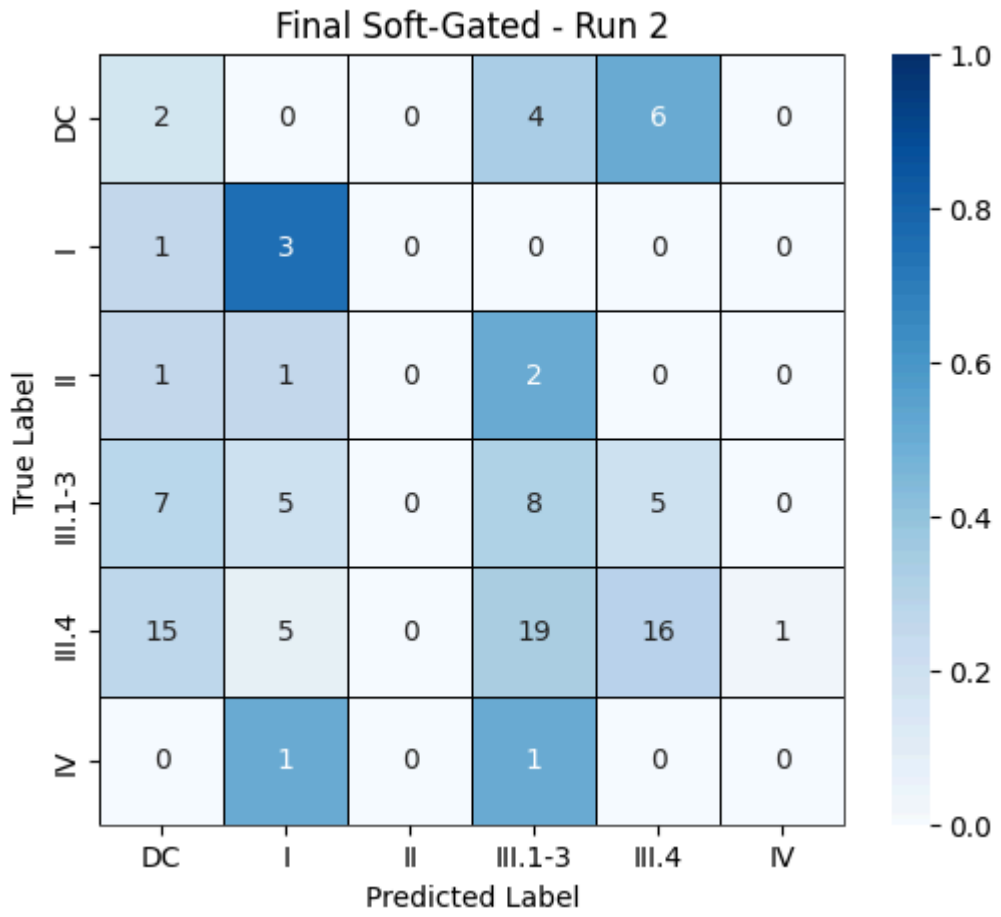

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.1690	0.4800	0.2500	25
I	0.1667	0.8571	0.2791	7
II	0.0000	0.0000	0.0000	8
III.1-3	0.2174	0.2041	0.2105	49
III.4	0.6538	0.3009	0.4121	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3010	206
macro avg	0.2012	0.3070	0.1920	206
weighted avg	0.4365	0.3010	0.3160	206

Aggregated Balanced Accuracy: 0.3070

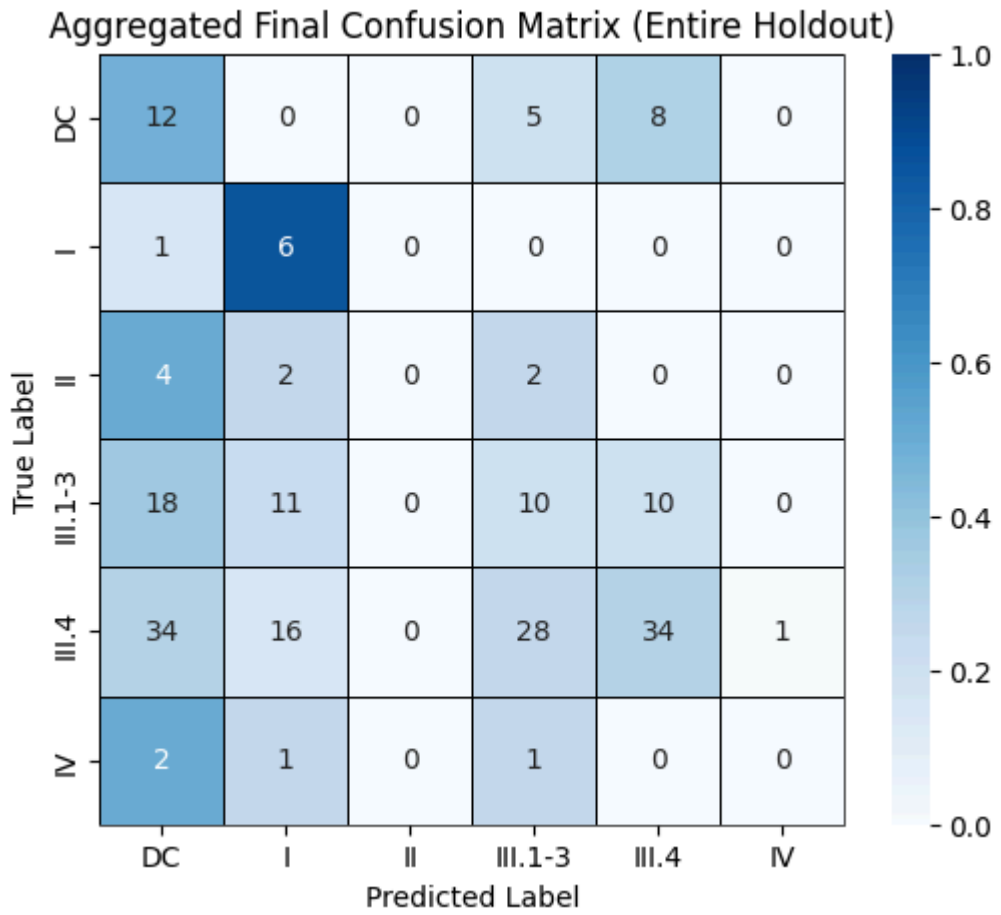

== Average Stage Balanced Accuracies ==
Stage 1 (I vs Rest): 0.8722
Stage 2 (DC vs Rest): 0.5644
Stage 3 (Multiclass): 0.3015
Final (Soft-Gated): 0.3076

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

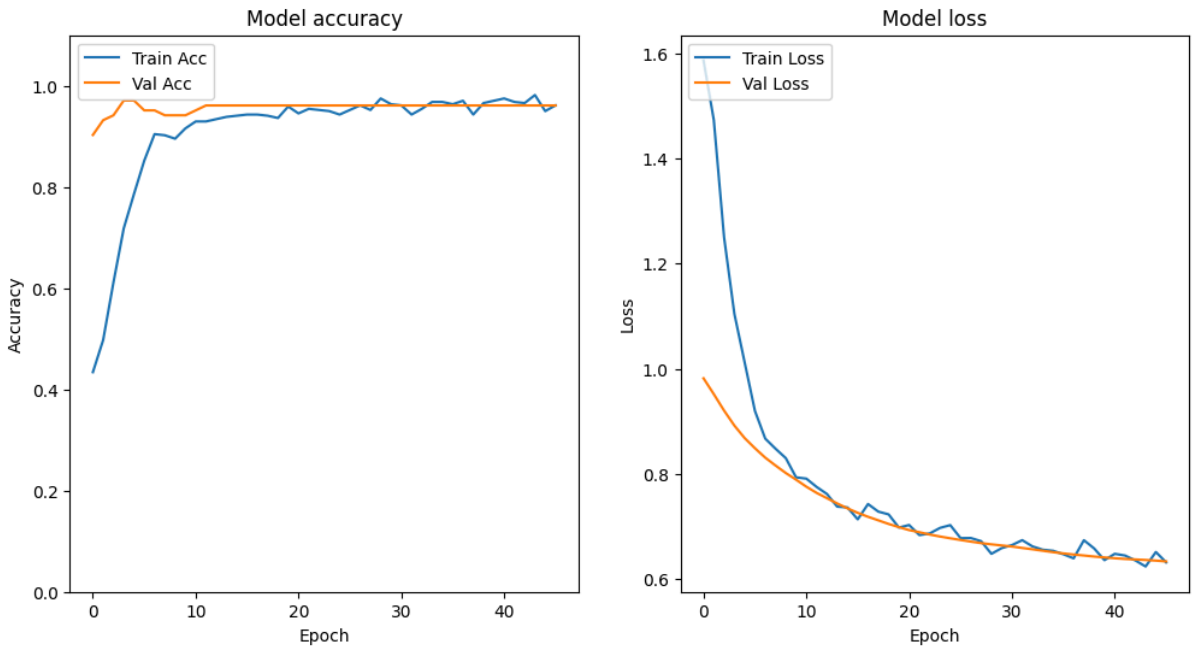

[Threshold Optimization] Best balanced_accuracy: 0.8838 at threshold=0.050
Optimal threshold (Stage 1 (I vs Rest)): 0.050

	precision	recall	f1-score	support
0	1.0000	0.8700	0.9305	100
1	0.1875	1.0000	0.3158	3
accuracy			0.8738	103
macro avg	0.5938	0.9350	0.6231	103
weighted avg	0.9763	0.8738	0.9126	103

Balanced Accuracy: 0.935

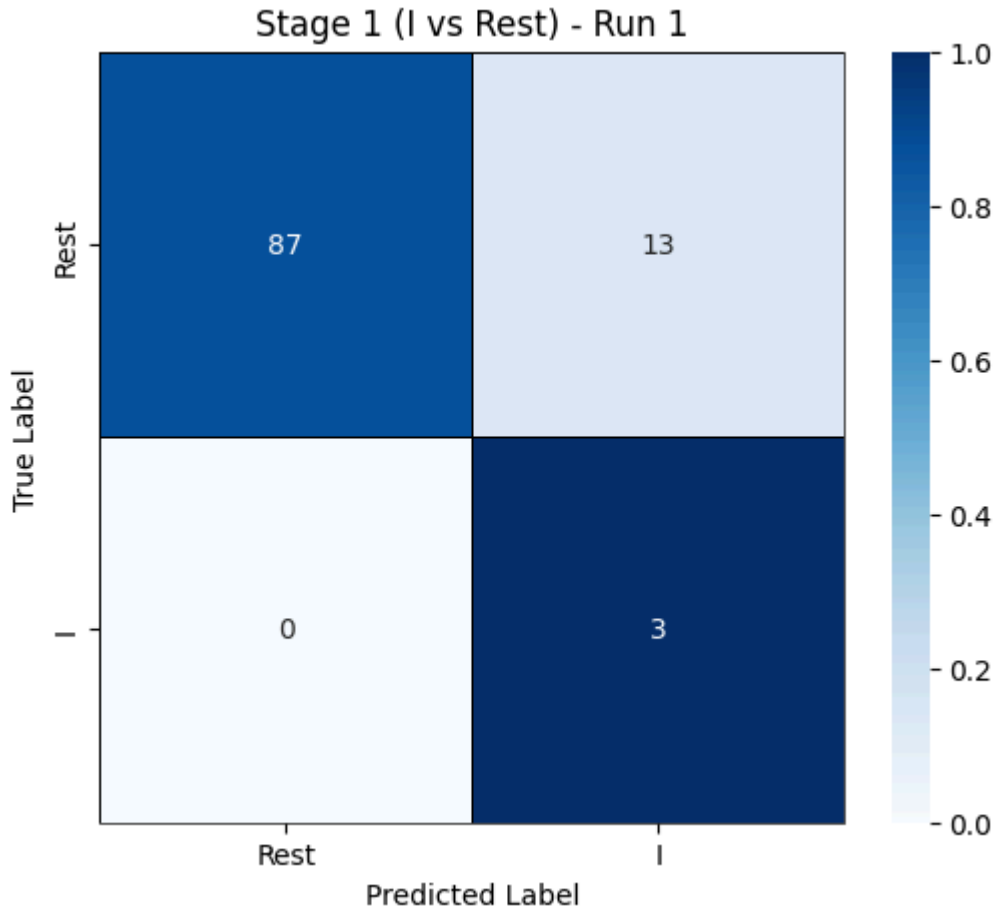

-- Stage 2 (DC vs Rest) --

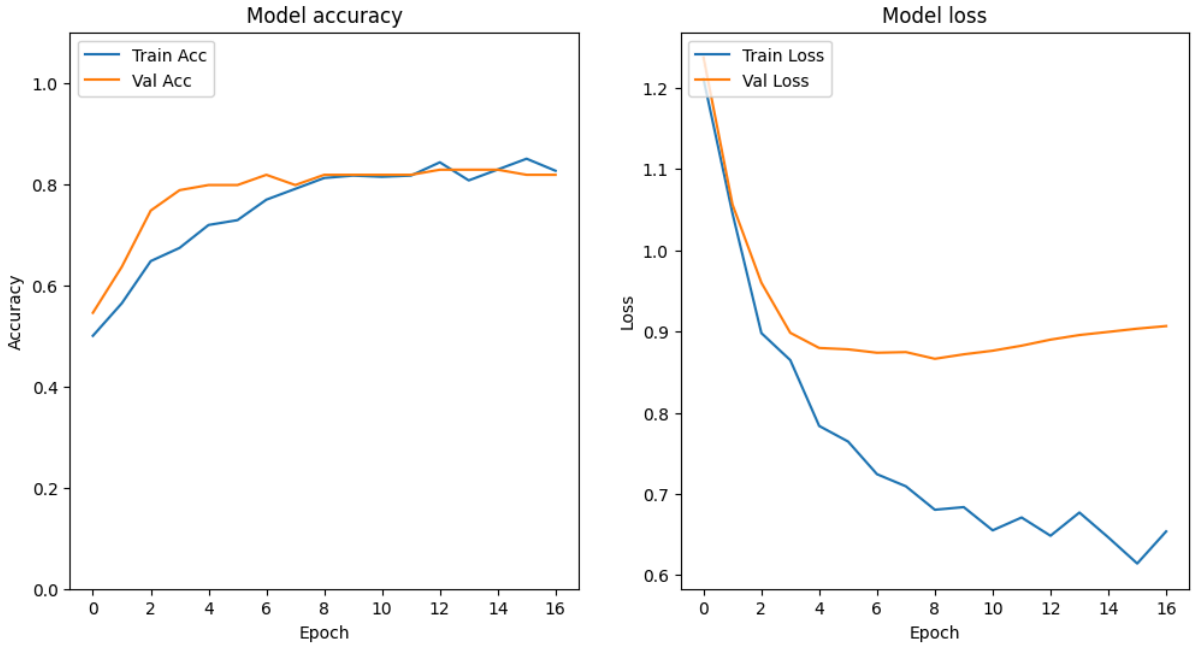

[Threshold Optimization] Best balanced_accuracy: 0.6351 at threshold=0.161
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.9219	0.6782	0.7815	87
1	0.2222	0.6154	0.3265	13
accuracy			0.6700	100
macro avg	0.5720	0.6468	0.5540	100
weighted avg	0.8309	0.6700	0.7223	100

Balanced Accuracy: 0.6467727674624226

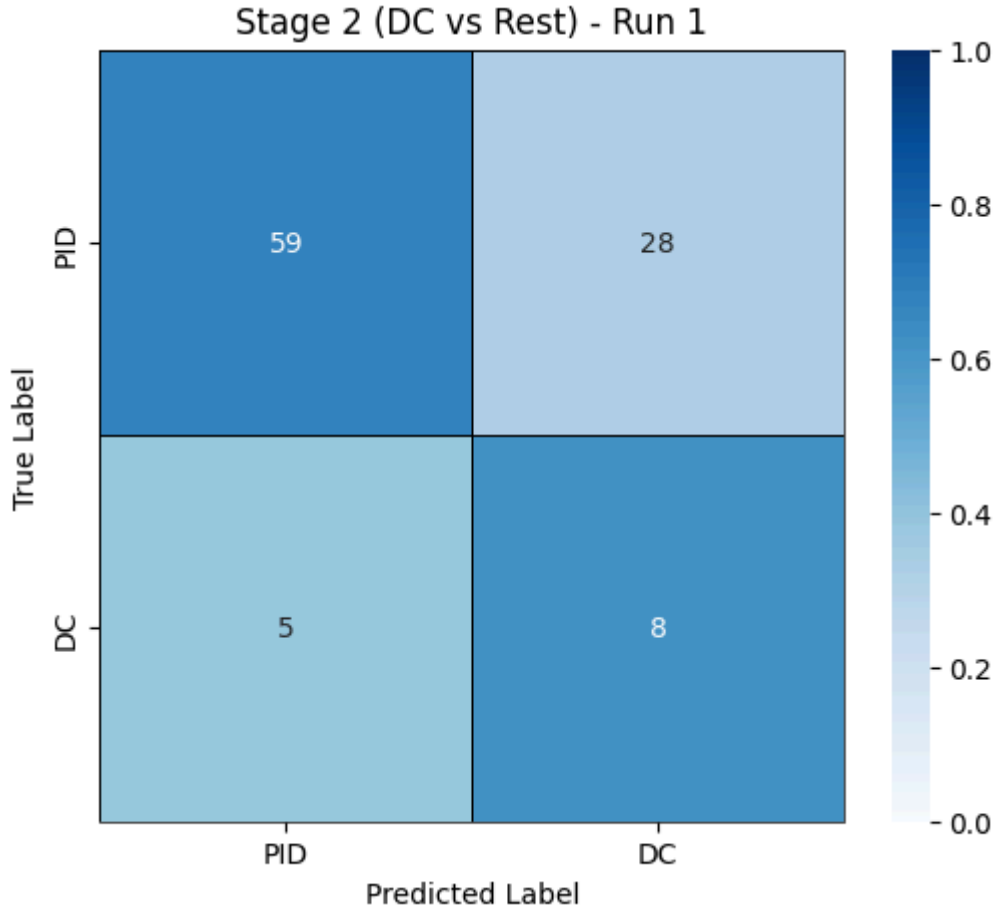

-- Stage 3 (Multiclass) --

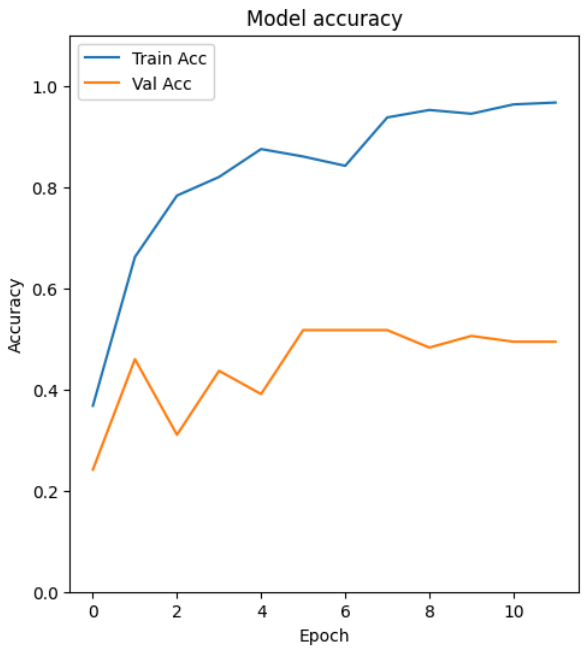

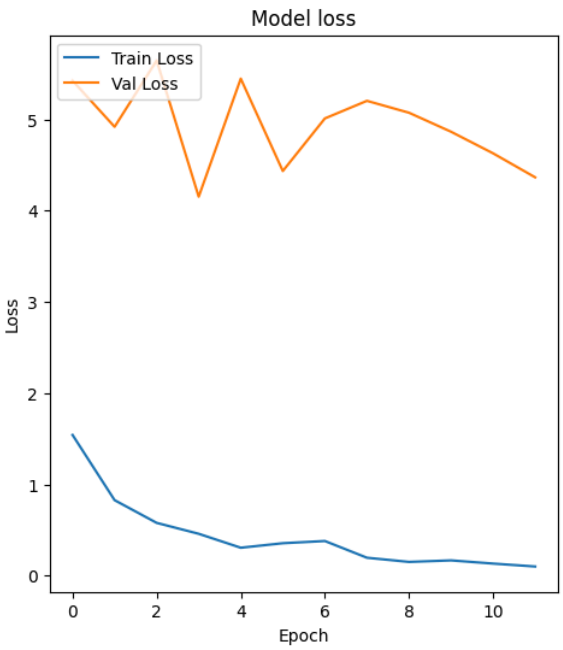

	precision	recall	f1-score	support
0	0.6667	0.5000	0.5714	4
1	0.2600	0.5417	0.3514	24
2	0.7000	0.3684	0.4828	57
3	0.0000	0.0000	0.0000	2
accuracy			0.4138	87
macro avg	0.4067	0.3525	0.3514	87
weighted avg	0.5610	0.4138	0.4395	87

Balanced Accuracy: 0.3525219298245614

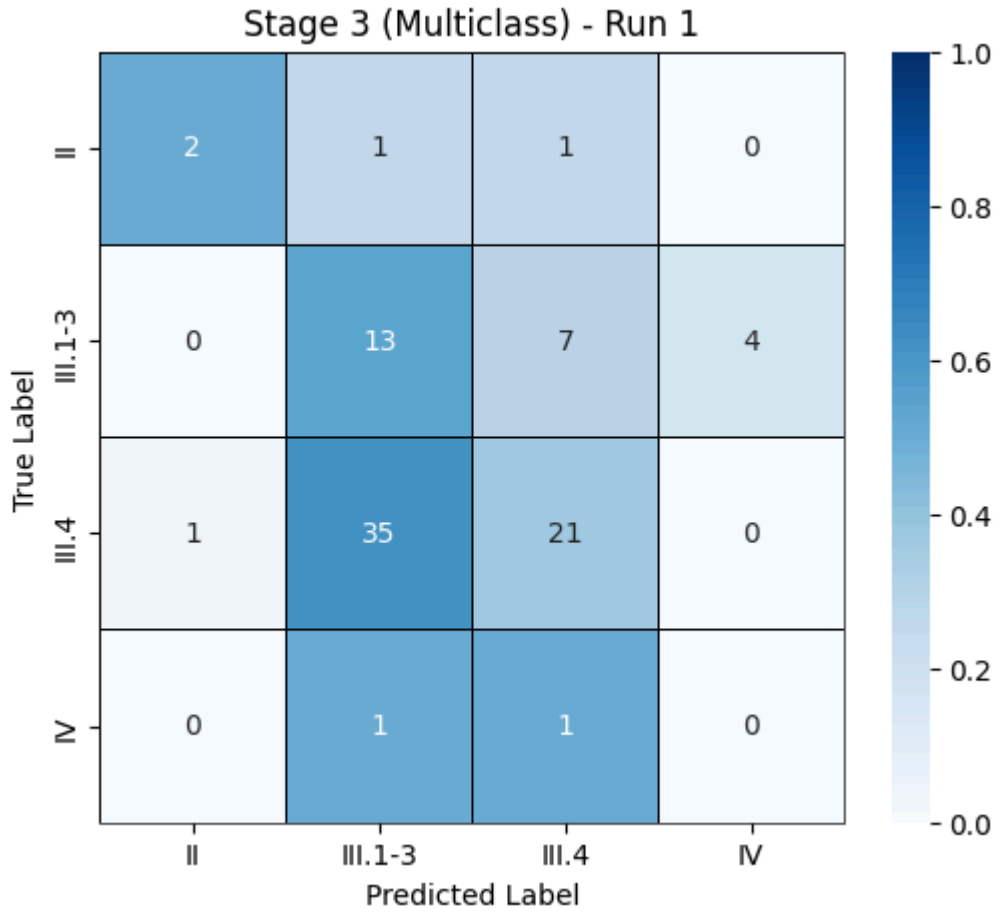

== Soft-Gated Overall (Test Set) - Run 1 ==

	precision	recall	f1-score	support
DC	0.2800	0.5385	0.3684	13
I	0.0612	1.0000	0.1154	3
II	0.0000	0.0000	0.0000	4
III.1-3	0.1304	0.1250	0.1277	24
III.4	0.6667	0.0702	0.1270	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.1650	103
macro avg	0.1897	0.2889	0.1231	103
weighted avg	0.4364	0.1650	0.1499	103

Balanced Accuracy: 0.28893949617633824

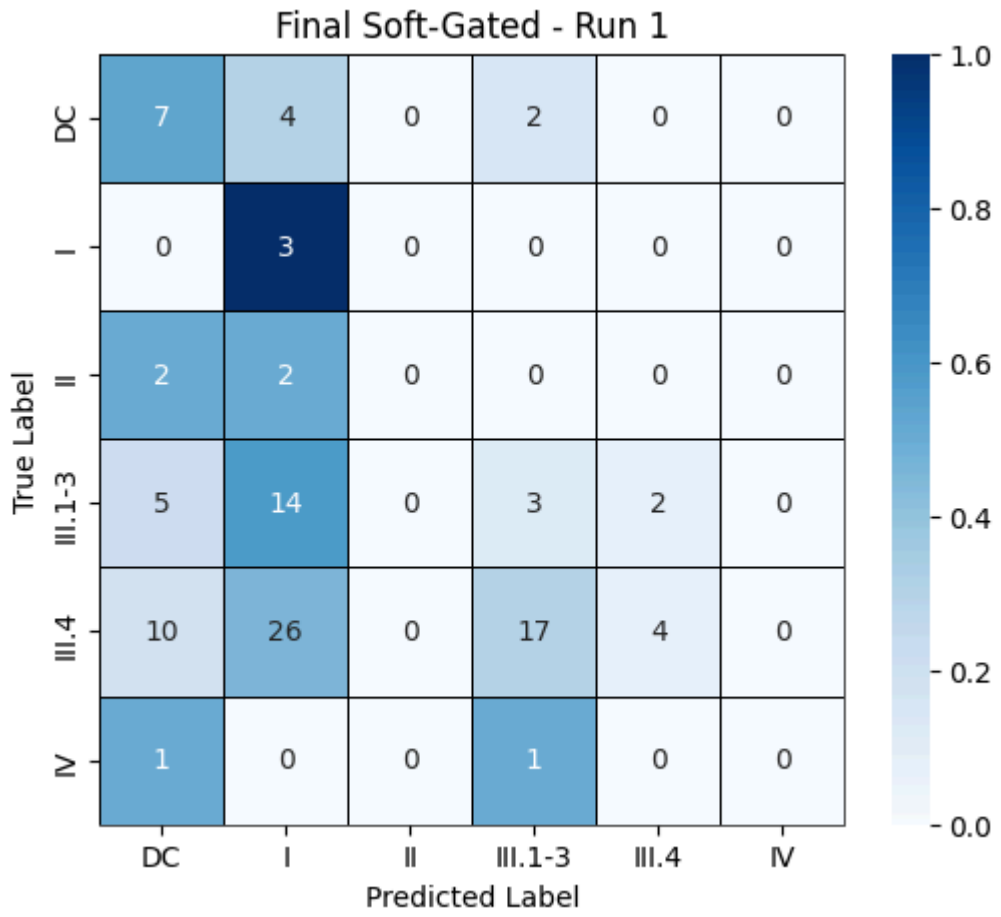

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

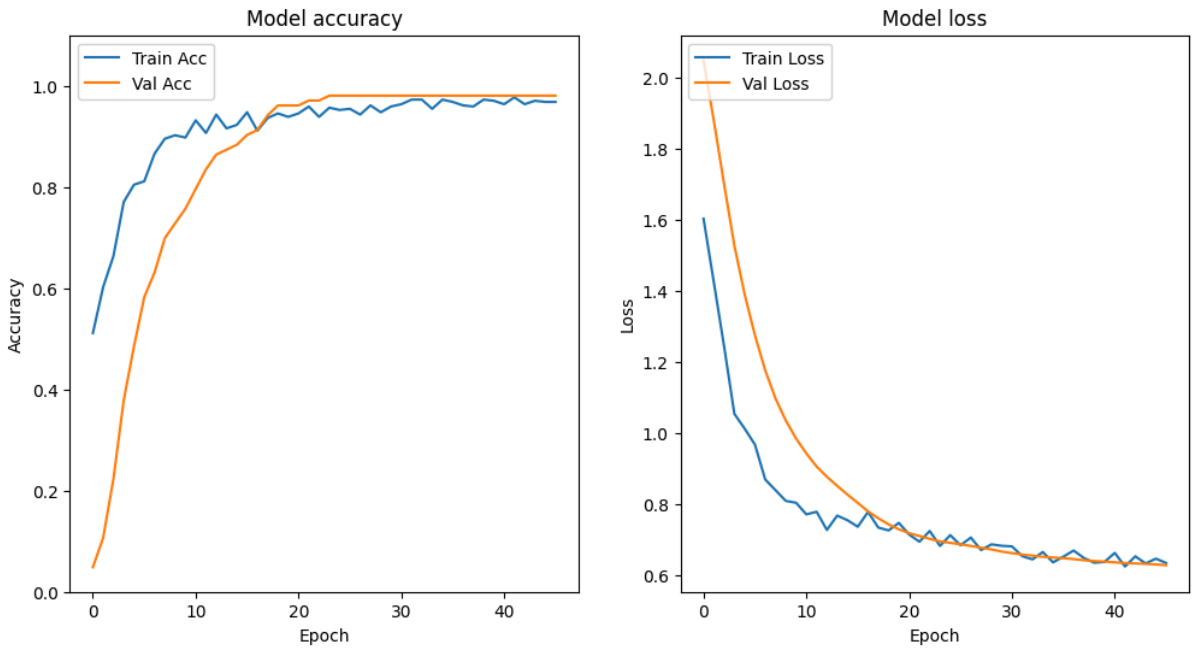

[Threshold Optimization] Best balanced_accuracy: 0.9900 at threshold=0.362
Optimal threshold (Stage 1 (I vs Rest)): 0.362

	precision	recall	f1-score	support
0	0.9894	0.9394	0.9637	99
1	0.3333	0.7500	0.4615	4
accuracy			0.9320	103
macro avg	0.6613	0.8447	0.7126	103
weighted avg	0.9639	0.9320	0.9442	103

Balanced Accuracy: 0.8446969696969697

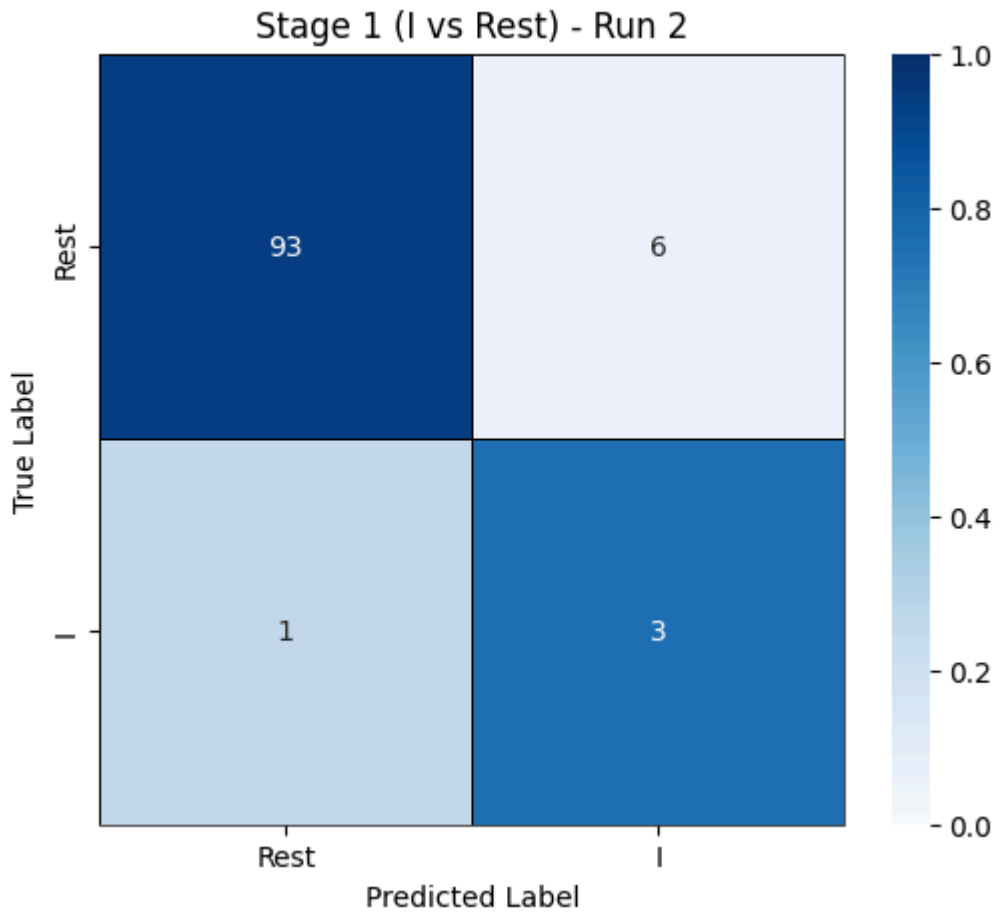

-- Stage 2 (DC vs Rest) --

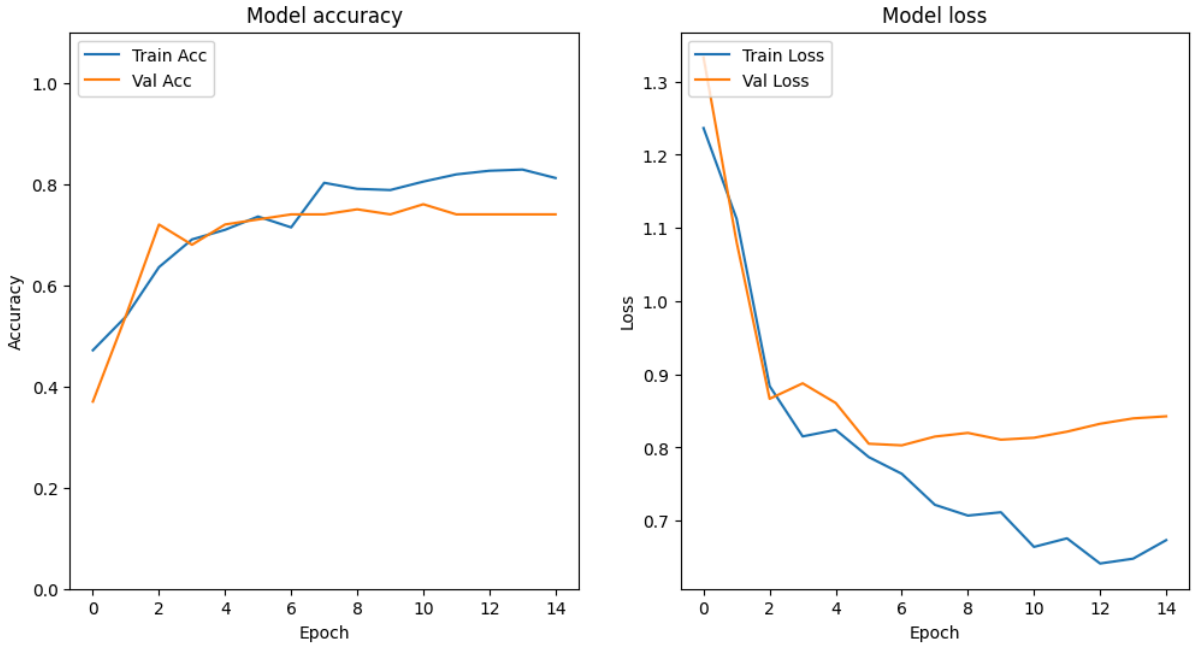

[Threshold Optimization] Best balanced_accuracy: 0.6220 at threshold=0.101
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.8448	0.5632	0.6759	87
1	0.0732	0.2500	0.1132	12
accuracy			0.5253	99
macro avg	0.4590	0.4066	0.3945	99
weighted avg	0.7513	0.5253	0.6077	99

Balanced Accuracy: 0.40660919540229884

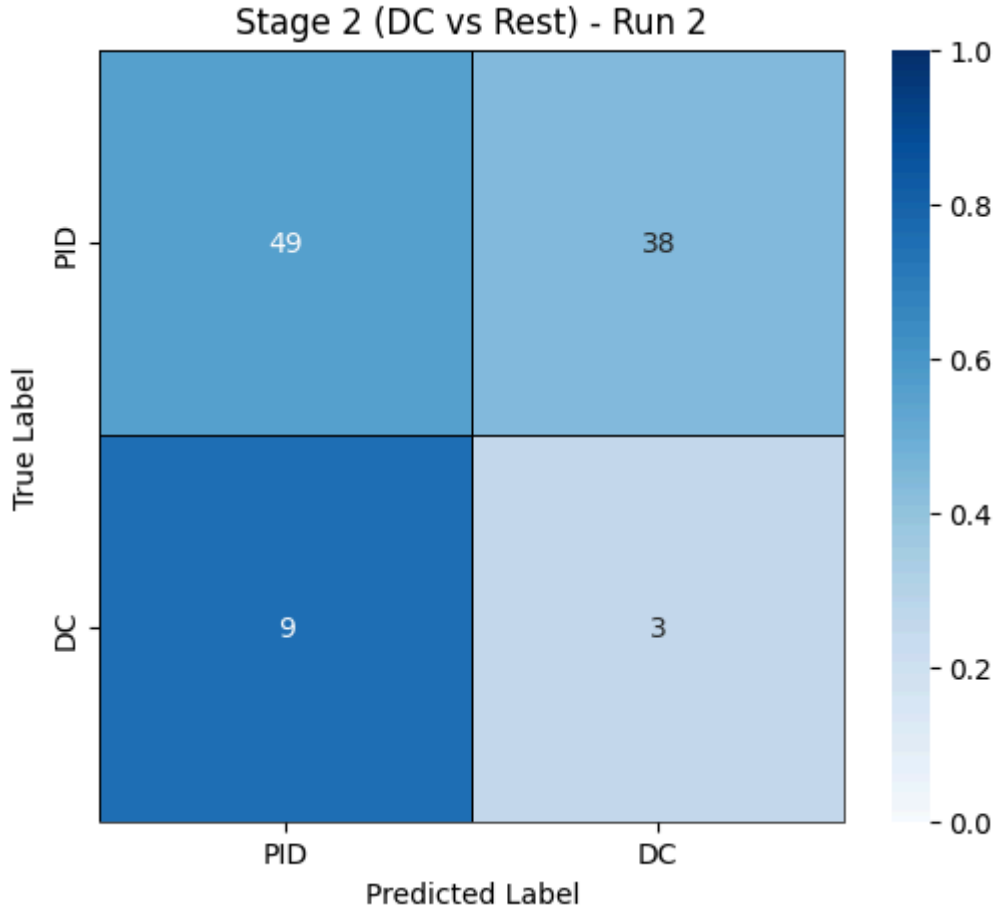

-- Stage 3 (Multiclass) --

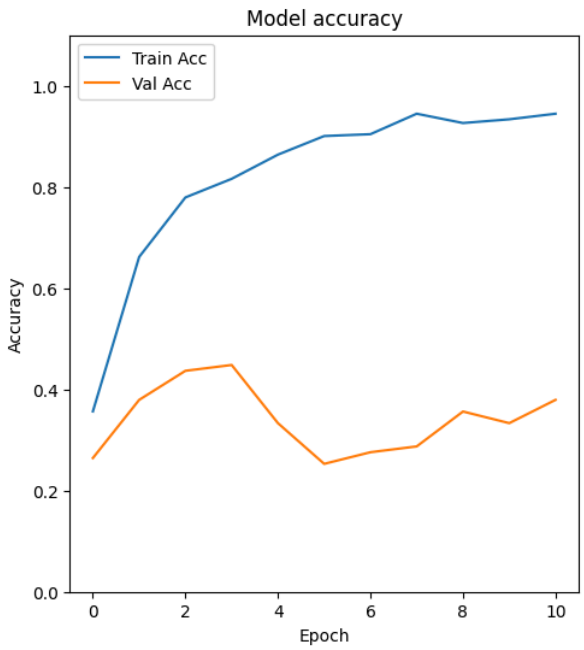

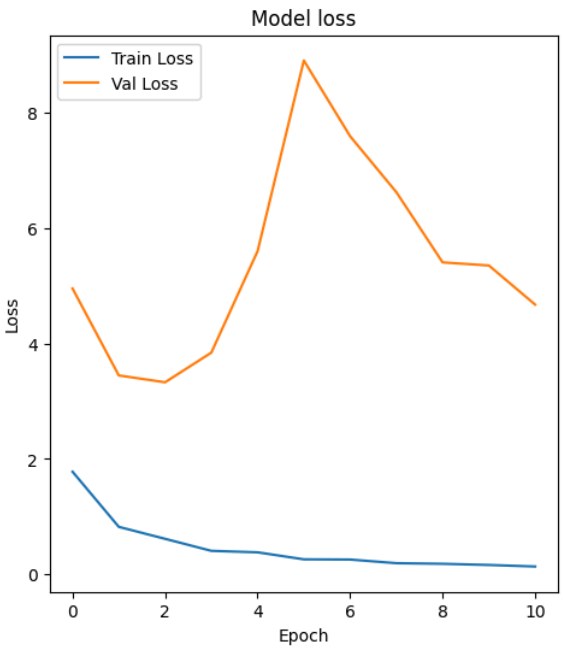

	precision	recall	f1-score	support
0	0.3333	0.2500	0.2857	4
1	0.3571	0.4000	0.3774	25
2	0.6545	0.6429	0.6486	56
3	0.0000	0.0000	0.0000	2
accuracy			0.5402	87
macro avg	0.3363	0.3232	0.3279	87
weighted avg	0.5393	0.5402	0.5391	87

Balanced Accuracy: 0.32321428571428573

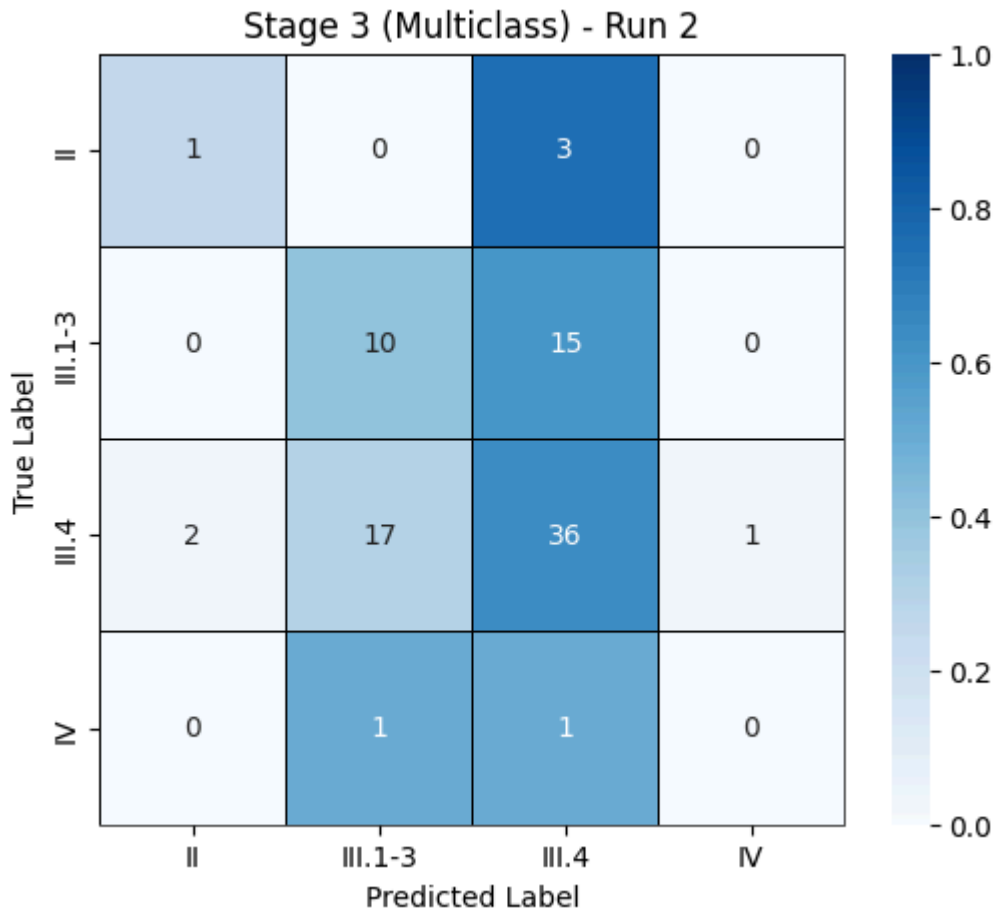

== Soft-Gated Overall (Test Set) - Run 2 ==

	precision	recall	f1-score	support
DC	0.0667	0.2500	0.1053	12
I	0.3000	0.7500	0.4286	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2667	0.1600	0.2000	25
III.4	0.6333	0.3393	0.4419	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.2816	103
macro avg	0.2111	0.2499	0.1959	103
weighted avg	0.4285	0.2816	0.3177	103

Balanced Accuracy: 0.24988095238095234

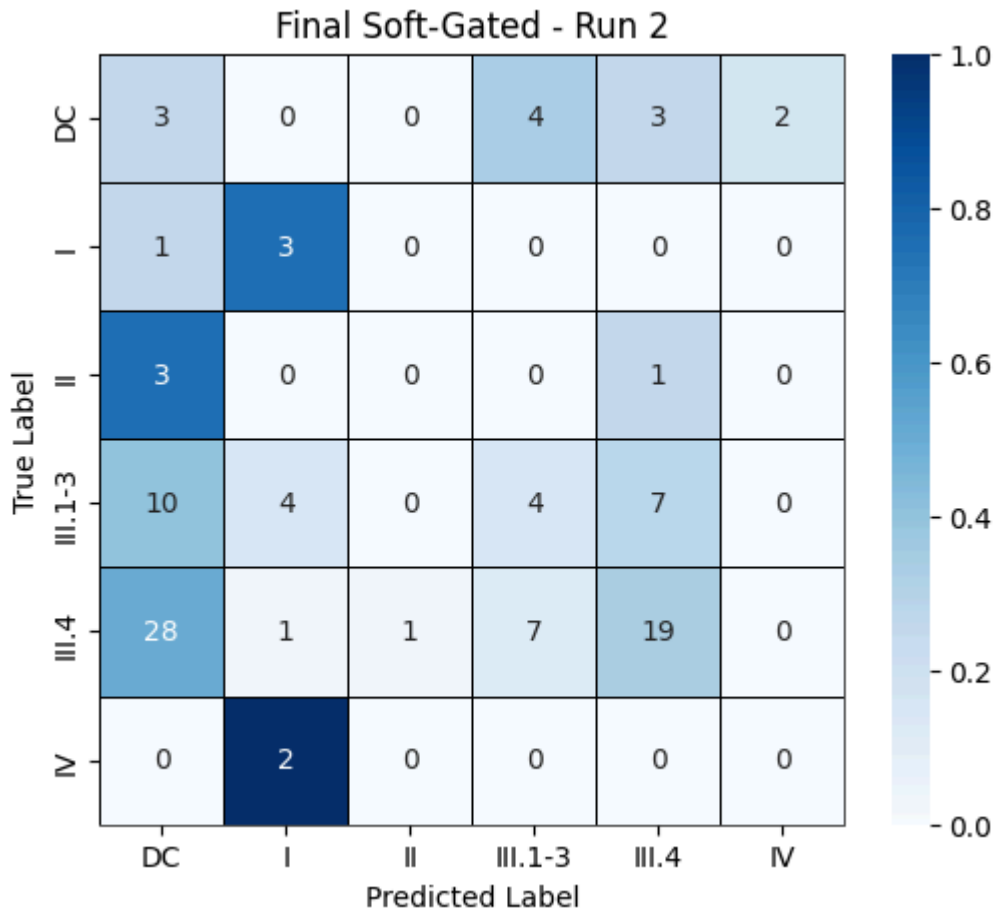

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.1429	0.4000	0.2105	25
I	0.1017	0.8571	0.1818	7
II	0.0000	0.0000	0.0000	8
III.1-3	0.1842	0.1429	0.1609	49
III.4	0.6389	0.2035	0.3087	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.2233	206
macro avg	0.1779	0.2673	0.1437	206
weighted avg	0.4151	0.2233	0.2394	206

Aggregated Balanced Accuracy: 0.2673

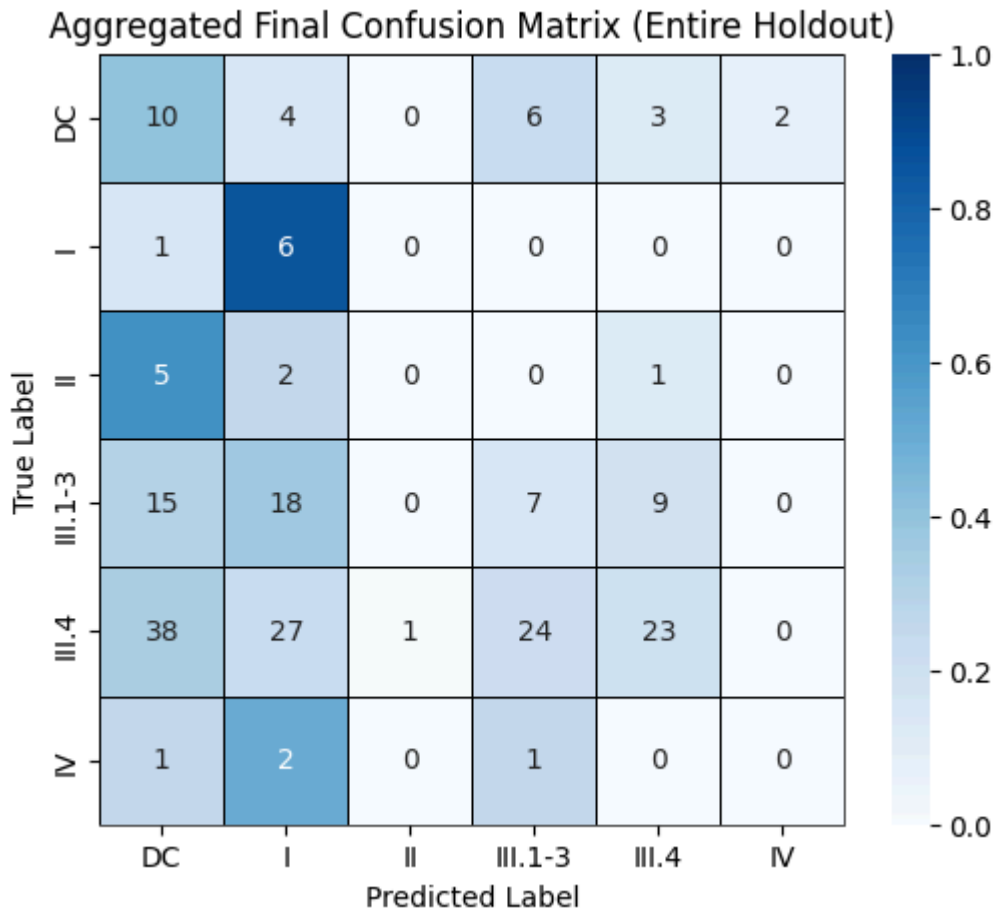

== Average Stage Balanced Accuracies ==
Stage 1 (I vs Rest): 0.8898
Stage 2 (DC vs Rest): 0.5267
Stage 3 (Multiclass): 0.3379
Final (Soft-Gated): 0.2694

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

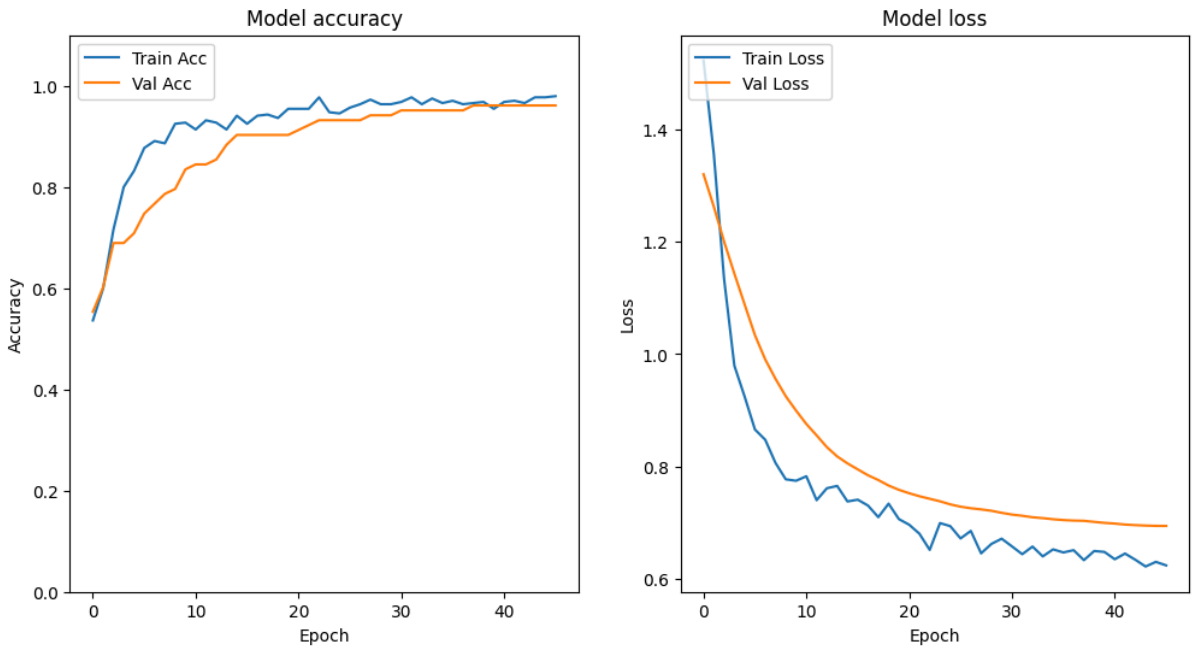

[Threshold Optimization] Best balanced_accuracy: 0.8737 at threshold=0.080
Optimal threshold (Stage 1 (I vs Rest)): 0.080

	precision	recall	f1-score	support
0	1.0000	0.8400	0.9130	100
1	0.1579	1.0000	0.2727	3
accuracy			0.8447	103
macro avg	0.5789	0.9200	0.5929	103
weighted avg	0.9755	0.8447	0.8944	103

Balanced Accuracy: 0.9199999999999999

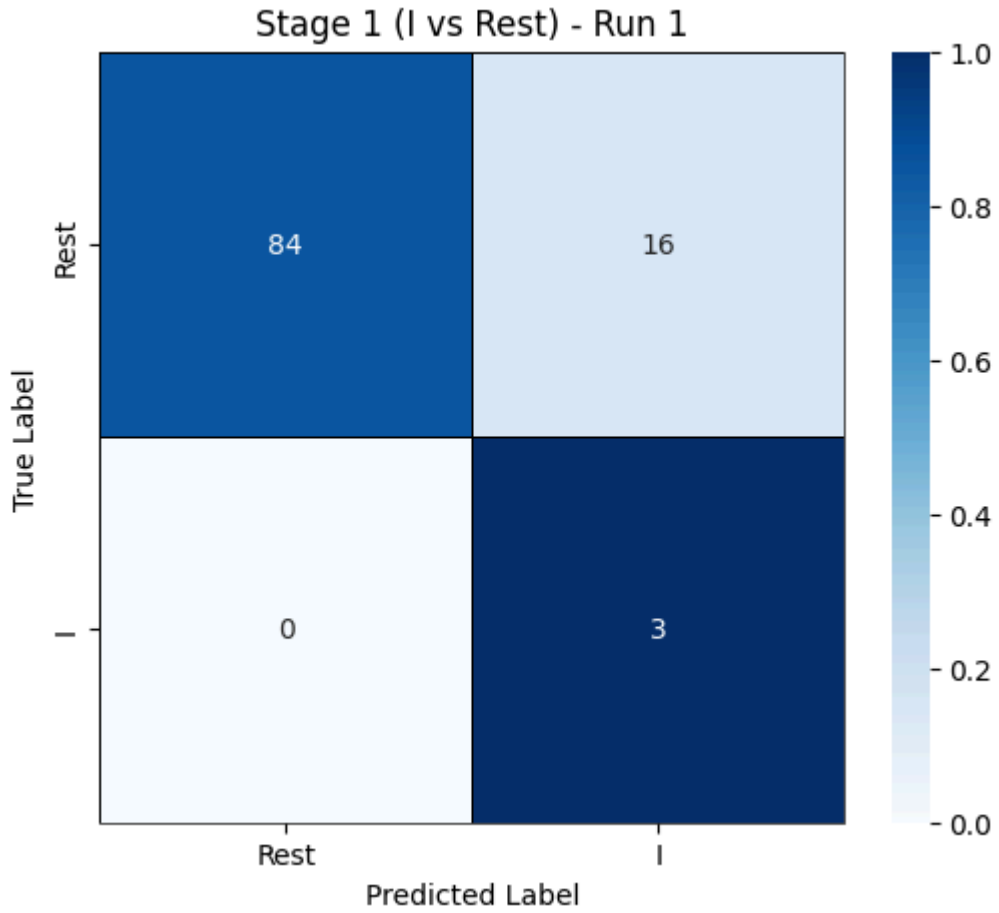

-- Stage 2 (DC vs Rest) --

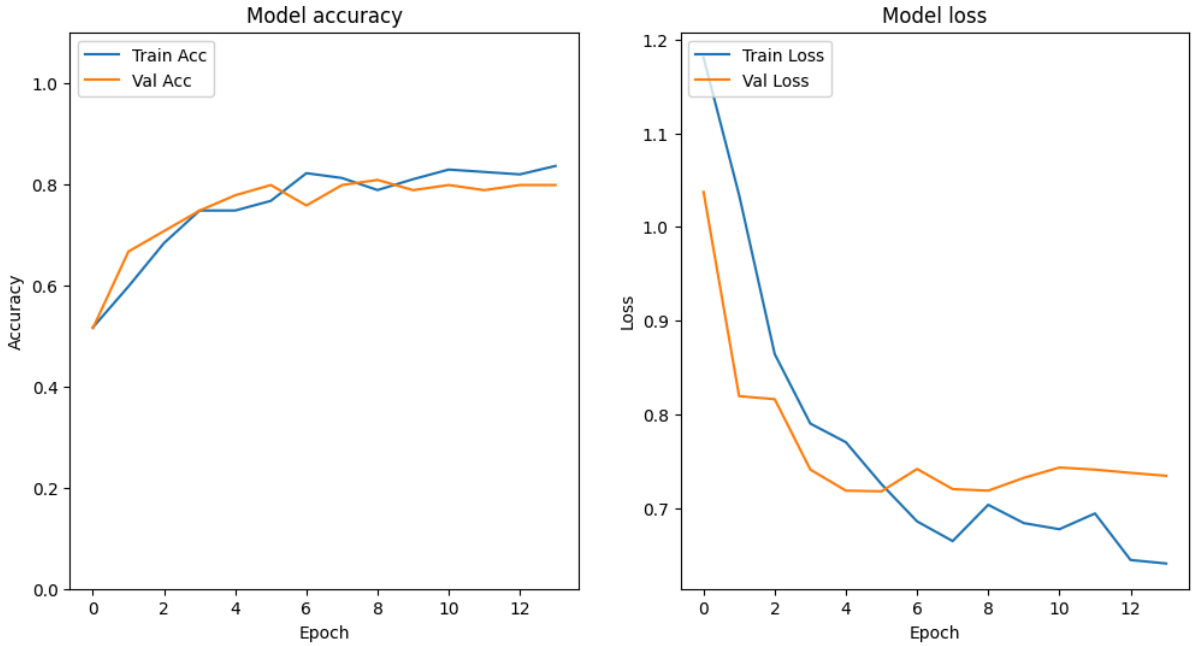

[Threshold Optimization] Best balanced_accuracy: 0.6351 at threshold=0.101
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.9615	0.5747	0.7194	87
1	0.2292	0.8462	0.3607	13
accuracy			0.6100	100
macro avg	0.5954	0.7104	0.5400	100
weighted avg	0.8663	0.6100	0.6728	100

Balanced Accuracy: 0.7104332449160036

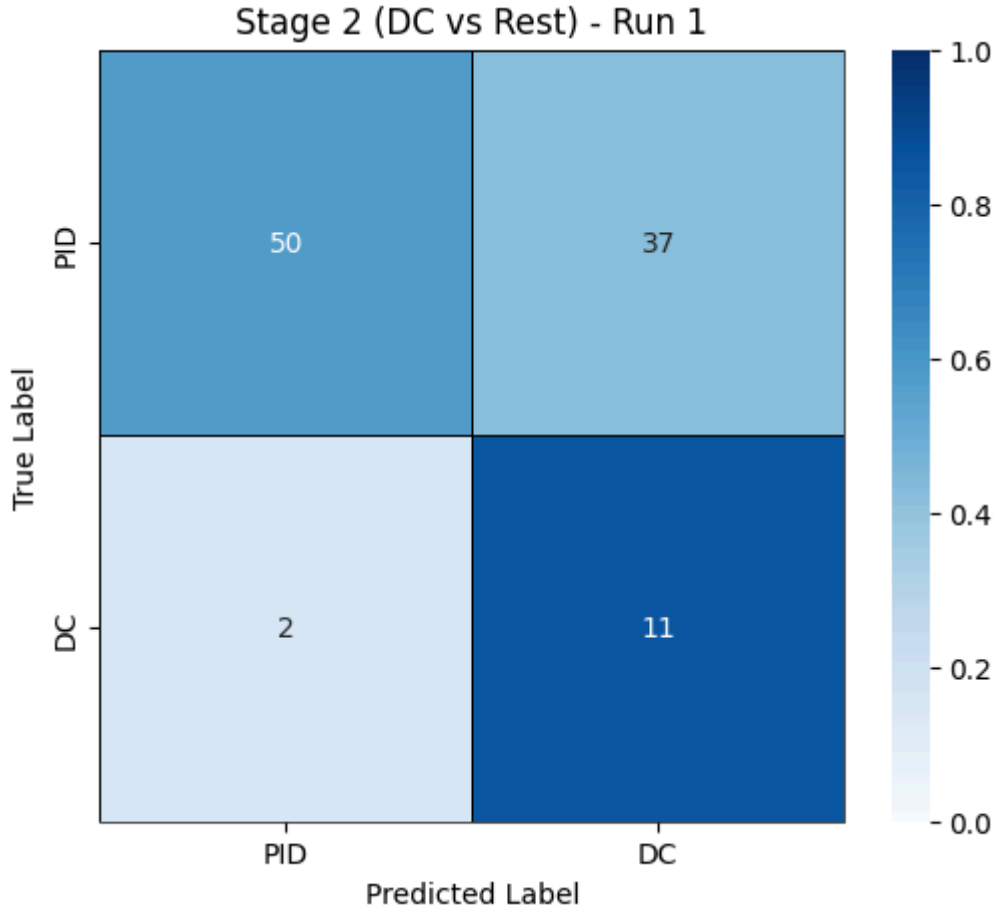

-- Stage 3 (Multiclass) --

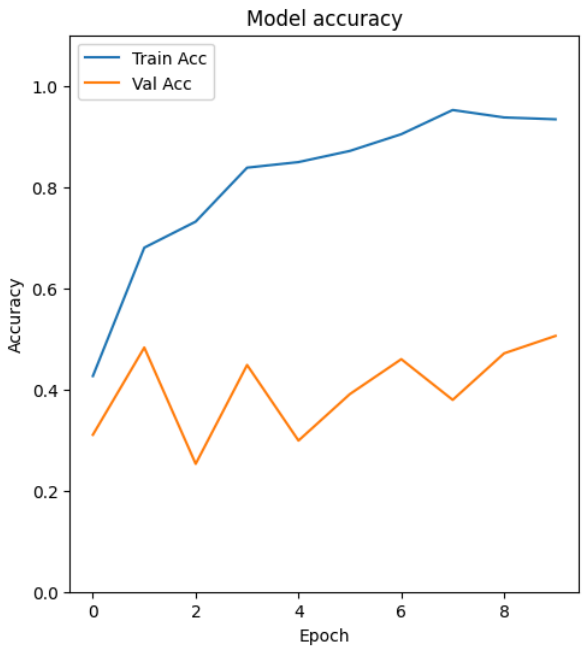

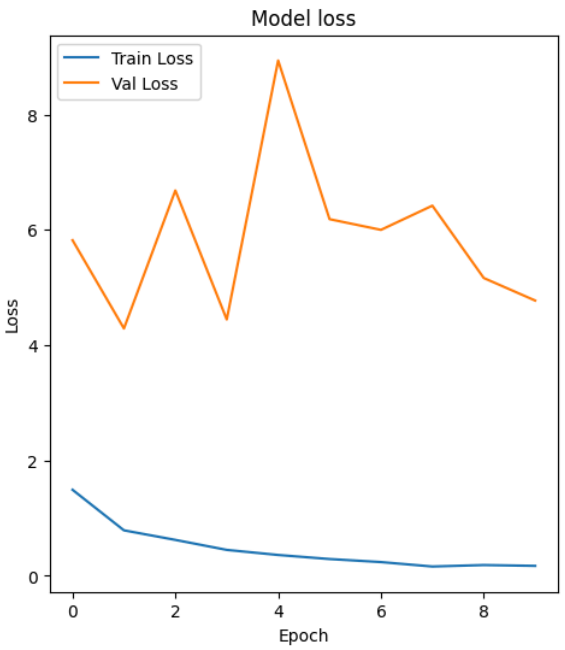

	precision	recall	f1-score	support
0	0.1429	0.5000	0.2222	4
1	0.2500	0.1250	0.1667	24
2	0.7000	0.6140	0.6542	57
3	0.0909	0.5000	0.1538	2
accuracy			0.4713	87
macro avg	0.2959	0.4348	0.2992	87
weighted avg	0.5362	0.4713	0.4883	87

Balanced Accuracy: 0.4347587719298246

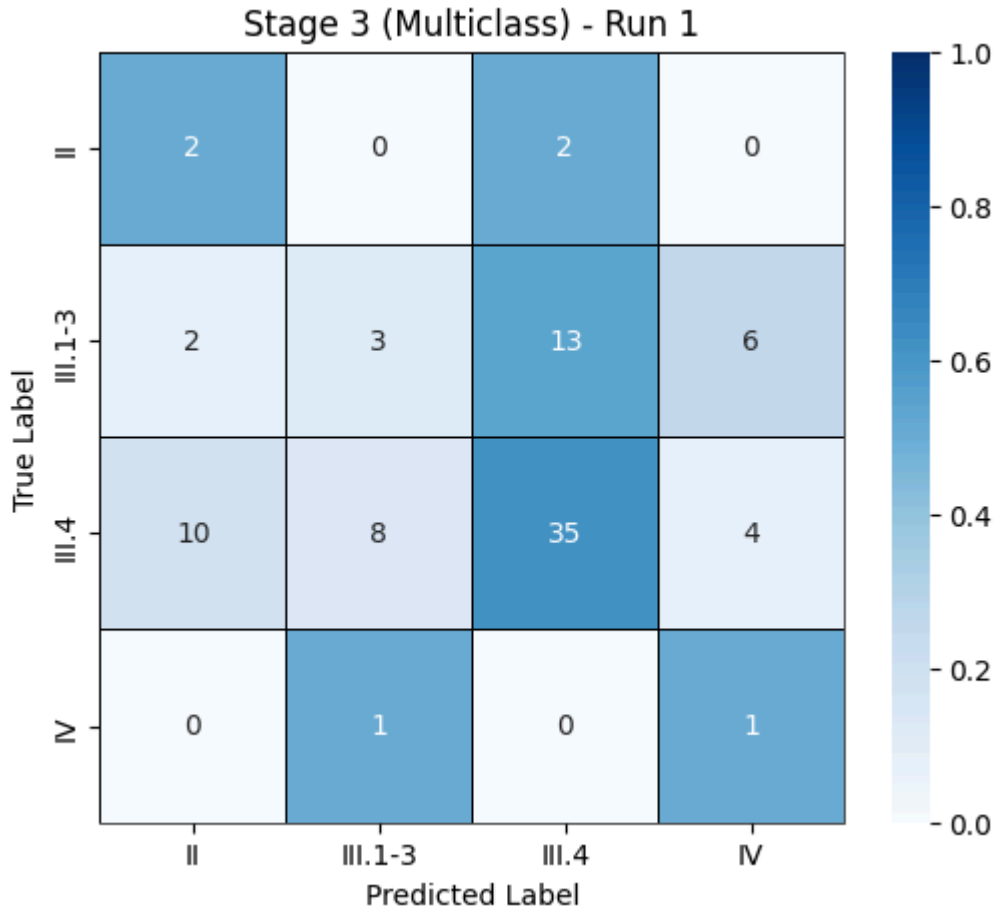

== Soft-Gated Overall (Test Set) - Run 1 ==

	precision	recall	f1-score	support
DC	0.2444	0.8462	0.3793	13
I	0.1000	1.0000	0.1818	3
II	0.1429	0.2500	0.1818	4
III.1-3	0.0000	0.0000	0.0000	24
III.4	0.8235	0.2456	0.3784	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.2816	103
macro avg	0.2185	0.3903	0.1869	103
weighted avg	0.4951	0.2816	0.2696	103

Balanced Accuracy: 0.3902946468735942

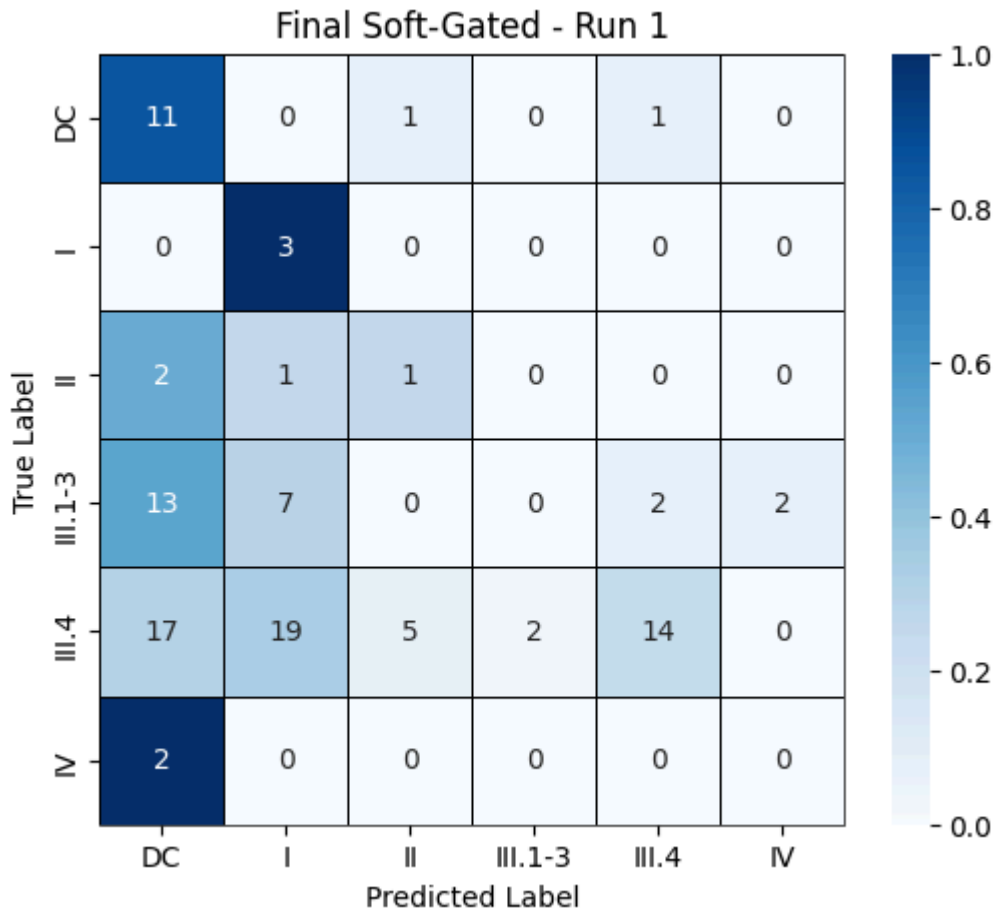

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

-- Stage 1 (I vs Rest) --

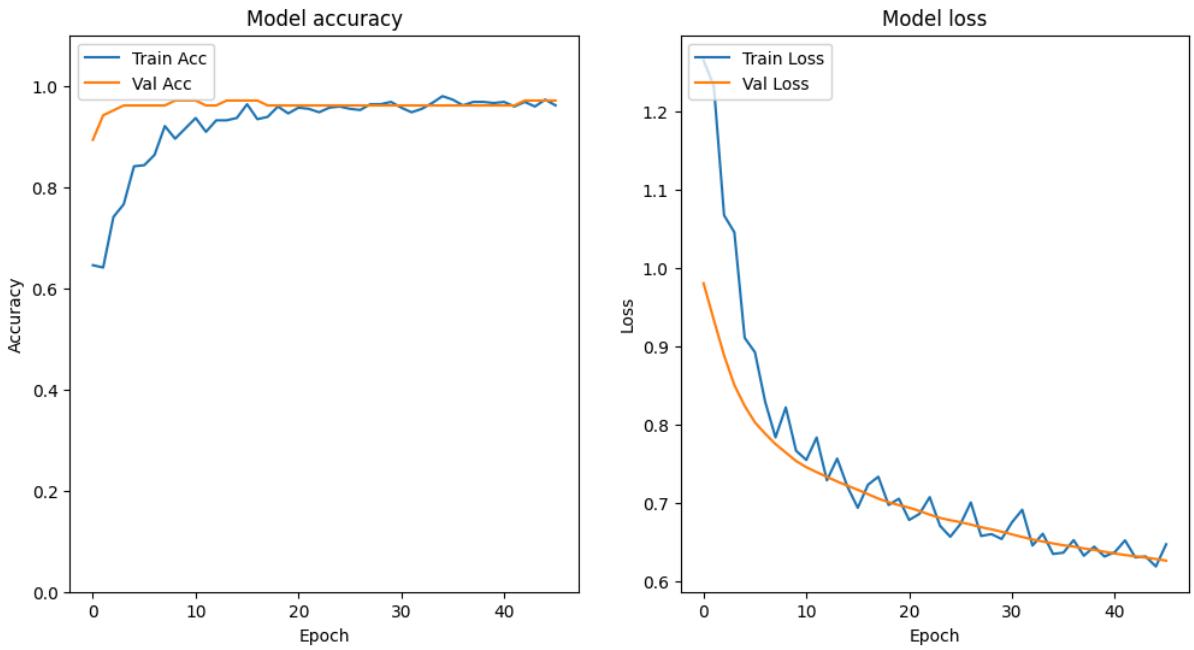

[Threshold Optimization] Best balanced_accuracy: 0.9550 at threshold=0.151
Optimal threshold (Stage 1 (I vs Rest)): 0.151

	precision	recall	f1-score	support
0	1.0000	0.7879	0.8814	99
1	0.1600	1.0000	0.2759	4
accuracy			0.7961	103
macro avg	0.5800	0.8939	0.5786	103
weighted avg	0.9674	0.7961	0.8578	103

Balanced Accuracy: 0.8939393939393939

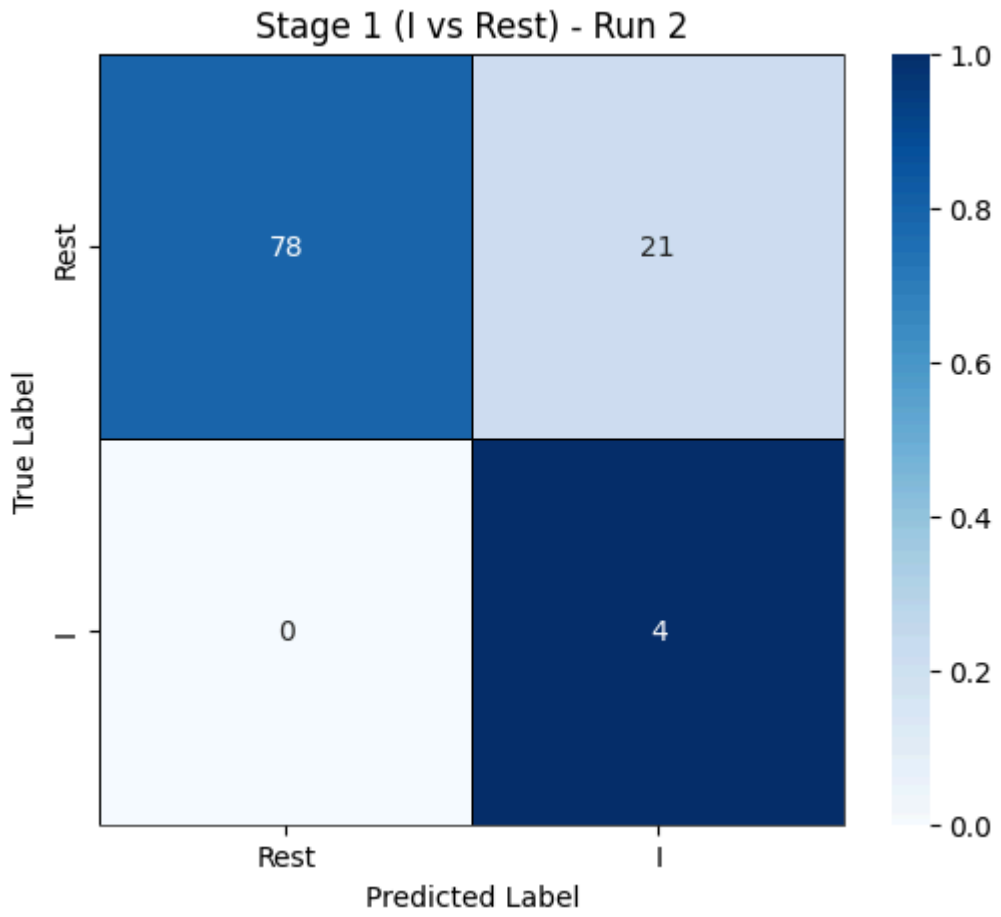

-- Stage 2 (DC vs Rest) --

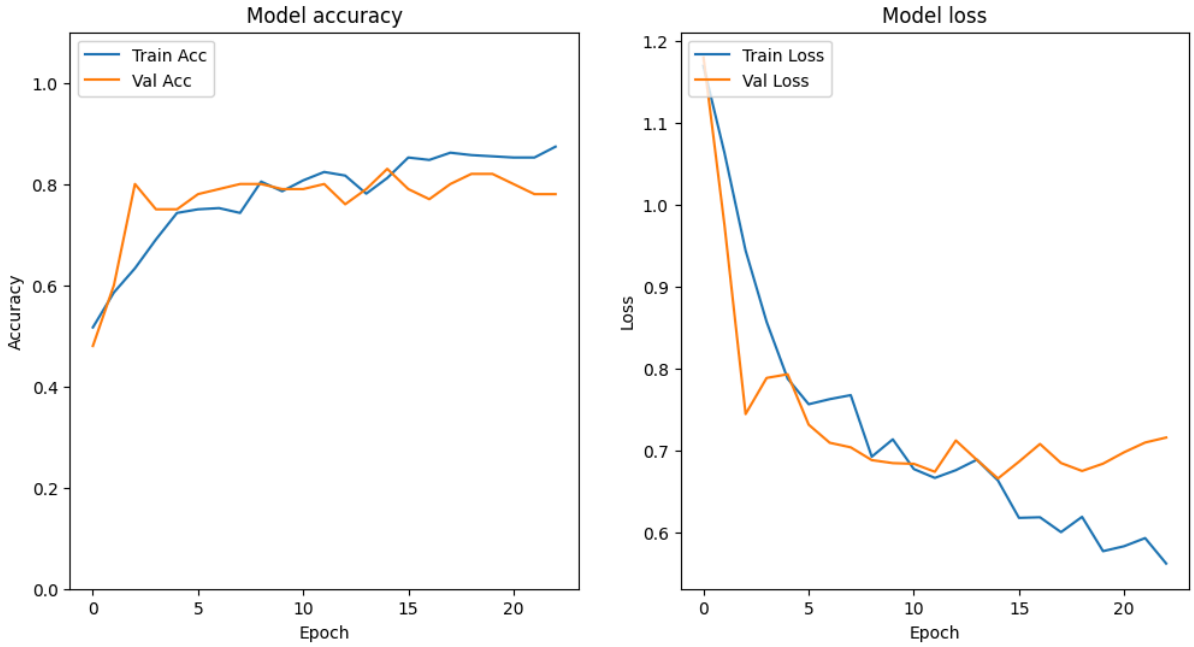

[Threshold Optimization] Best balanced_accuracy: 0.7776 at threshold=0.146
Optimal threshold (Stage 2 (DC vs Rest)): 0.200

	precision	recall	f1-score	support

0	0.9167	0.7586	0.8302	87
1	0.2222	0.5000	0.3077	12
accuracy			0.7273	99
macro avg	0.5694	0.6293	0.5689	99
weighted avg	0.8325	0.7273	0.7669	99

Balanced Accuracy: 0.6293103448275862

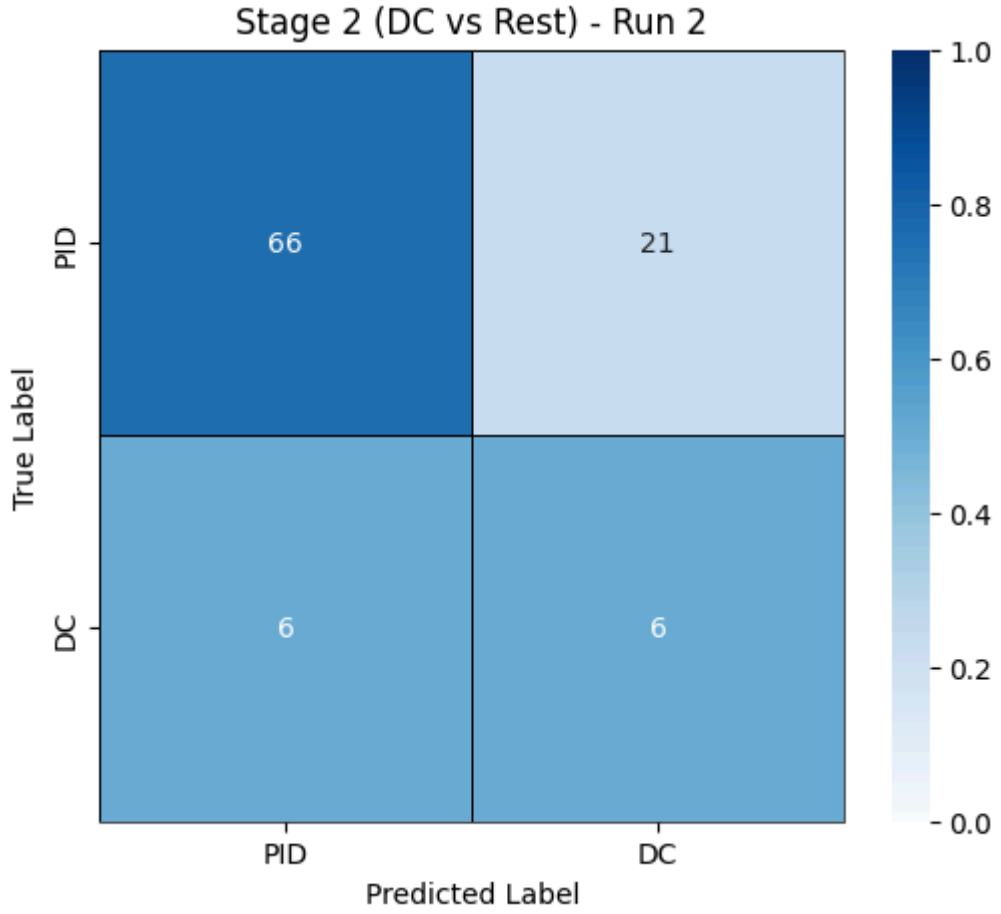

-- Stage 3 (Multiclass) --

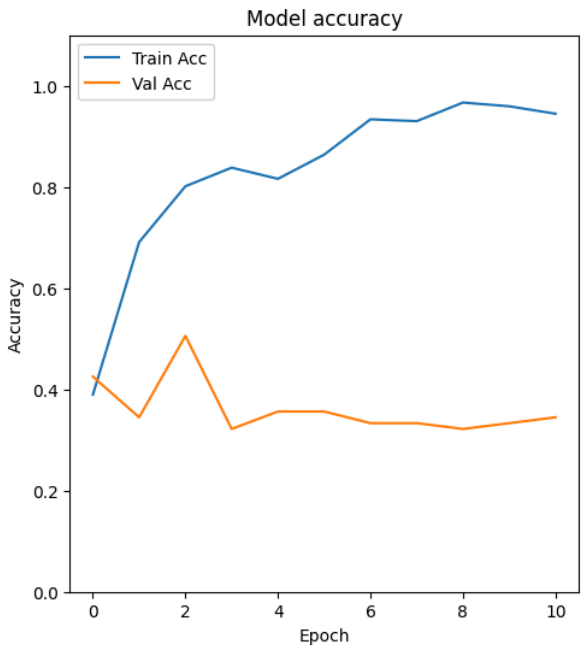

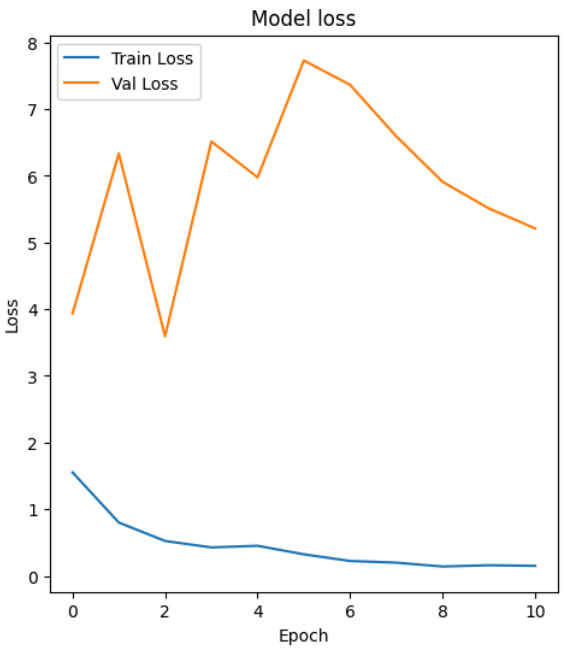

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	4
1	0.4074	0.4400	0.4231	25
2	0.7917	0.6786	0.7308	56
3	0.0000	0.0000	0.0000	2
accuracy			0.5632	87
macro avg	0.2998	0.2796	0.2885	87
weighted avg	0.6266	0.5632	0.5920	87

Balanced Accuracy: 0.27964285714285714

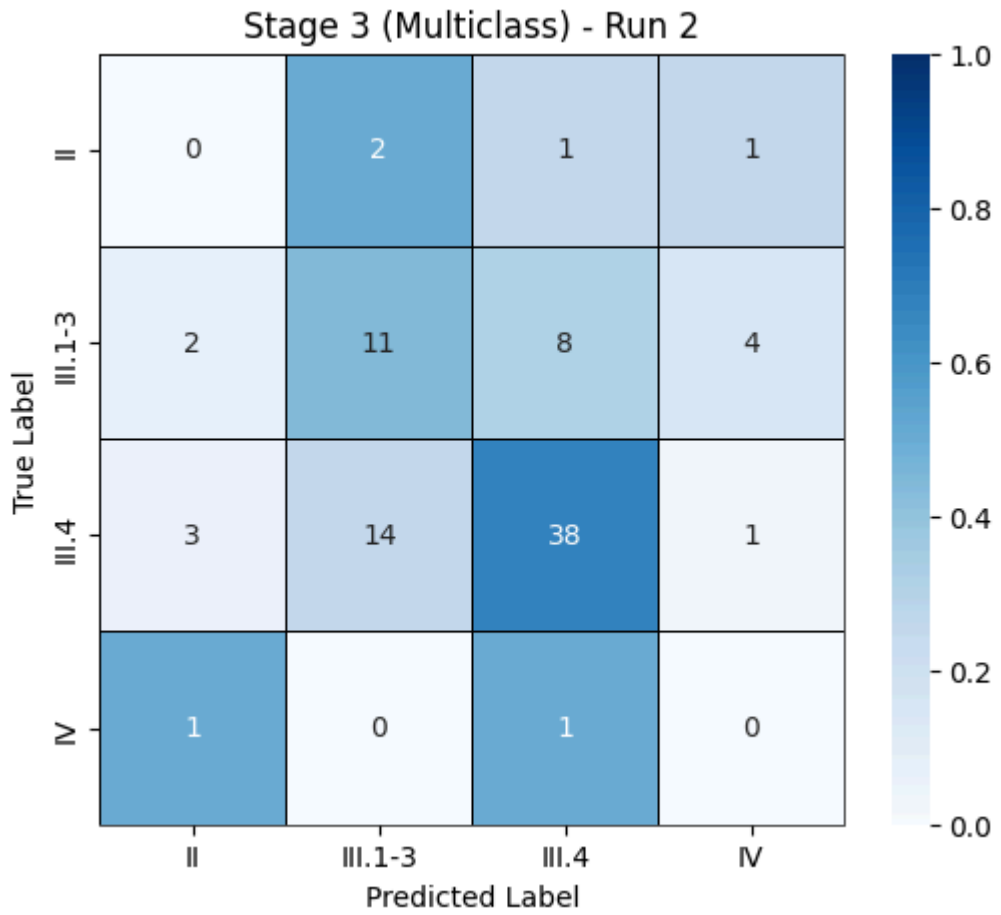

== Soft-Gated Overall (Test Set) - Run 2 ==

	precision	recall	f1-score	support
DC	0.2143	0.5000	0.3000	12
I	0.1538	1.0000	0.2667	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.3125	0.2000	0.2439	25
III.4	0.7778	0.3750	0.5060	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.3495	103
macro avg	0.2431	0.3458	0.2194	103
weighted avg	0.5297	0.3495	0.3796	103

Balanced Accuracy: 0.3458333333333334

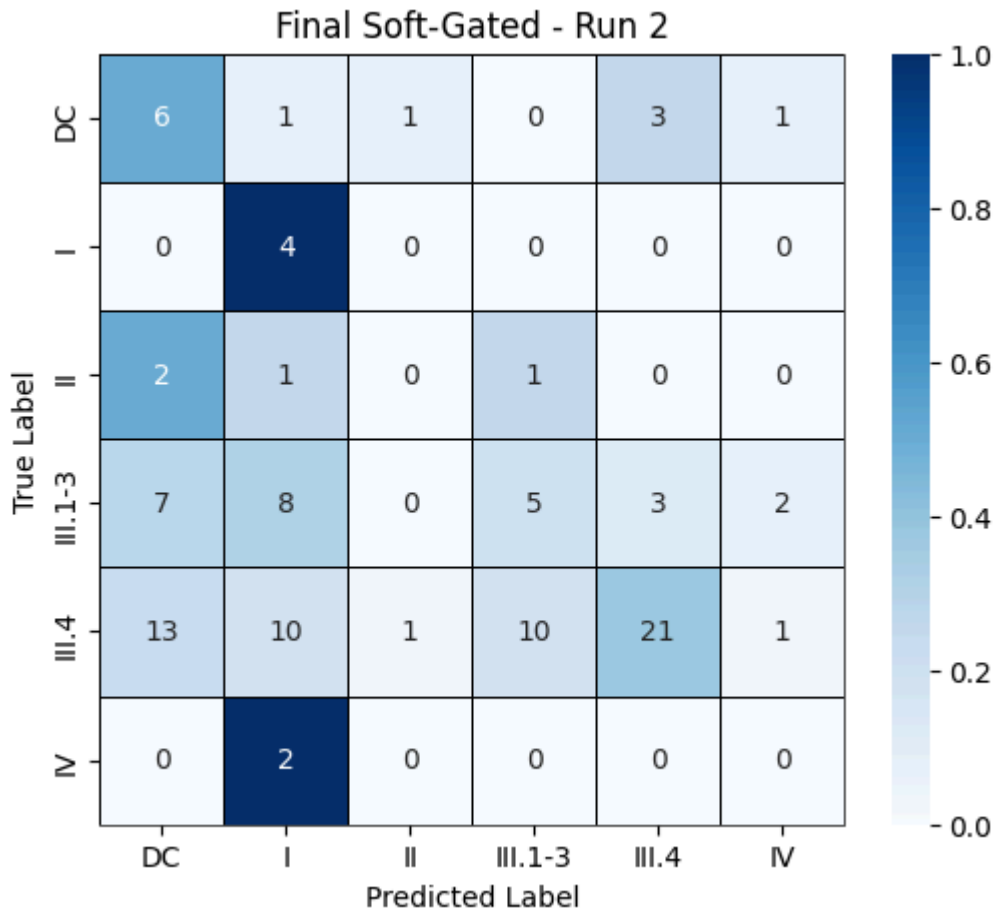

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.2329	0.6800	0.3469	25
I	0.1250	1.0000	0.2222	7
II	0.1111	0.1250	0.1176	8
III.1-3	0.2778	0.1020	0.1493	49
III.4	0.7955	0.3097	0.4459	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3155	206
macro avg	0.2570	0.3695	0.2137	206
weighted avg	0.5392	0.3155	0.3343	206

Aggregated Balanced Accuracy: 0.3695

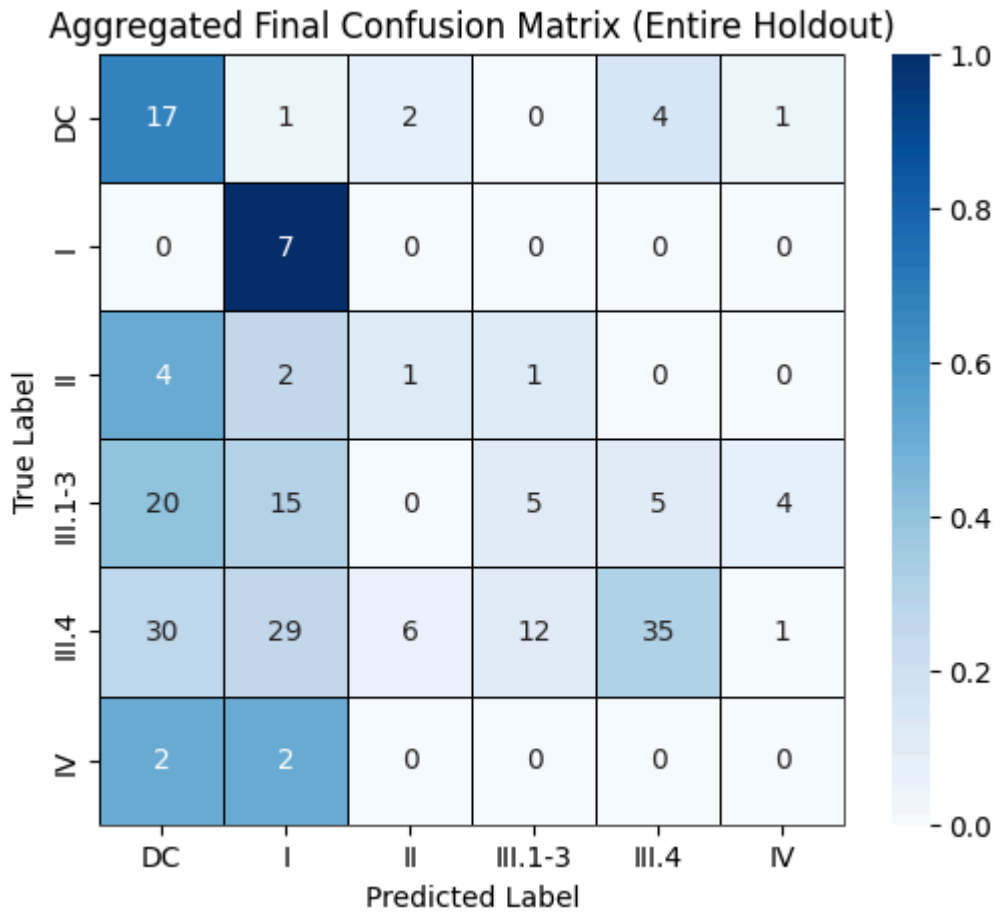

== Average Stage Balanced Accuracies ==
Stage 1 (I vs Rest): 0.9070
Stage 2 (DC vs Rest): 0.6699
Stage 3 (Multiclass): 0.3572
Final (Soft-Gated): 0.3681

[0.8673232323232323, 0.5627763041556144, 0.2800031328320802, 0.24578171091445425]
[0.861919191919192, 0.5250884173297966, 0.36178258145363407, 0.2895755824453675]
[0.8722474747474747, 0.5644341290893015, 0.30152882205513787, 0.3070182409246885]
[0.8898484848484849, 0.5266909814323607, 0.33786810776942355, 0.2672566371681416]
[0.906969696969697, 0.6698717948717949, 0.3572008145363409, 0.3694625549334778]

0.880 ± 0.007
0.570 ± 0.024
0.328 ± 0.014
0.296 ± 0.019

FLAT
==
RUN 1: Using Split A for validation, Split B for testing
==
These features will be dropped:
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']
4/4 ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ **0s** 21ms/step

== Test Set Evaluation - Run 1 ==

	precision	recall	f1-score	support
DC	0.5714	0.3077	0.4000	13
I	0.1250	1.0000	0.2222	3
II	0.2500	0.2500	0.2500	4
III.1-3	0.1930	0.4583	0.2716	24
III.4	1.0000	0.1228	0.2188	57
IV	0.2500	0.5000	0.3333	2
accuracy			0.2621	103
macro avg	0.3982	0.4398	0.2827	103
weighted avg	0.6887	0.2621	0.2575	103

Balanced Accuracy: 0.43980544309491676
Confusion Matrix:
[[4 0 0 9 0 0]
[0 3 0 0 0 0]
[0 0 1 2 0 1]
[0 9 2 11 0 2]
[3 12 1 34 7 0]
[0 0 0 1 0 1]]

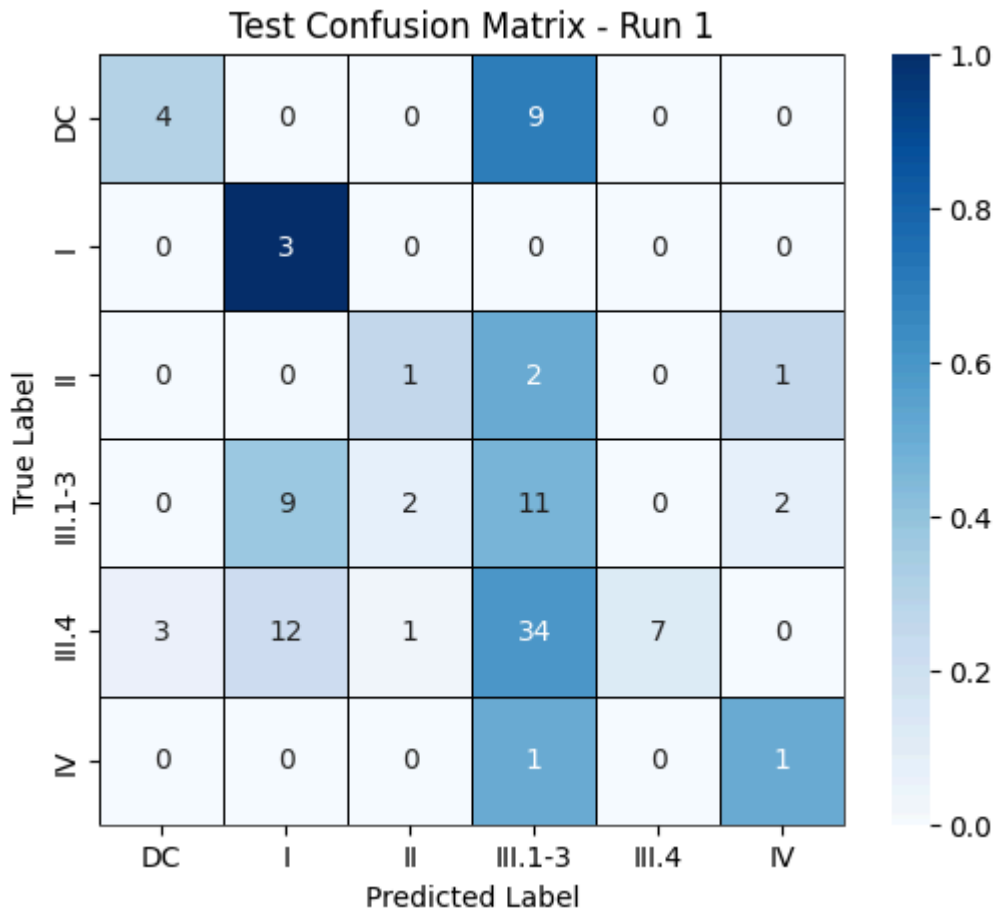

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ━━━━━━━━━━━━━━━━━━━━ **0s** 20ms/step

== Test Set Evaluation - Run 2 ==

	precision	recall	f1-score	support
DC	0.5000	0.2500	0.3333	12
I	0.5000	1.0000	0.6667	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2857	0.4800	0.3582	25
III.4	0.6364	0.5000	0.5600	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.4563	103
macro avg	0.3203	0.3717	0.3197	103
weighted avg	0.4930	0.4563	0.4561	103

Balanced Accuracy: 0.37166666666666665

Confusion Matrix:
[[3 0 0 3 4 2]
 [0 4 0 0 0 0]
 [0 0 0 1 3 0]
 [1 3 0 12 9 0]
 [2 0 0 25 28 1]
 [0 1 0 1 0 0]]

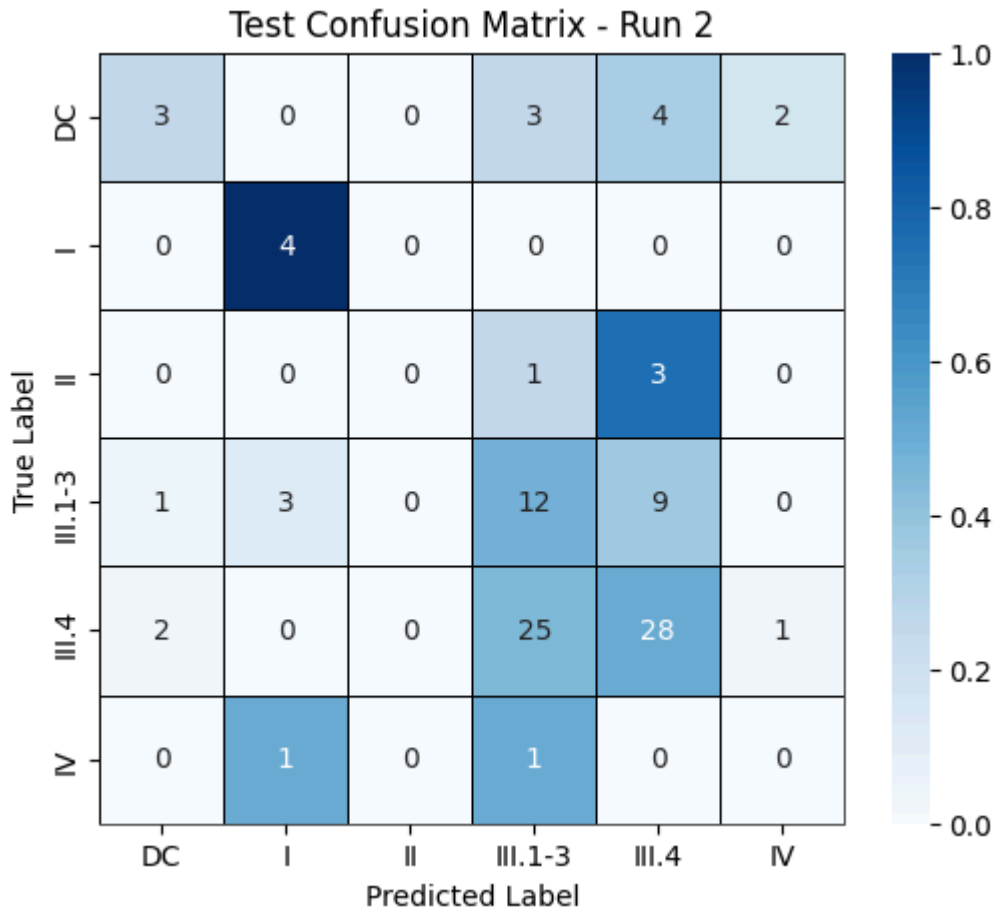

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.5385	0.2800	0.3684	25
I	0.2188	1.0000	0.3590	7
II	0.2500	0.1250	0.1667	8
III.1-3	0.2323	0.4694	0.3108	49
III.4	0.6863	0.3097	0.4268	113
IV	0.1429	0.2500	0.1818	4
accuracy			0.3592	206

macro avg	0.3448	0.4057	0.3023	206
weighted avg	0.5170	0.3592	0.3750	206

Aggregated Balanced Accuracy: 0.4057
Aggregated Confusion Matrix:
[[7 0 0 12 4 2]
 [0 7 0 0 0 0]
 [0 0 1 3 3 1]
 [1 12 2 23 9 2]
 [5 12 1 59 35 1]
 [0 1 0 2 0 1]]

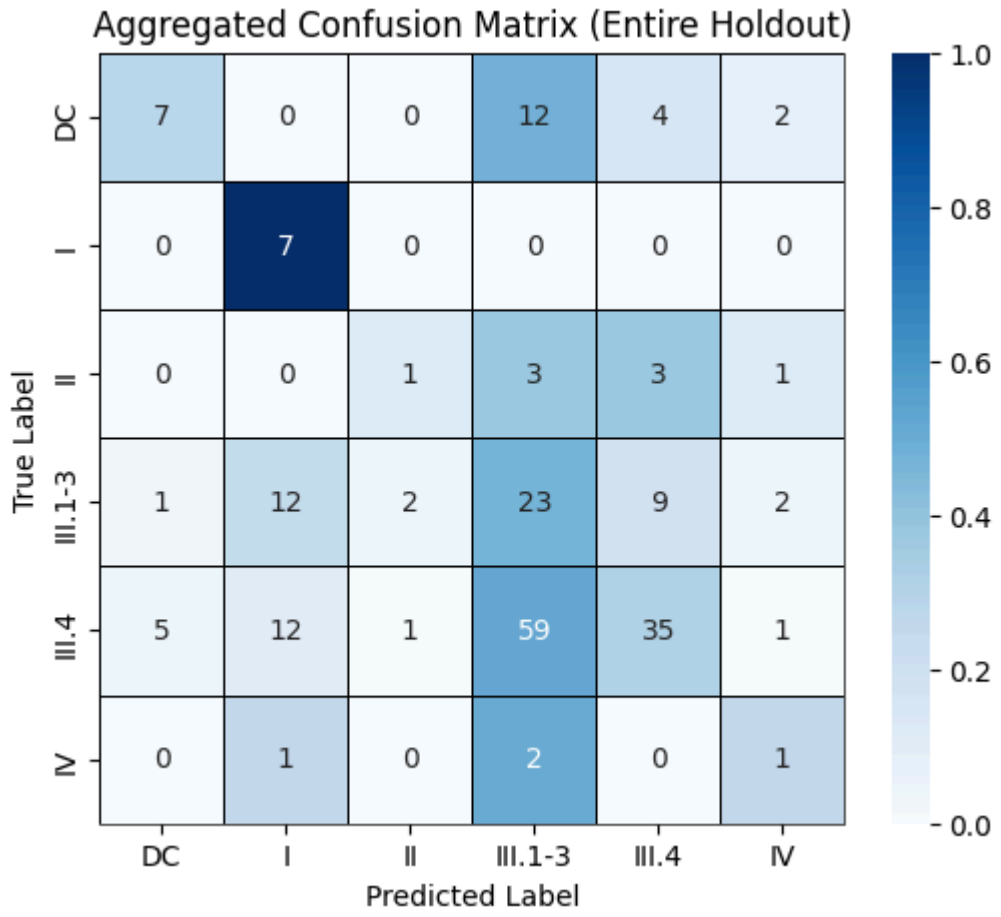

== Average Balanced Accuracy Across Runs ==
Run 1: 0.4398
Run 2: 0.3717
Average: 0.4057

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']
4/4 ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ **0s** 21ms/step

== Test Set Evaluation - Run 1 ==

	precision	recall	f1-score	support
DC	0.0000	0.0000	0.0000	13
I	0.3750	1.0000	0.5455	3
II	0.4000	0.5000	0.4444	4
III.1-3	0.2759	0.3333	0.3019	24
III.4	0.5667	0.5965	0.5812	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.4563	103
macro avg	0.2696	0.4050	0.3122	103
weighted avg	0.4043	0.4563	0.4251	103

Balanced Accuracy: 0.4049707602339181
Confusion Matrix:
[[0 0 0 1 12 0]
[0 3 0 0 0 0]
[0 0 2 1 1 0]
[0 3 1 8 11 1]
[0 2 2 19 34 0]
[0 0 0 0 2 0]]

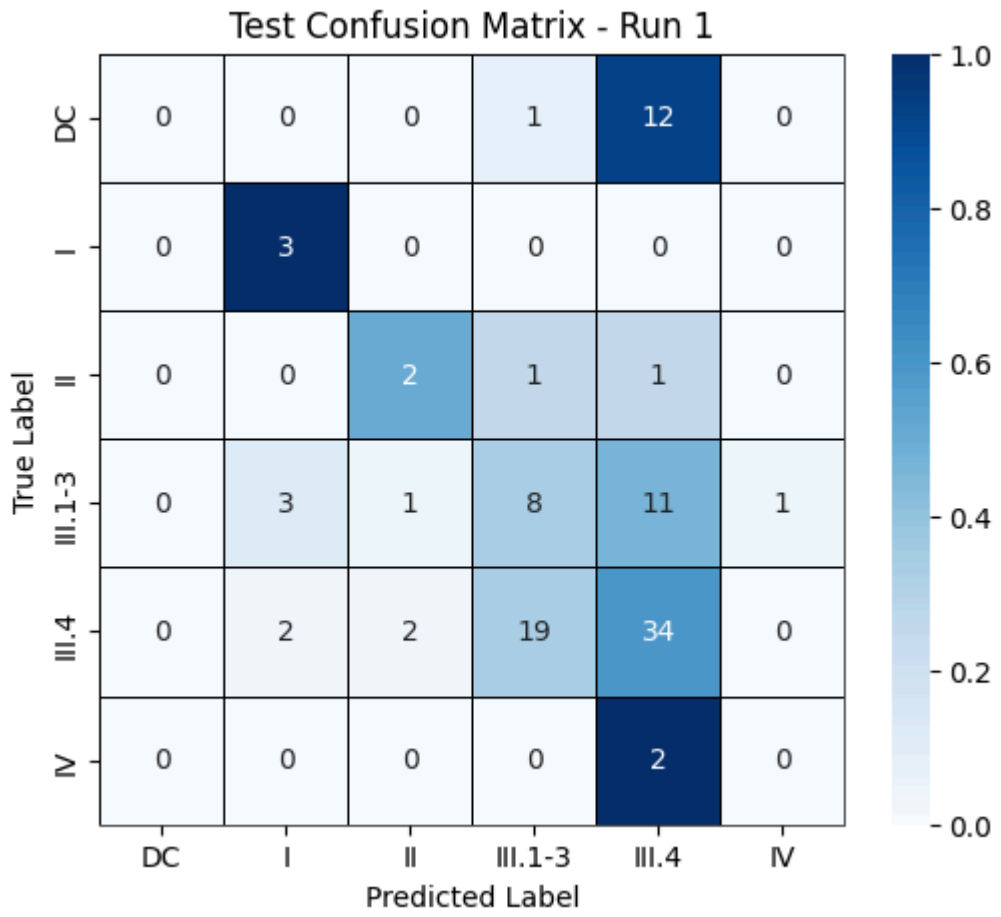

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ———————————————————— **0s** 24ms/step

== Test Set Evaluation - Run 2 ==

	precision	recall	f1-score	support
DC	0.5714	0.3333	0.4211	12
I	0.1905	1.0000	0.3200	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2778	0.4000	0.3279	25
III.4	0.6111	0.1964	0.2973	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.2816	103
macro avg	0.2751	0.3216	0.2277	103
weighted avg	0.4736	0.2816	0.3027	103

Balanced Accuracy: 0.32162698412698415

Confusion Matrix:
[[4 1 1 1 4 1]
 [0 4 0 0 0 0]
 [0 0 0 1 2 1]
 [1 8 4 10 1 1]
 [2 6 11 24 11 2]
 [0 2 0 0 0 0]]

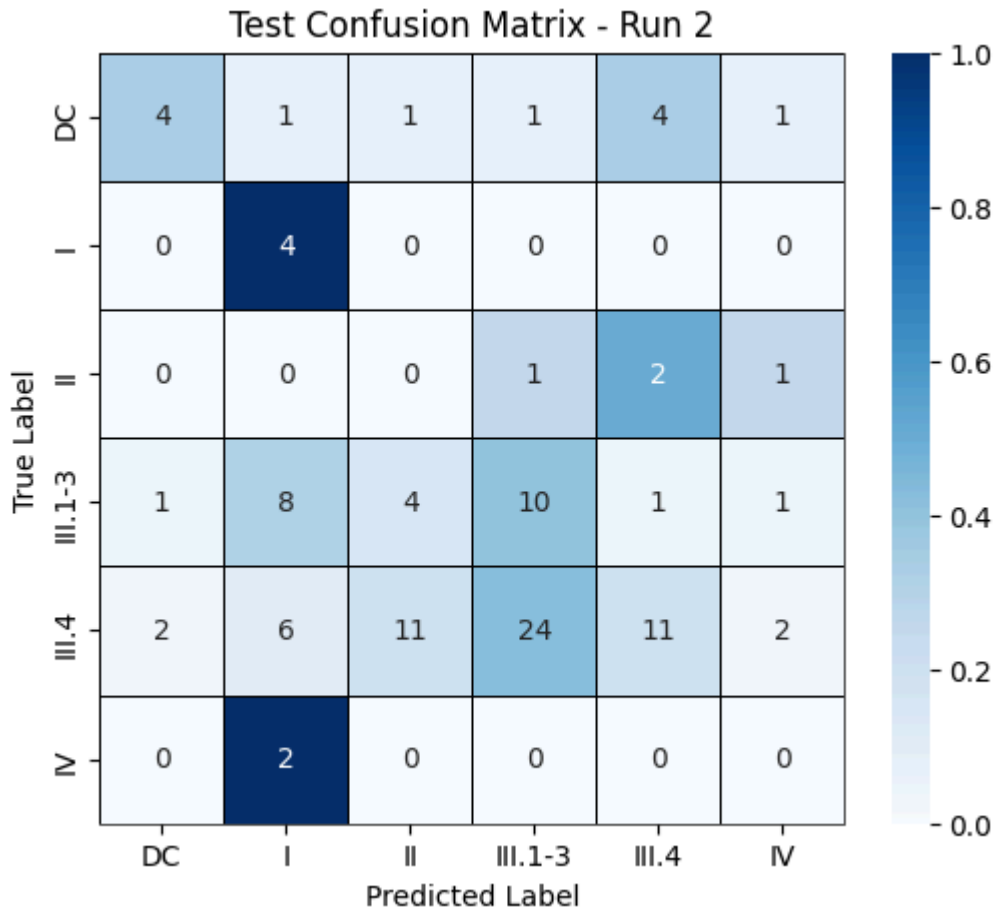

==========
AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET
==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.5714	0.1600	0.2500	25
I	0.2414	1.0000	0.3889	7
II	0.0952	0.2500	0.1379	8
III.1-3	0.2769	0.3673	0.3158	49
III.4	0.5769	0.3982	0.4712	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3689	206

macro avg	0.2936	0.3626	0.2606	206
weighted avg	0.4636	0.3689	0.3825	206

Aggregated Balanced Accuracy: 0.3626
Aggregated Confusion Matrix:
[[4 1 1 2 16 1]
 [0 7 0 0 0 0]
 [0 0 2 2 3 1]
 [1 11 5 18 12 2]
 [2 8 13 43 45 2]
 [0 2 0 0 2 0]]

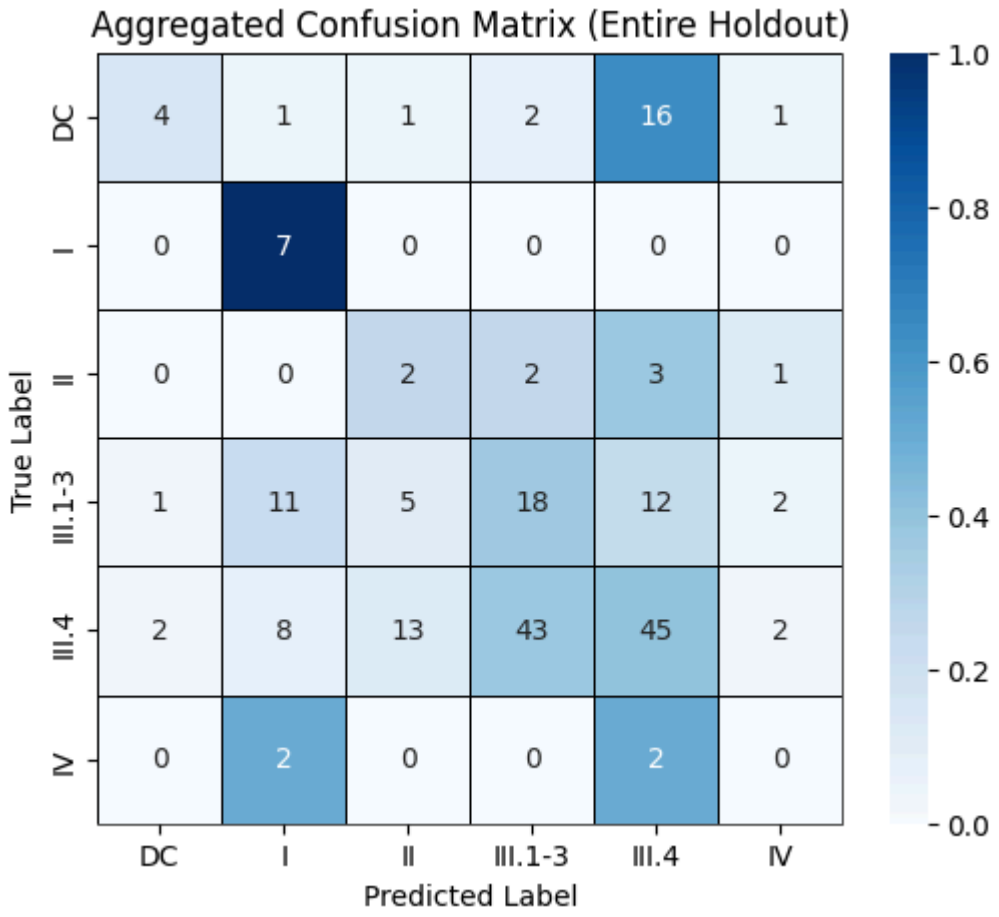

== Average Balanced Accuracy Across Runs ==
Run 1: 0.4050
Run 2: 0.3216
Average: 0.3633

==
RUN 1: Using Split A for validation, Split B for testing
==
These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']
4/4 ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ **0s** 21ms/step

== Test Set Evaluation - Run 1 ==

	precision	recall	f1-score	support
DC	1.0000	0.1538	0.2667	13
I	0.2500	1.0000	0.4000	3
II	0.5000	0.5000	0.5000	4
III.1-3	0.1940	0.5417	0.2857	24
III.4	0.6471	0.1930	0.2973	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.3010	103
macro avg	0.4318	0.3981	0.2916	103
weighted avg	0.5562	0.3010	0.2958	103

Balanced Accuracy: 0.3980825461088619
Confusion Matrix:
[[2 0 0 9 2 0]
[0 3 0 0 0 0]
[0 0 2 2 0 0]
[0 5 1 13 4 1]
[0 4 1 41 11 0]
[0 0 0 2 0 0]]

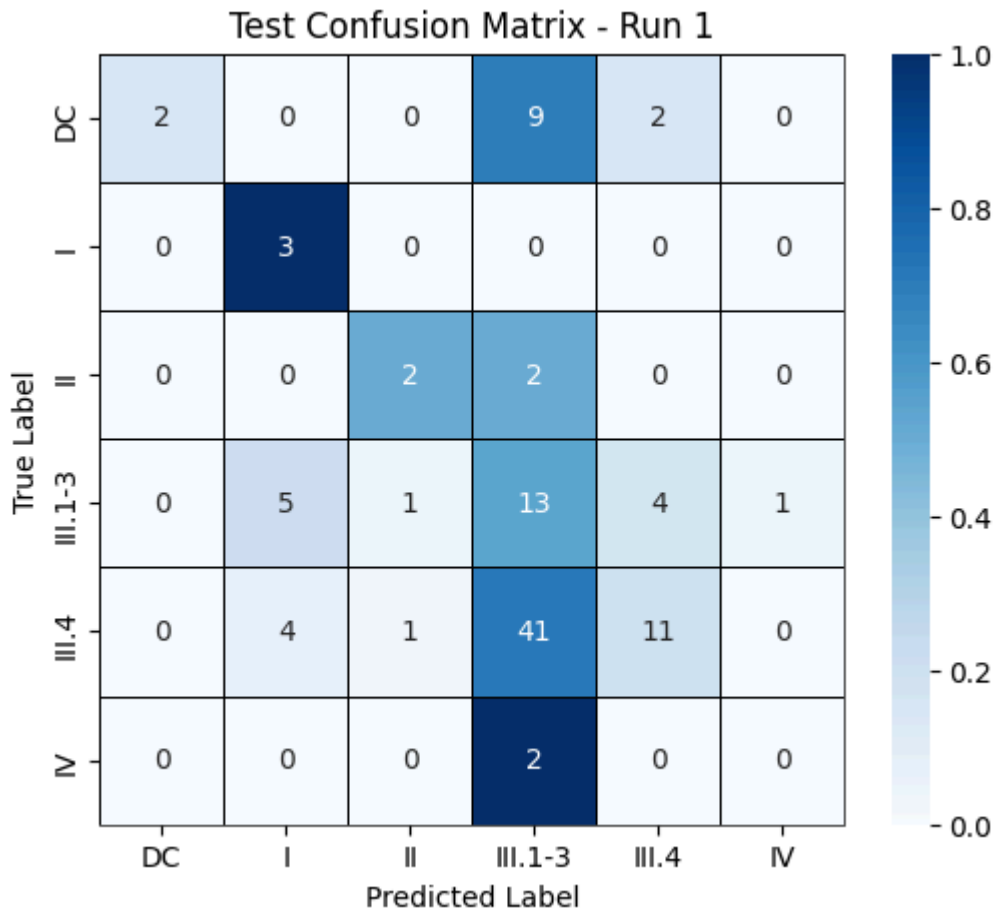

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ━━━━━━━━━━━━━━━━━━━━ **0s** 21ms/step

== Test Set Evaluation - Run 2 ==

	precision	recall	f1-score	support
DC	0.5000	0.0833	0.1429	12
I	0.2353	1.0000	0.3810	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2593	0.5600	0.3544	25
III.4	0.7059	0.2143	0.3288	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.3010	103
macro avg	0.2834	0.3096	0.2012	103
weighted avg	0.5141	0.3010	0.2962	103

Balanced Accuracy: 0.3096031746031746

Confusion Matrix:
[[1 0 1 3 4 3]
 [0 4 0 0 0 0]
 [1 0 0 3 0 0]
 [0 6 1 14 1 3]
 [0 5 3 34 12 2]
 [0 2 0 0 0 0]]

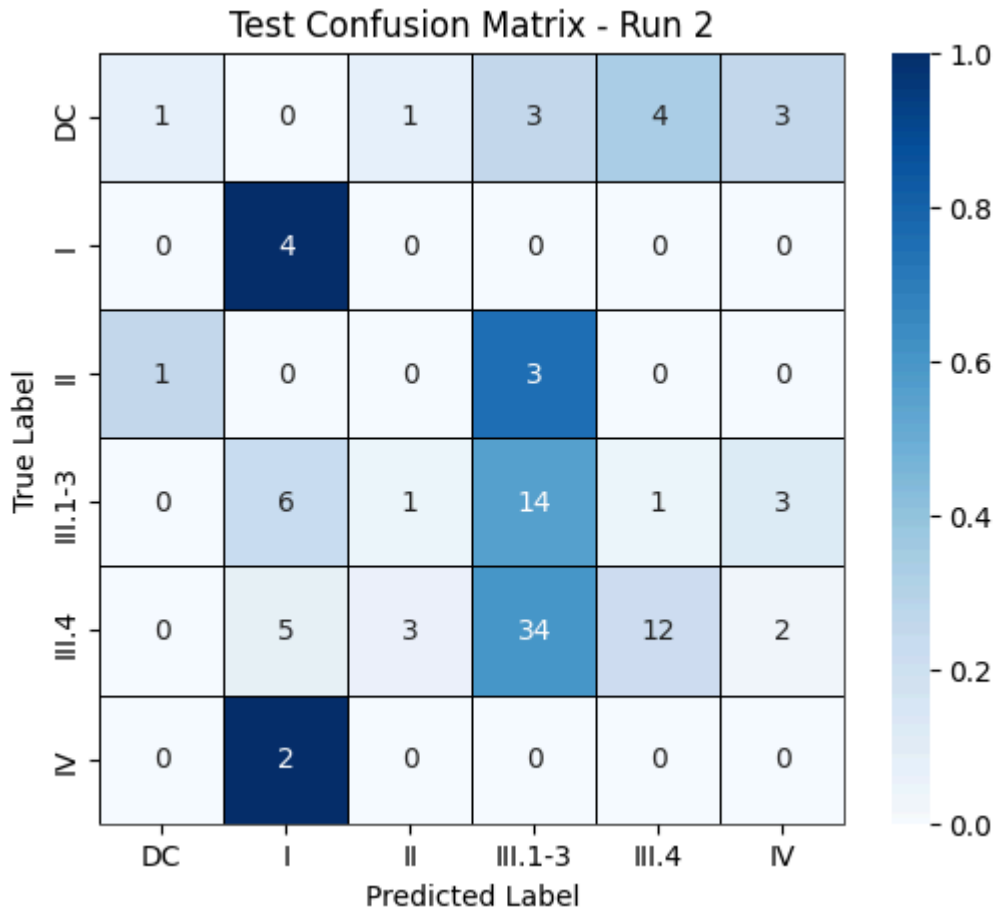

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.7500	0.1200	0.2069	25
I	0.2414	1.0000	0.3889	7
II	0.2222	0.2500	0.2353	8
III.1-3	0.2231	0.5510	0.3176	49
III.4	0.6765	0.2035	0.3129	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3010	206

macro avg	0.3522	0.3541	0.2436	206
weighted avg	0.5320	0.3010	0.2947	206

Aggregated Balanced Accuracy: 0.3541
Aggregated Confusion Matrix:
[[3 0 1 12 6 3]
 [0 7 0 0 0 0]
 [1 0 2 5 0 0]
 [0 11 2 27 5 4]
 [0 9 4 75 23 2]
 [0 2 0 2 0 0]]

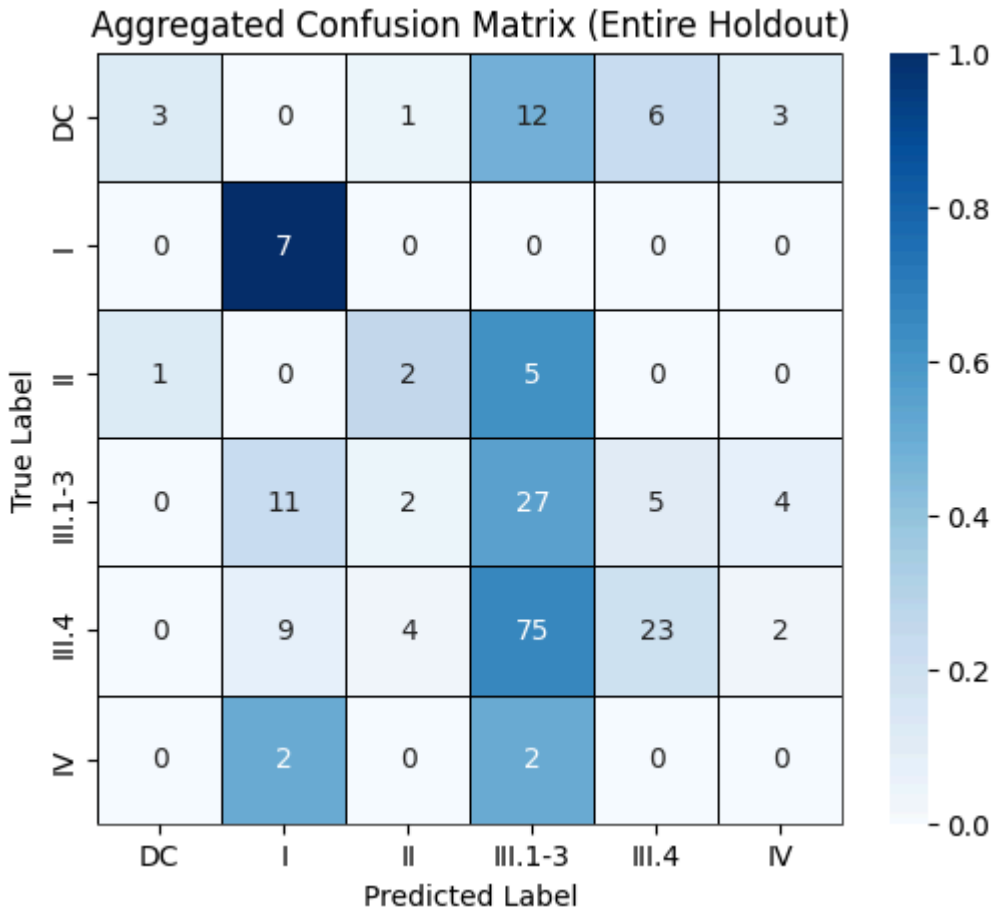

== Average Balanced Accuracy Across Runs ==
Run 1: 0.3981
Run 2: 0.3096
Average: 0.3538

==
RUN 1: Using Split A for validation, Split B for testing
==
These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']
4/4 ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ **0s** 26ms/step

== Test Set Evaluation - Run 1 ==

	precision	recall	f1-score	support
DC	0.0000	0.0000	0.0000	13
I	0.3750	1.0000	0.5455	3
II	0.4000	0.5000	0.4444	4
III.1-3	0.1803	0.4583	0.2588	24
III.4	0.5172	0.2632	0.3488	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.3010	103
macro avg	0.2454	0.3702	0.2663	103
weighted avg	0.3547	0.3010	0.2865	103

Balanced Accuracy: 0.37024853801169594
Confusion Matrix:
```
[[ 0  0  0  8  5  0]
 [ 0  3  0  0  0  0]
 [ 0  0  2  2  0  0]
 [ 0  3  1 11  9  0]
 [ 0  2  2 38 15  0]
 [ 0  0  0  2  0  0]]
```

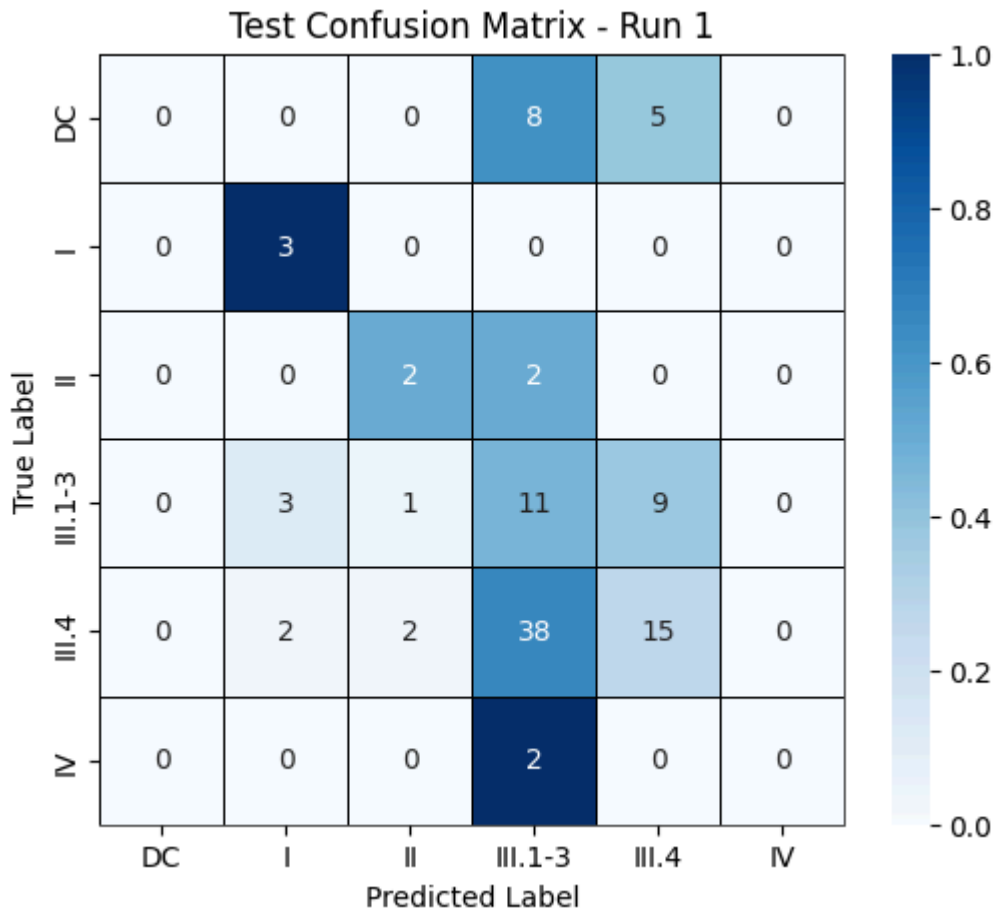

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ━━━━━━━━━━━━━━━━━━━━ **0s** 24ms/step

== Test Set Evaluation - Run 2 ==

	precision	recall	f1-score	support
DC	0.4000	0.1667	0.2353	12
I	0.2857	1.0000	0.4444	4
II	0.3333	0.2500	0.2857	4
III.1-3	0.2167	0.5200	0.3059	25
III.4	0.9286	0.2321	0.3714	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.3204	103
macro avg	0.3607	0.3615	0.2738	103
weighted avg	0.6281	0.3204	0.3320	103

Balanced Accuracy: 0.36146825396825394

Confusion Matrix:
[[2 0 1 6 0 3]
 [0 4 0 0 0 0]
 [0 0 1 3 0 0]
 [2 6 0 13 1 3]
 [1 2 1 38 13 1]
 [0 2 0 0 0 0]]

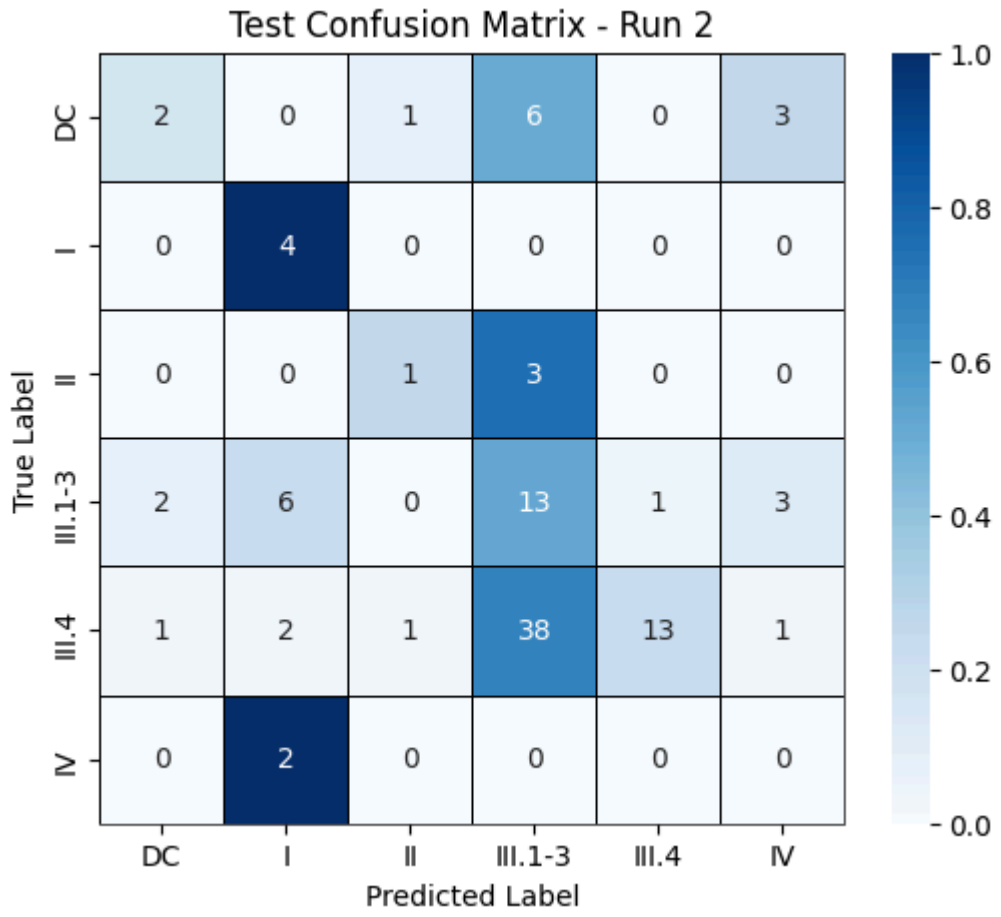

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.4000	0.0800	0.1333	25
I	0.3182	1.0000	0.4828	7
II	0.3750	0.3750	0.3750	8
III.1-3	0.1983	0.4898	0.2824	49
III.4	0.6512	0.2478	0.3590	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3107	206

macro avg	0.3238	0.3654	0.2721	206
weighted avg	0.4783	0.3107	0.3112	206

Aggregated Balanced Accuracy: 0.3654
Aggregated Confusion Matrix:
[[2 0 1 14 5 3]
[0 7 0 0 0 0]
[0 0 3 5 0 0]
[2 9 1 24 10 3]
[1 4 3 76 28 1]
[0 2 0 2 0 0]]

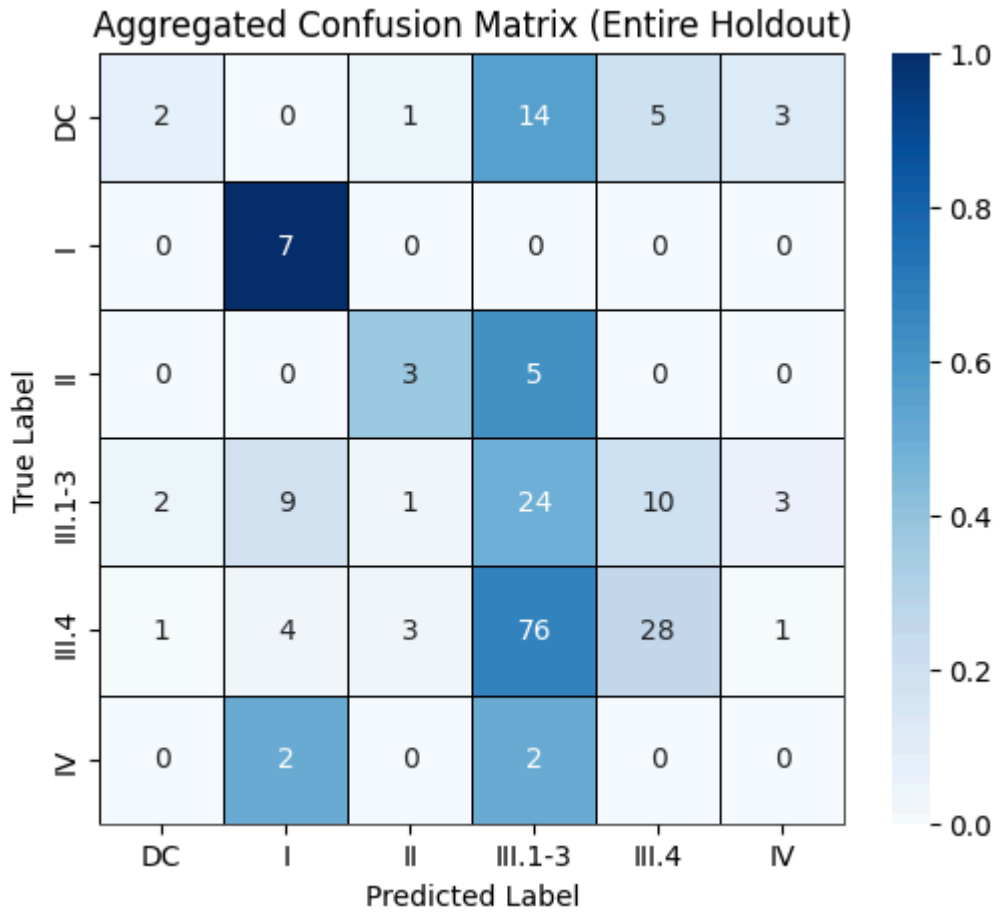

== Average Balanced Accuracy Across Runs ==
Run 1: 0.3702
Run 2: 0.3615
Average: 0.3659

==========
RUN 1: Using Split A for validation, Split B for testing
==========
These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']
4/4 ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ **0s** 25ms/step

== Test Set Evaluation - Run 1 ==

	precision	recall	f1-score	support
DC	0.0000	0.0000	0.0000	13
I	0.1364	1.0000	0.2400	3
II	0.2857	0.5000	0.3636	4
III.1-3	0.1778	0.3333	0.2319	24
III.4	0.5185	0.2456	0.3333	57
IV	0.0000	0.0000	0.0000	2
accuracy			0.2621	103
macro avg	0.1864	0.3465	0.1948	103
weighted avg	0.3434	0.2621	0.2596	103

Balanced Accuracy: 0.34649122807017546
Confusion Matrix:
[[0 0 1 5 7 0]
[0 3 0 0 0 0]
[0 1 2 1 0 0]
[0 8 2 8 5 1]
[1 10 2 30 14 0]
[0 0 0 1 1 0]]

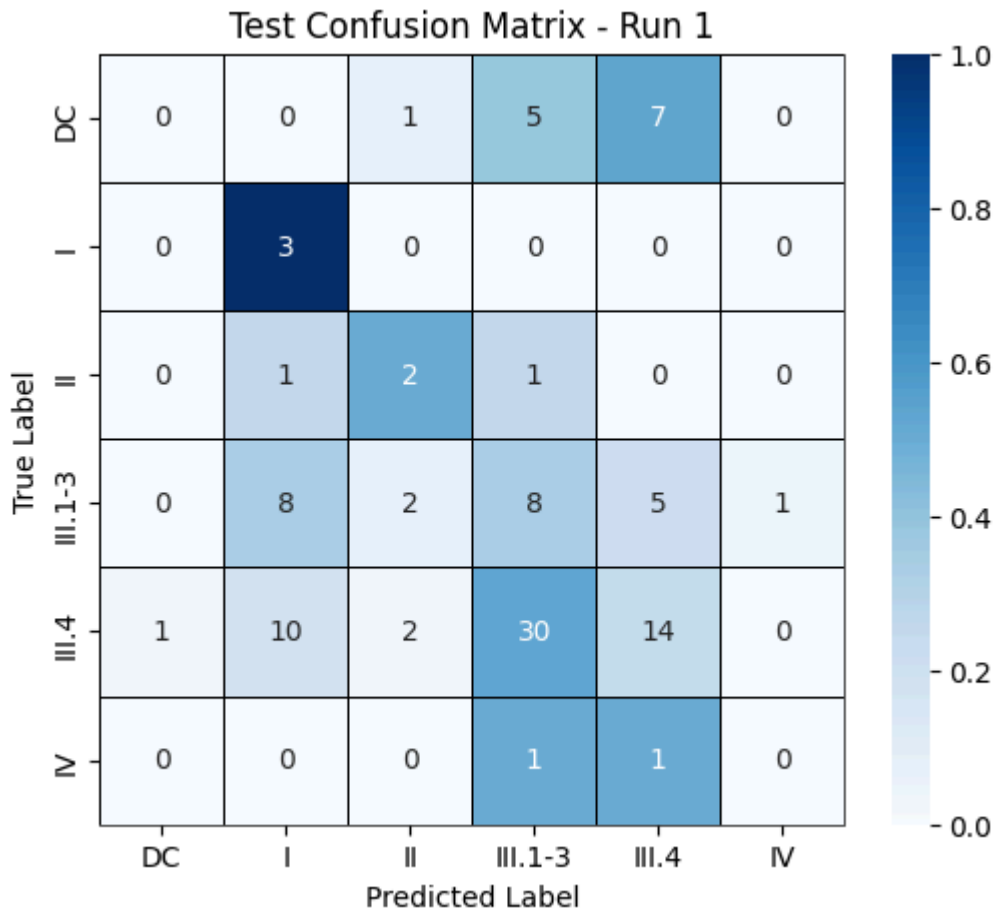

==========

RUN 2: Using Split B for validation, Split A for testing (SWAP)

==========

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ━━━━━━━━━━━━━━━━━━━━ **0s** 24ms/step

== Test Set Evaluation - Run 2 ==

	precision	recall	f1-score	support
DC	1.0000	0.3333	0.5000	12
I	0.2667	1.0000	0.4211	4
II	0.0000	0.0000	0.0000	4
III.1-3	0.2414	0.2800	0.2593	25
III.4	0.6667	0.6071	0.6355	56
IV	0.0000	0.0000	0.0000	2
accuracy			0.4757	103
macro avg	0.3625	0.3701	0.3026	103
weighted avg	0.5479	0.4757	0.4831	103

Balanced Accuracy: 0.37007936507936506

Confusion Matrix:
[[4 0 0 1 5 2]
 [0 4 0 0 0 0]
 [0 0 0 4 0 0]
 [0 5 0 7 12 1]
 [0 4 1 17 34 0]
 [0 2 0 0 0 0]]

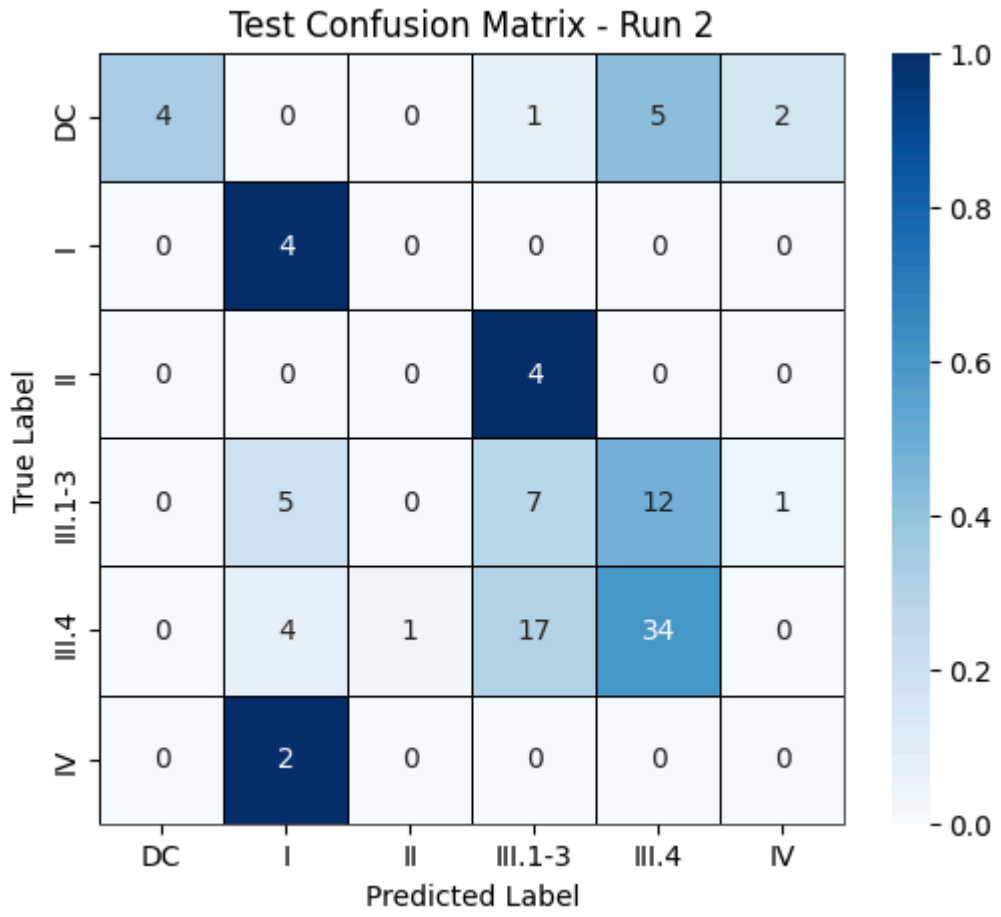

==========

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

==========

== Aggregated Classification Report ==

	precision	recall	f1-score	support
DC	0.8000	0.1600	0.2667	25
I	0.1892	1.0000	0.3182	7
II	0.2500	0.2500	0.2500	8
III.1-3	0.2027	0.3061	0.2439	49
III.4	0.6154	0.4248	0.5026	113
IV	0.0000	0.0000	0.0000	4
accuracy			0.3689	206

macro avg	0.3429	0.3568	0.2636	206
weighted avg	0.4990	0.3689	0.3866	206

Aggregated Balanced Accuracy: 0.3568
Aggregated Confusion Matrix:
[[4 0 1 6 12 2]
 [0 7 0 0 0 0]
 [0 1 2 5 0 0]
 [0 13 2 15 17 2]
 [1 14 3 47 48 0]
 [0 2 0 1 1 0]]

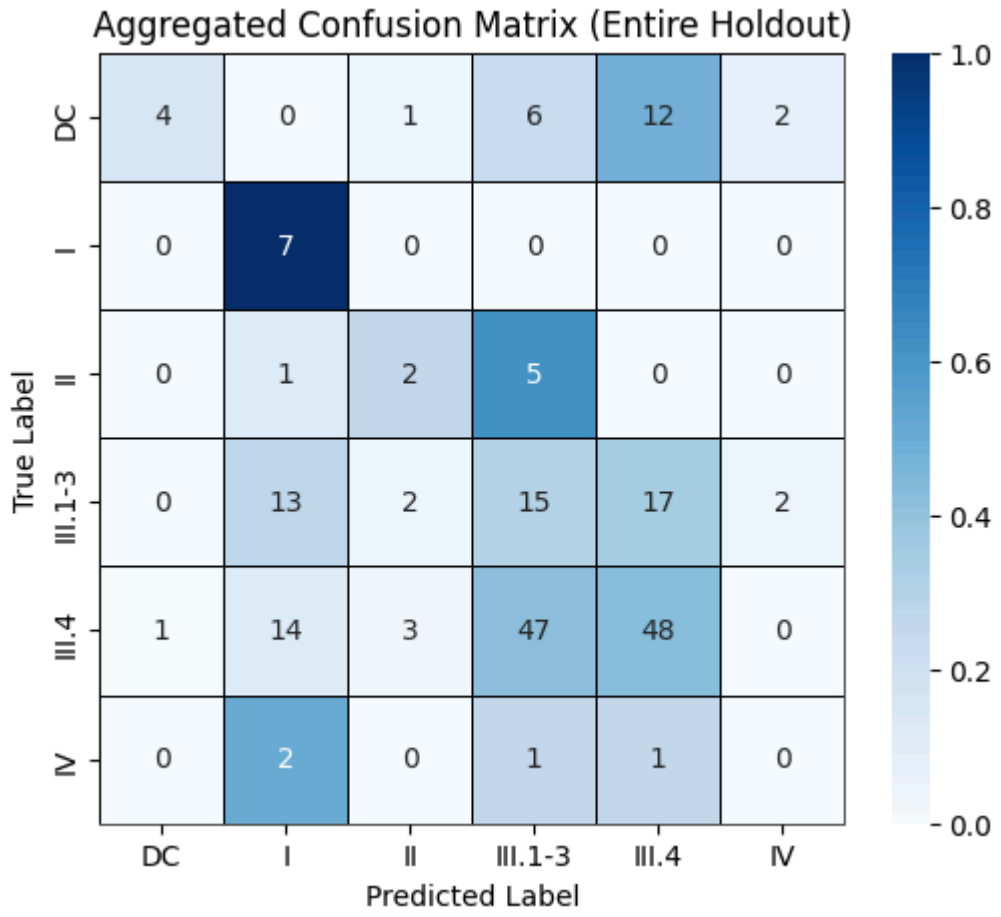

== Average Balanced Accuracy Across Runs ==
Run 1: 0.3465
Run 2: 0.3701
Average: 0.3583

== Average Balanced Accuracy Across Runs ==
Run 1: 0.3494
Run 2: 0.2712
Average: 0.3103

[0.40568704472939615, 0.36259617121184756, 0.35409337186201917, 0.3654305881644693, 0.3568168683402564]

0.369 ± 0.008