

HIERARCHICAL

=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

-----

STAGE 1 (I vs Rest) - AGGREGATED

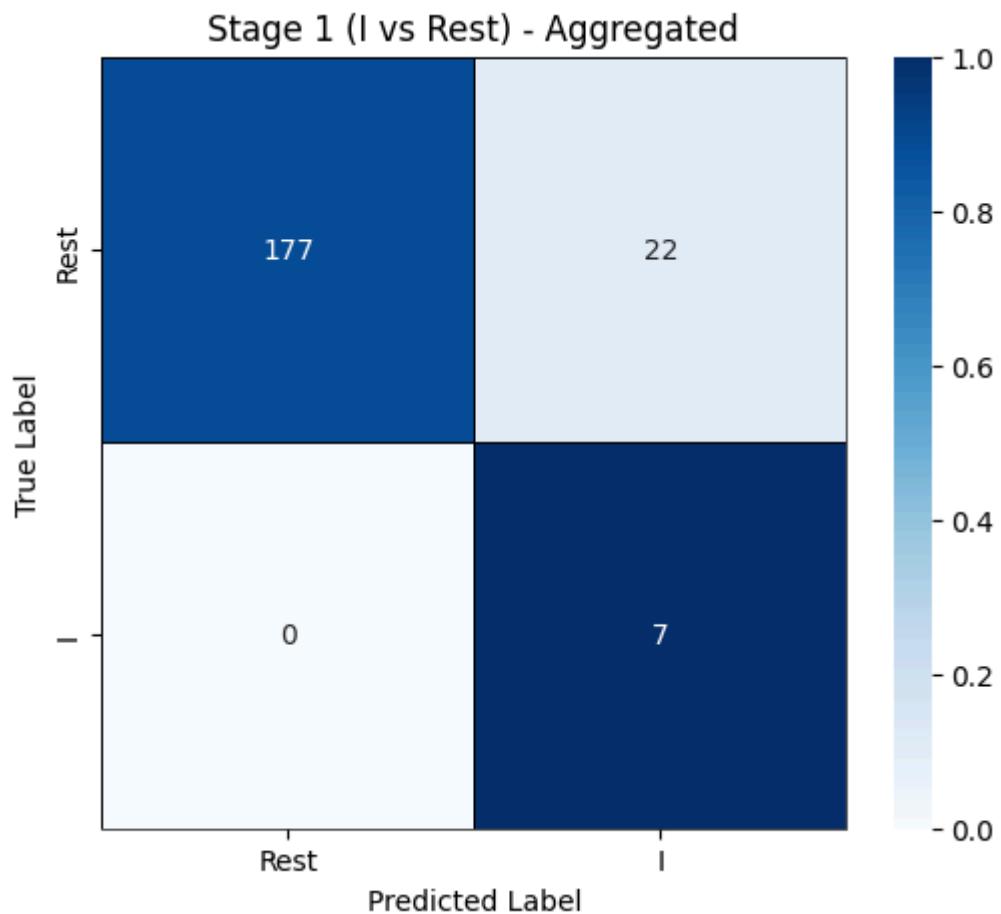
-----

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.0000	0.8894	0.9415	199
1	0.2414	1.0000	0.3889	7

accuracy		0.8932	206
macro avg	0.6207	0.9447	0.6652
weighted avg	0.9742	0.8932	0.9227
			206

Balanced Accuracy: 0.9447



-----

STAGE 2 (DC vs Rest) - AGGREGATED

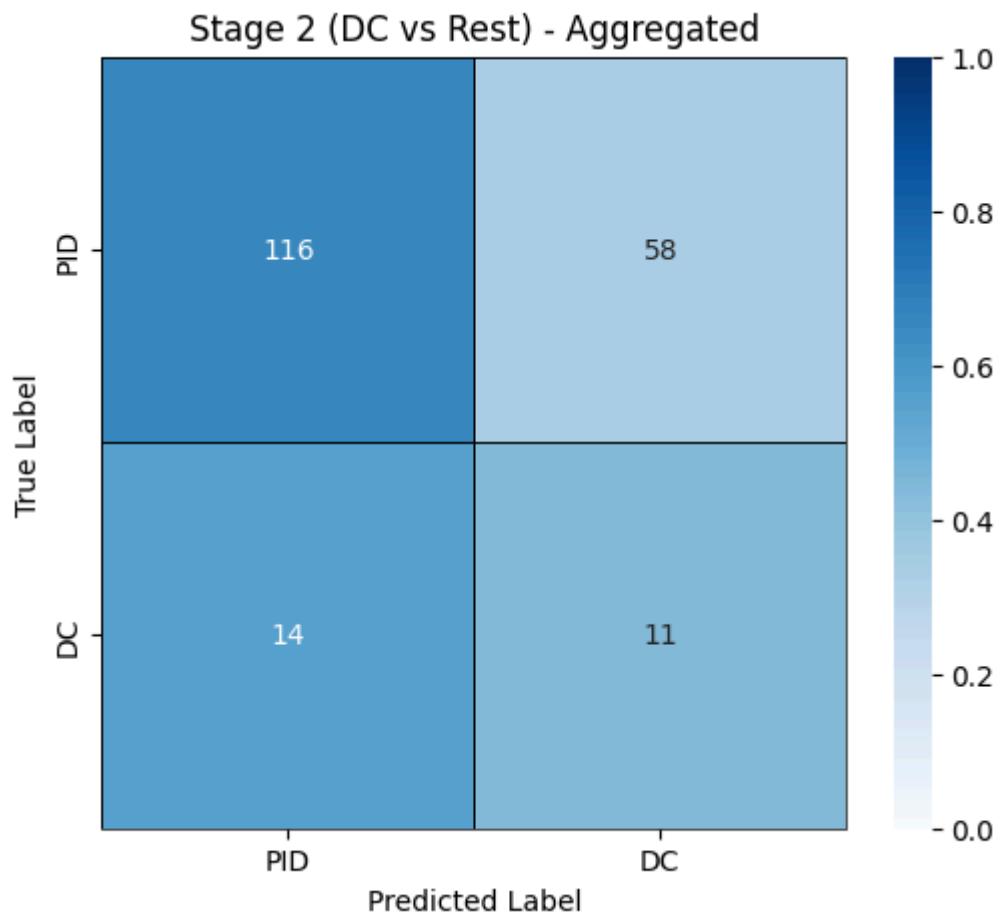
-----

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.8923	0.6667	0.7632	174
1	0.1594	0.4400	0.2340	25

accuracy		0.6382	199	
macro avg	0.5259	0.5533	0.4986	199
weighted avg	0.8002	0.6382	0.6967	199

Balanced Accuracy: 0.5533



---

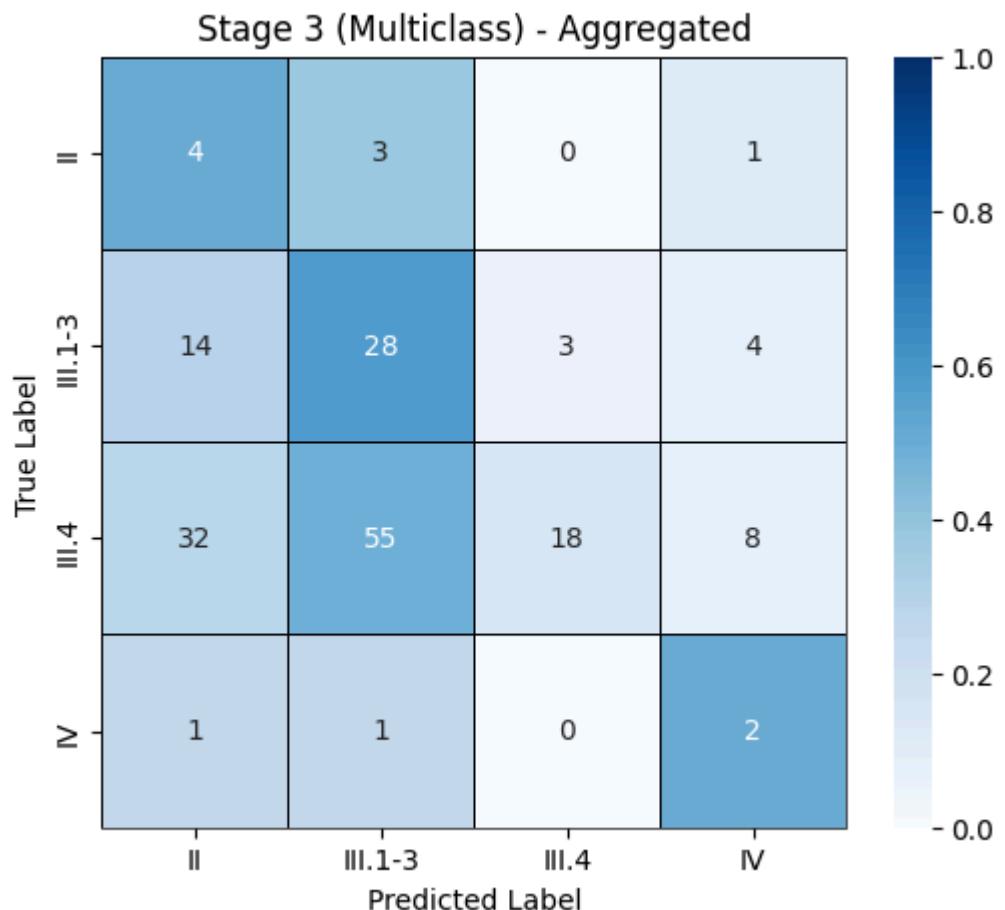
### STAGE 3 (Multiclass) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.0784	0.5000	0.1356	8
1	0.3218	0.5714	0.4118	49
2	0.8571	0.1593	0.2687	113
3	0.1333	0.5000	0.2105	4

accuracy		0.2989	174	
macro avg	0.3477	0.4327	0.2566	174
weighted avg	0.6540	0.2989	0.3015	174

Balanced Accuracy: 0.4327



---

**SOFT-GATED OVERALL - AGGREGATED**

---

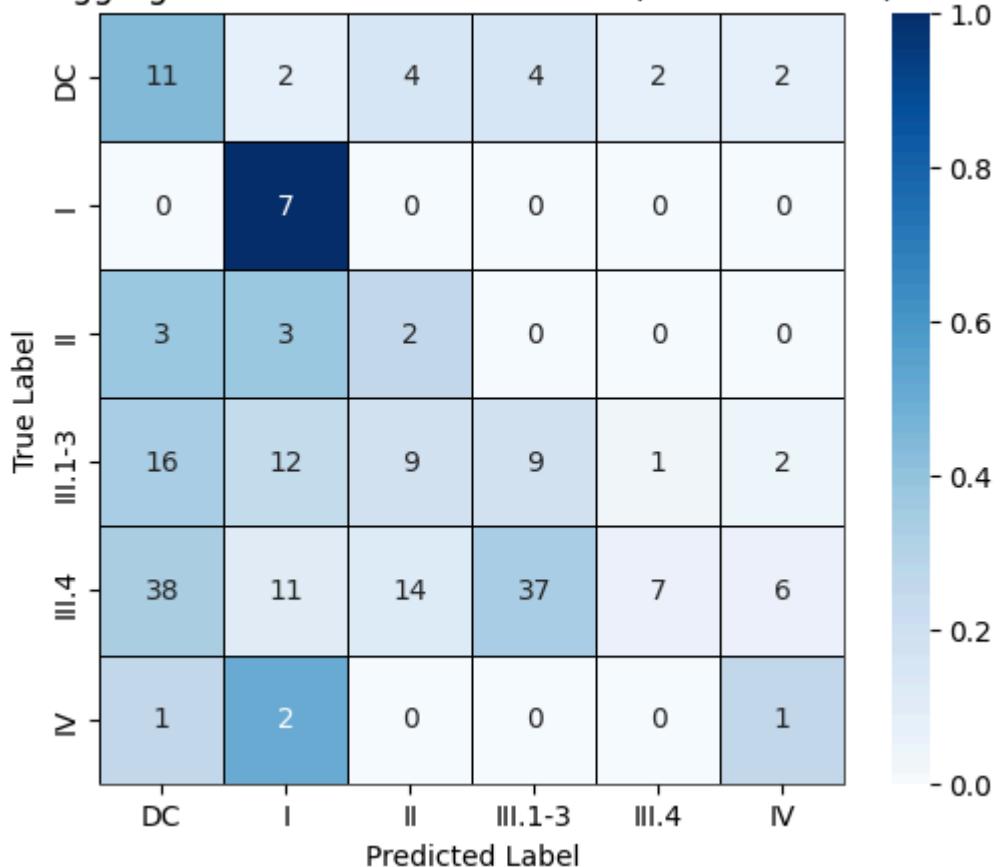
== Aggregated Classification Report ==  
precision recall f1-score support

DC	0.1594	0.4400	0.2340	25
I	0.1892	1.0000	0.3182	7
II	0.0690	0.2500	0.1081	8
III.1-3	0.1800	0.1837	0.1818	49
III.4	0.7000	0.0619	0.1138	113
IV	0.0909	0.2500	0.1333	4

accuracy		0.1796	206	
macro avg	0.2314	0.3643	0.1816	206
weighted avg	0.4570	0.1796	0.1517	206

Aggregated Balanced Accuracy: 0.3643

Aggregated Final Confusion Matrix (Entire Holdout)



== Average Stage Balanced Accuracies ==

Stage 1 (I vs Rest): 0.9444

Stage 2 (DC vs Rest): 0.5481

Stage 3 (Multiclass): 0.4318

Final (Soft-Gated): 0.3624

---

=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

---



---

STAGE 1 (I vs Rest) - AGGREGATED

---

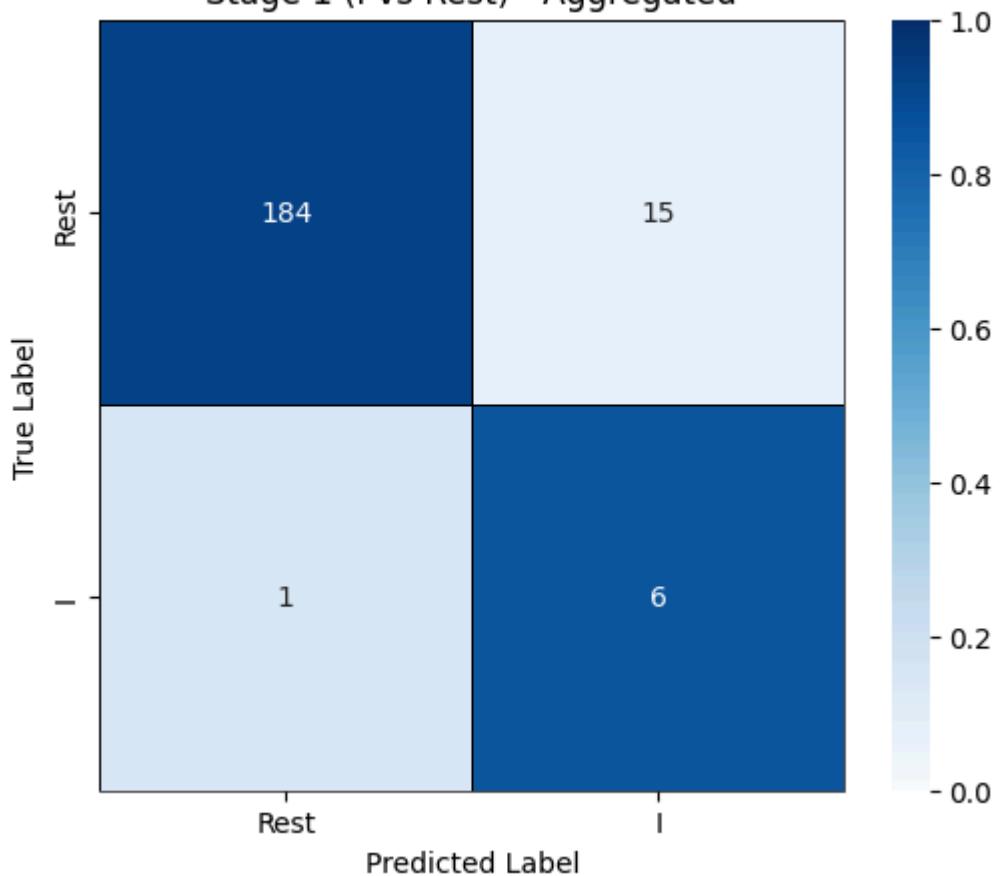
	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9946	0.9246	0.9583	199
1	0.2857	0.8571	0.4286	7

accuracy		0.9223	206	
macro avg	0.6402	0.8909	0.6935	206
weighted avg	0.9705	0.9223	0.9403	206

Balanced Accuracy: 0.8909

Stage 1 (I vs Rest) - Aggregated



---

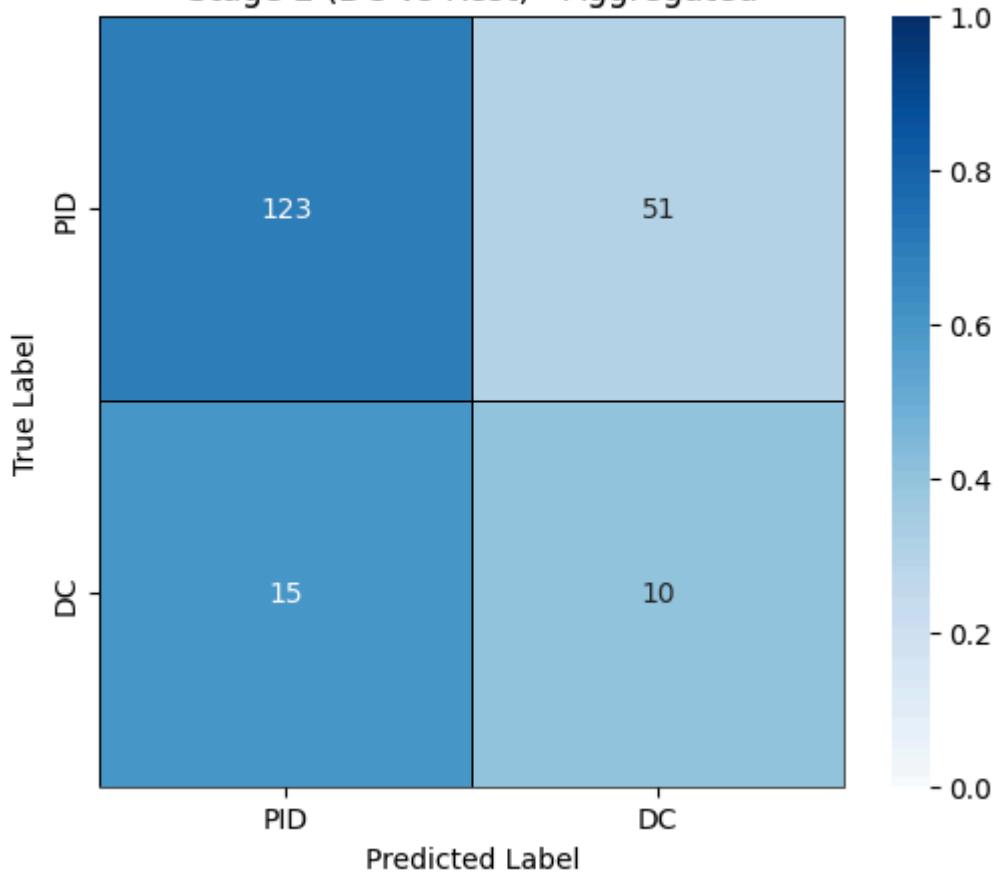
STAGE 2 (DC vs Rest) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.8913	0.7069	0.7885	174
1	0.1639	0.4000	0.2326	25
accuracy		0.6683	199	
macro avg	0.5276	0.5534	0.5105	199
weighted avg	0.7999	0.6683	0.7186	199

Balanced Accuracy: 0.5534

Stage 2 (DC vs Rest) - Aggregated



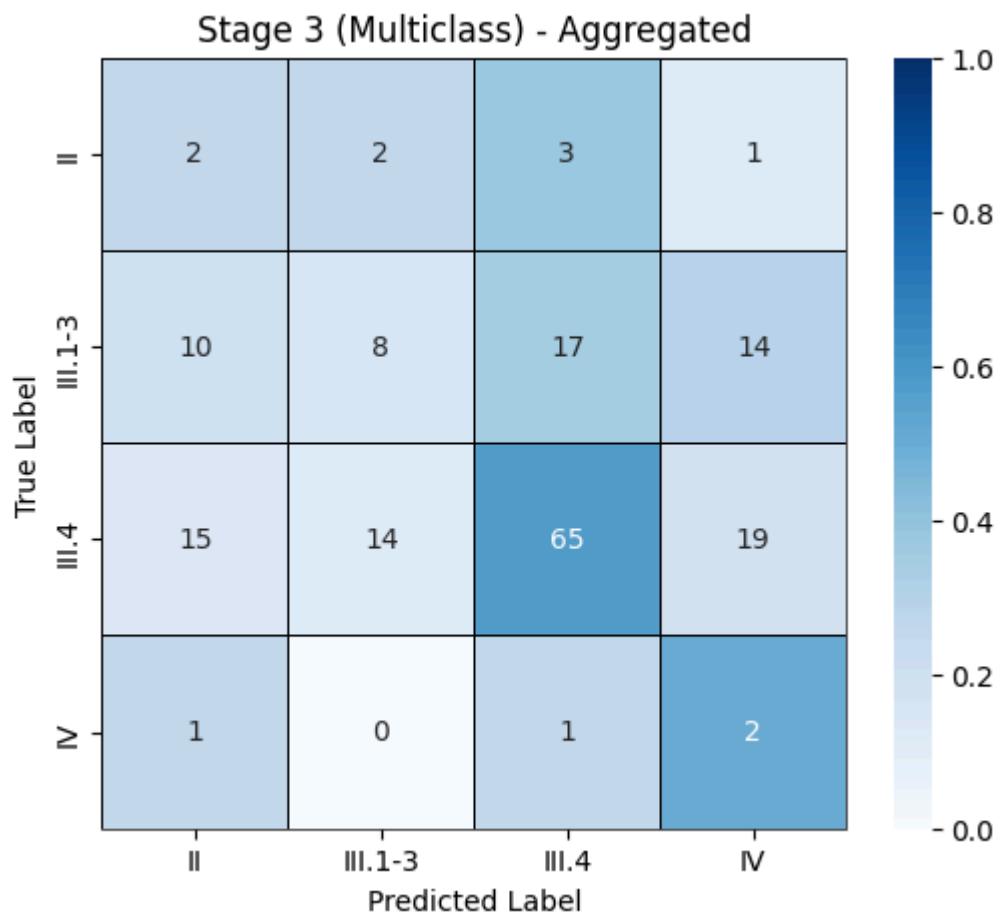
---

STAGE 3 (Multiclass) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.0714	0.2500	0.1111	8
1	0.3333	0.1633	0.2192	49
2	0.7558	0.5752	0.6533	113
3	0.0556	0.5000	0.1000	4
accuracy		0.4425	174	
macro avg	0.3040	0.3721	0.2709	174
weighted avg	0.5893	0.4425	0.4934	174

Balanced Accuracy: 0.3721




---

SOFT-GATED OVERALL - AGGREGATED

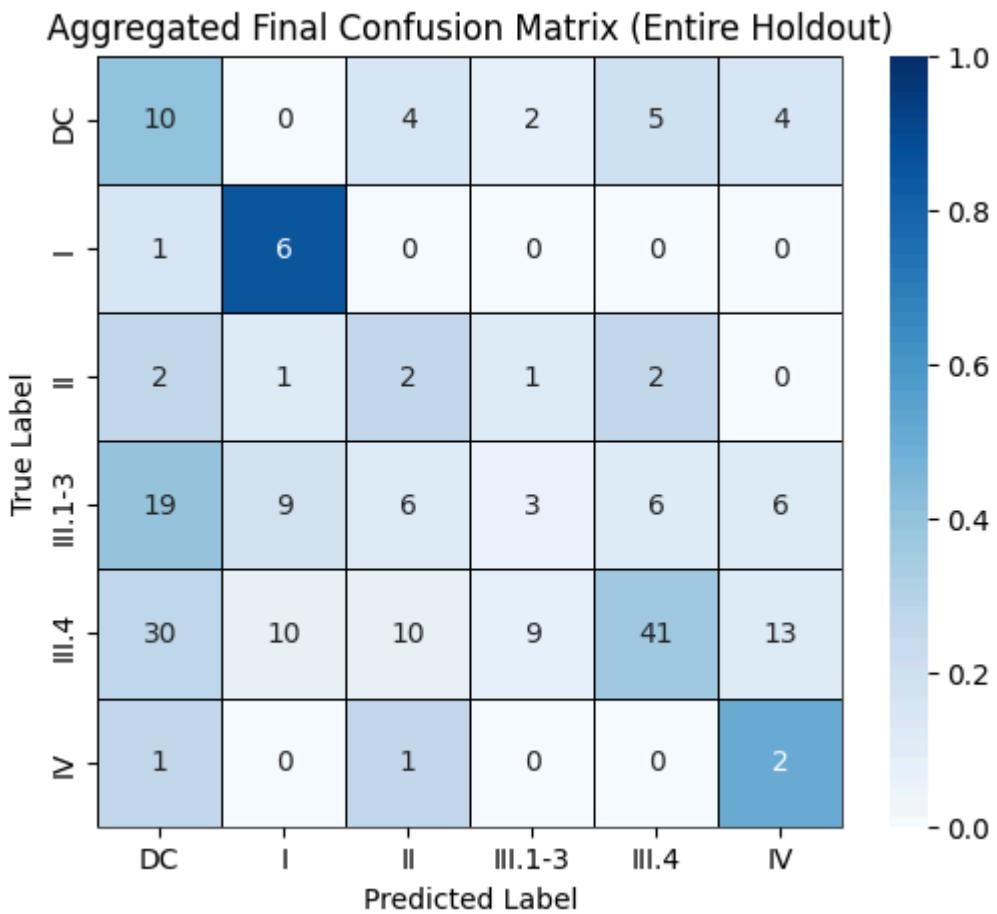
---

== Aggregated Classification Report ==  
precision recall f1-score support

DC	0.1587	0.4000	0.2273	25
I	0.2308	0.8571	0.3636	7
II	0.0870	0.2500	0.1290	8
III.1-3	0.2000	0.0612	0.0938	49
III.4	0.7593	0.3628	0.4910	113
IV	0.0800	0.5000	0.1379	4

accuracy	0.3107	206		
macro avg	0.2526	0.4052	0.2404	206
weighted avg	0.4961	0.3107	0.3393	206

Aggregated Balanced Accuracy: 0.4052



== Average Stage Balanced Accuracies ==

Stage 1 (I vs Rest): 0.8999

Stage 2 (DC vs Rest): 0.5538

Stage 3 (Multiclass): 0.3724

Final (Soft-Gated): 0.4082

---

=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

---



---

STAGE 1 (I vs Rest) - AGGREGATED

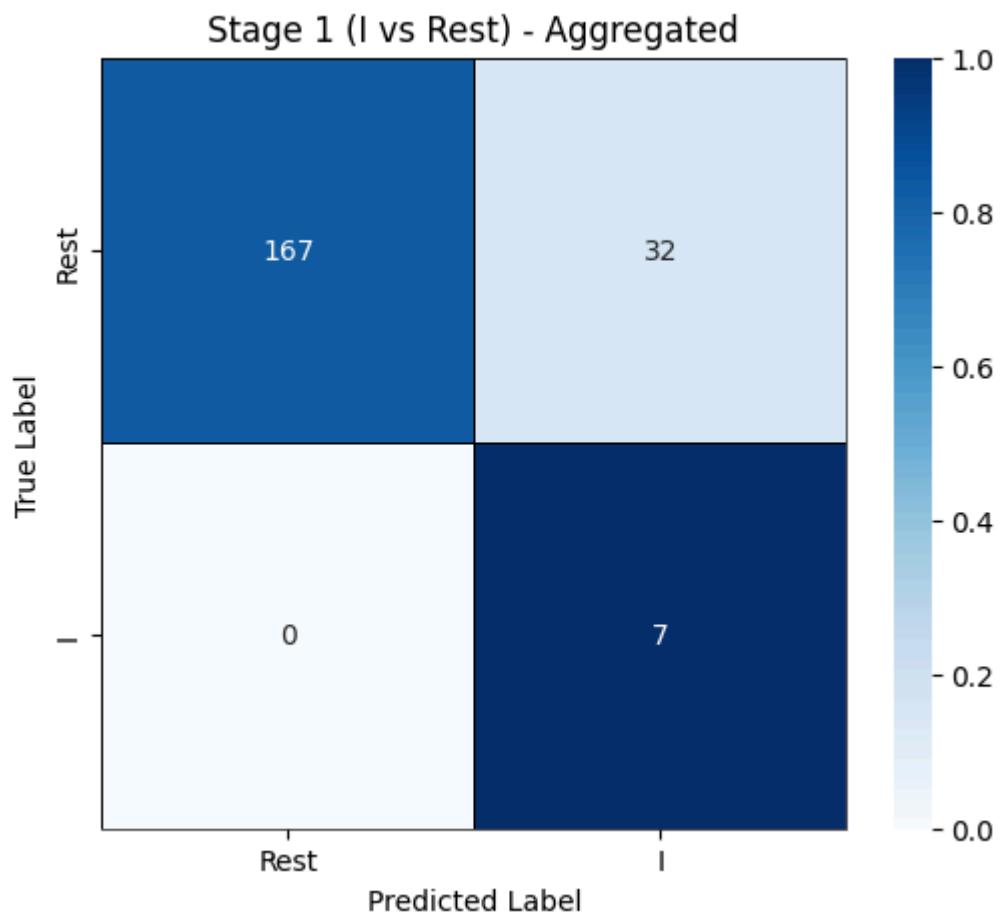
---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.0000	0.8392	0.9126	199
1	0.1795	1.0000	0.3043	7

accuracy		0.8447	206	
macro avg	0.5897	0.9196	0.6085	206
weighted avg	0.9721	0.8447	0.8919	206

Balanced Accuracy: 0.9196



---

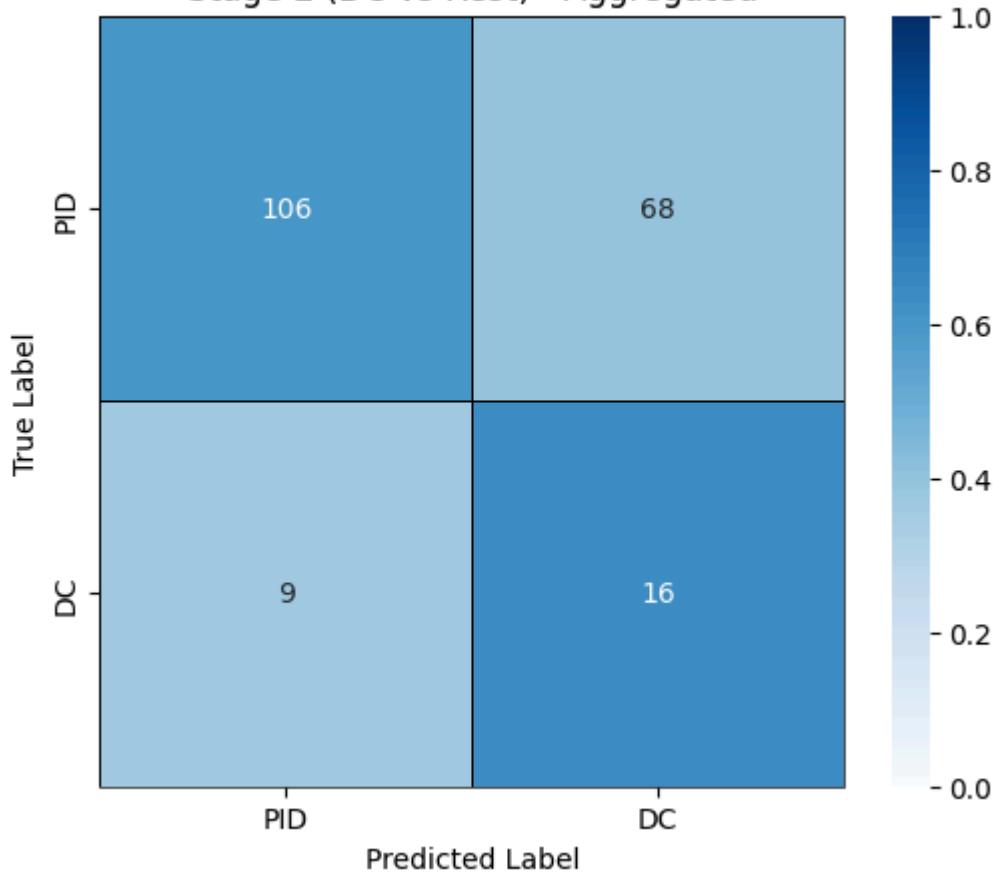
STAGE 2 (DC vs Rest) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.9217	0.6092	0.7336	174
1	0.1905	0.6400	0.2936	25
accuracy		0.6131		199
macro avg	0.5561	0.6246	0.5136	199
weighted avg	0.8299	0.6131	0.6783	199

Balanced Accuracy: 0.6246

Stage 2 (DC vs Rest) - Aggregated



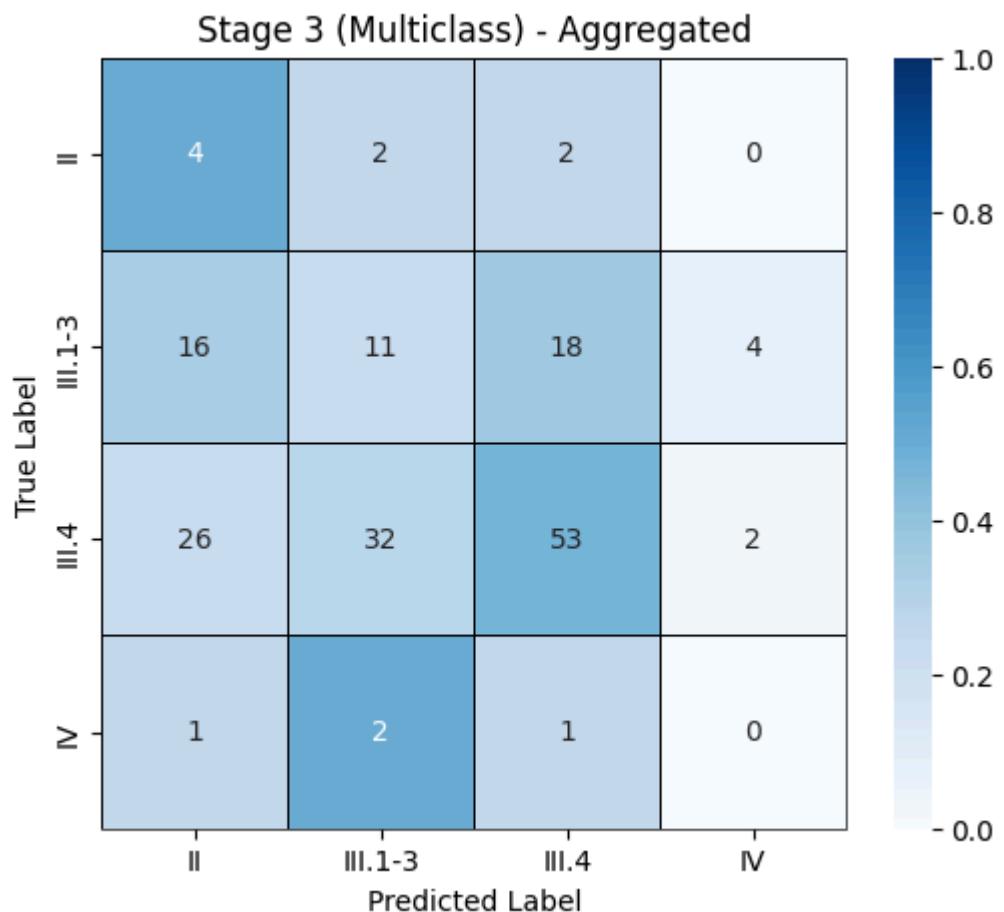
---

STAGE 3 (Multiclass) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.0851	0.5000	0.1455	8
1	0.2340	0.2245	0.2292	49
2	0.7162	0.4690	0.5668	113
3	0.0000	0.0000	0.0000	4
accuracy		0.3908	174	
macro avg	0.2588	0.2984	0.2354	174
weighted avg	0.5350	0.3908	0.4393	174

Balanced Accuracy: 0.2984




---

#### SOFT-GATED OVERALL - AGGREGATED

---

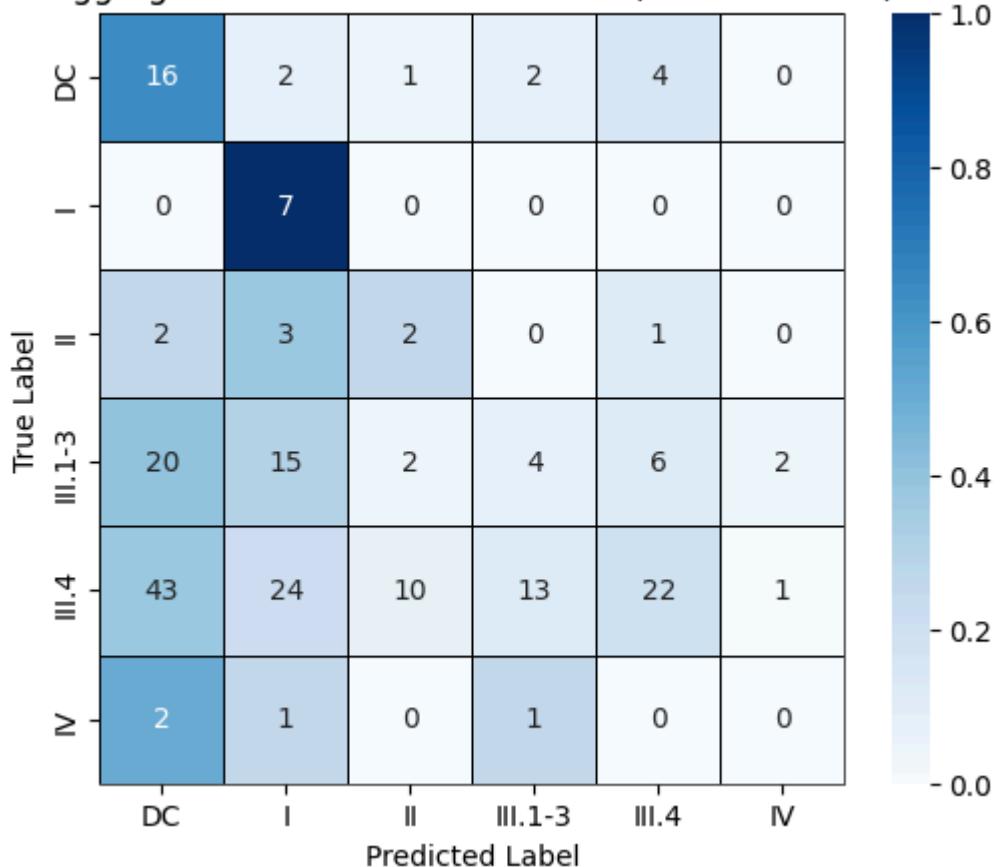
== Aggregated Classification Report ==  
precision recall f1-score support

DC	0.1928	0.6400	0.2963	25
I	0.1346	1.0000	0.2373	7
II	0.1333	0.2500	0.1739	8
III.1-3	0.2000	0.0816	0.1159	49
III.4	0.6667	0.1947	0.3014	113
IV	0.0000	0.0000	0.0000	4

accuracy                    0.2476    206  
macro avg    0.2212    0.3611    0.1875    206  
weighted avg    0.4464    0.2476    0.2437    206

Aggregated Balanced Accuracy: 0.3611

Aggregated Final Confusion Matrix (Entire Holdout)



== Average Stage Balanced Accuracies ==

Stage 1 (I vs Rest): 0.9196

Stage 2 (DC vs Rest): 0.6203

Stage 3 (Multiclass): 0.2982

Final (Soft-Gated): 0.3597

---

#### ===== AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET =====

---



---

##### STAGE 1 (I vs Rest) - AGGREGATED

---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.9946	0.9246	0.9583	199
1	0.2857	0.8571	0.4286	7

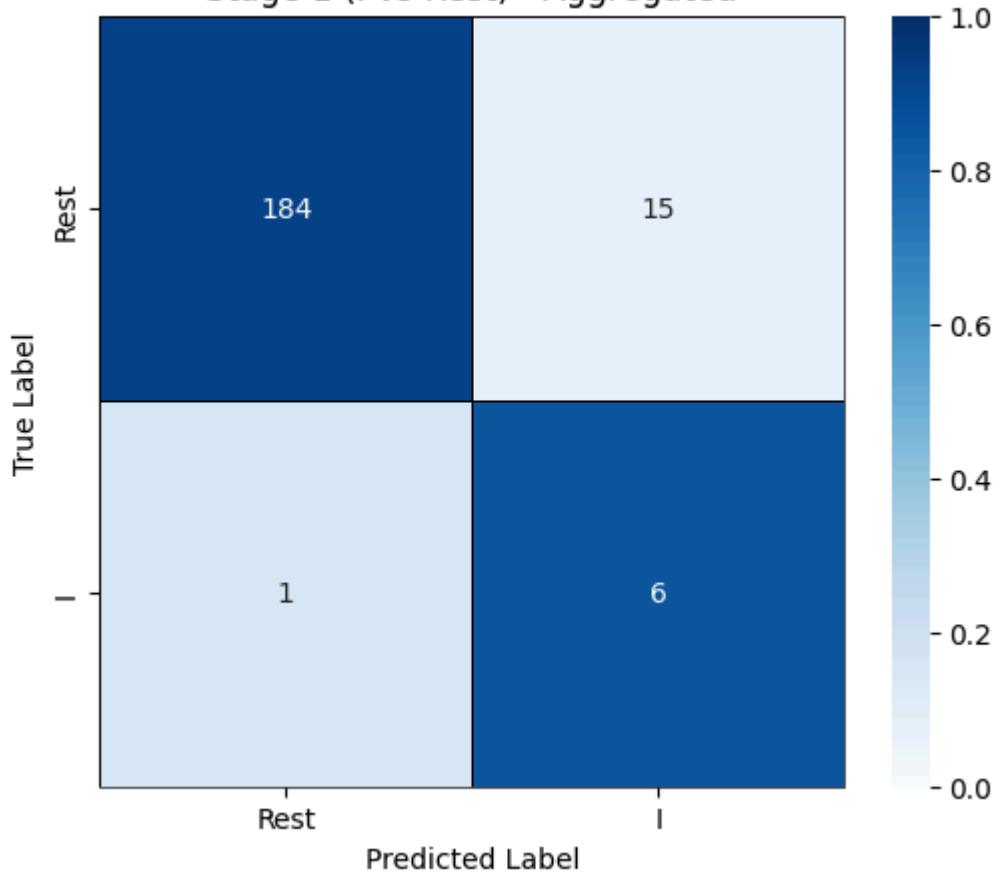
accuracy		0.9223	206	
----------	--	--------	-----	--

macro avg	0.6402	0.8909	0.6935	206
-----------	--------	--------	--------	-----

weighted avg	0.9705	0.9223	0.9403	206
--------------	--------	--------	--------	-----

Balanced Accuracy: 0.8909

Stage 1 (I vs Rest) - Aggregated



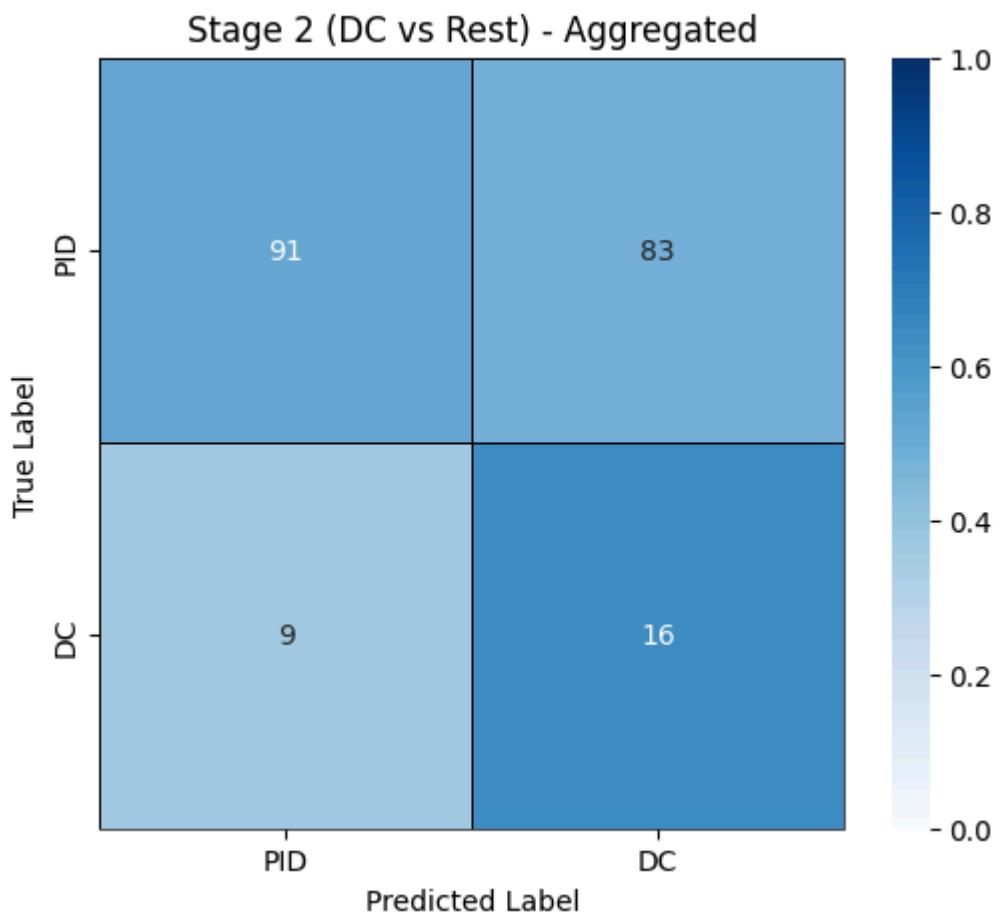
---

STAGE 2 (DC vs Rest) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.9100	0.5230	0.6642	174
1	0.1616	0.6400	0.2581	25
accuracy		0.5377		199
macro avg	0.5358	0.5815	0.4611	199
weighted avg	0.8160	0.5377	0.6132	199

Balanced Accuracy: 0.5815



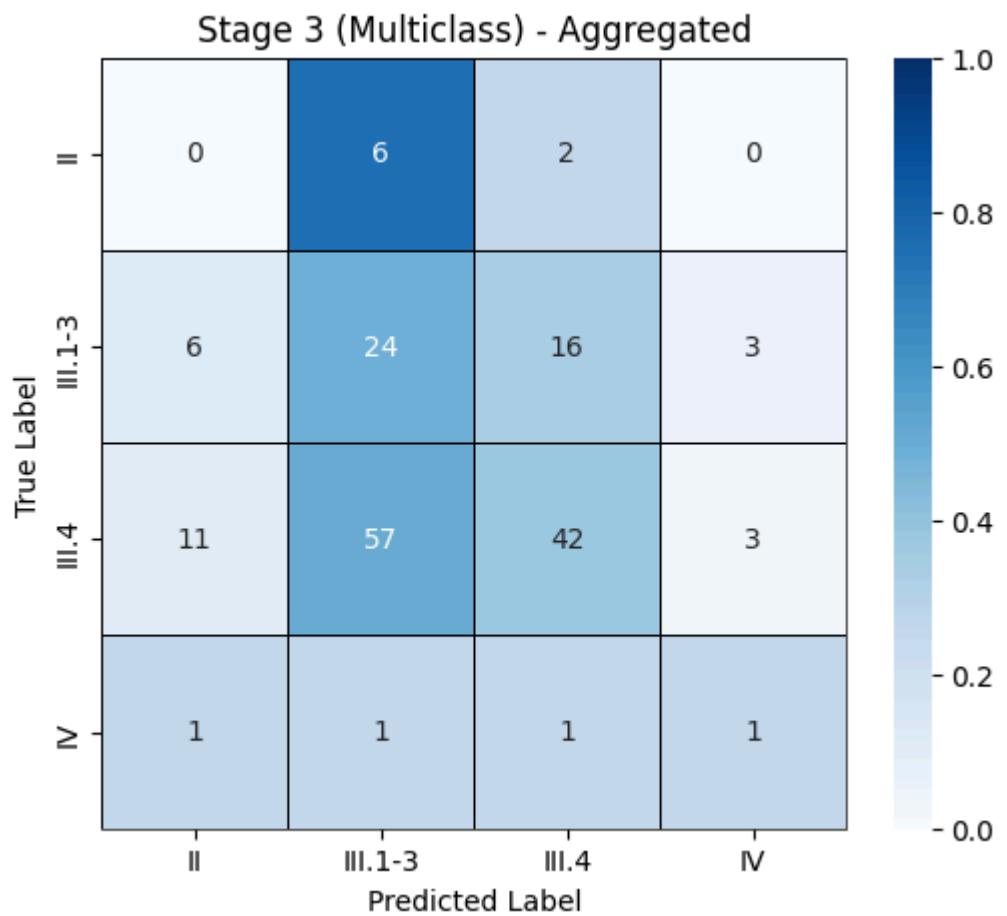
#### STAGE 3 (Multiclass) - AGGREGATED

	precision	recall	f1-score	support
0	0.0000	0.0000	0.0000	8
1	0.2727	0.4898	0.3504	49
2	0.6885	0.3717	0.4828	113
3	0.1429	0.2500	0.1818	4

accuracy	0.3851	174		
macro avg	0.2760	0.2779	0.2537	174
weighted avg	0.5272	0.3851	0.4164	174

Balanced Accuracy: 0.2779




---

SOFT-GATED OVERALL - AGGREGATED

---

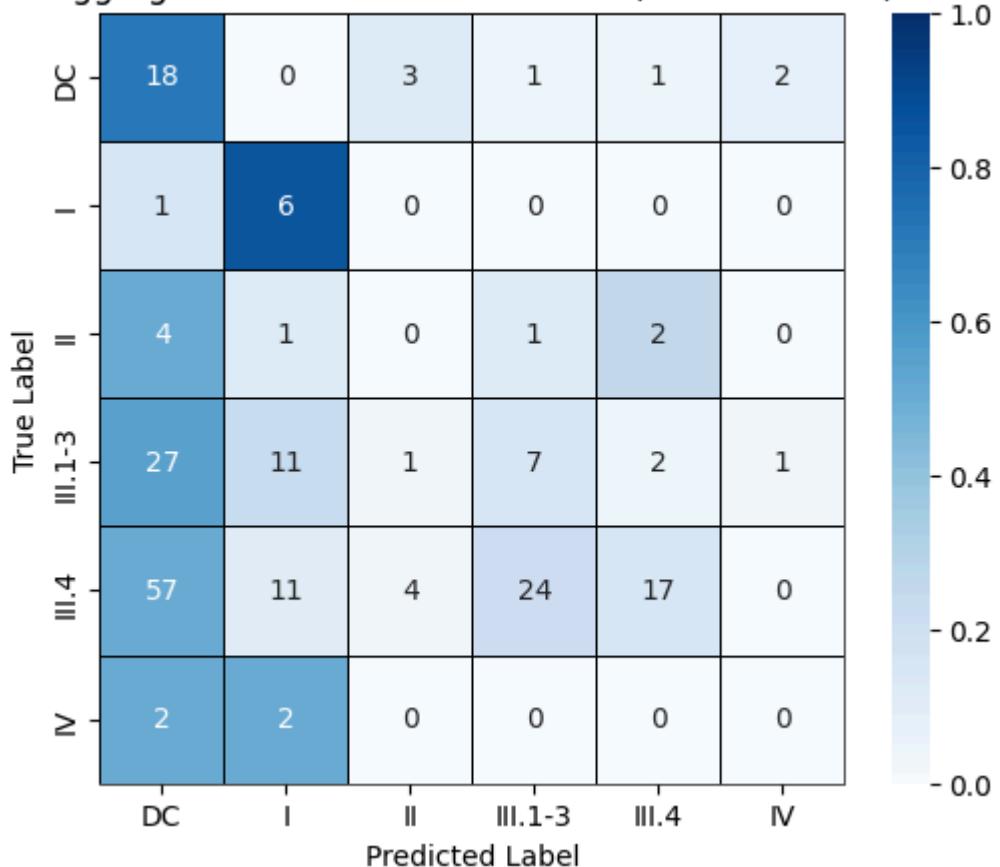
== Aggregated Classification Report ==  
precision recall f1-score support

DC	0.1651	0.7200	0.2687	25
I	0.1935	0.8571	0.3158	7
II	0.0000	0.0000	0.0000	8
III.1-3	0.2121	0.1429	0.1707	49
III.4	0.7727	0.1504	0.2519	113
IV	0.0000	0.0000	0.0000	4

accuracy		0.2330	206	
macro avg	0.2239	0.3117	0.1678	206
weighted avg	0.5009	0.2330	0.2221	206

Aggregated Balanced Accuracy: 0.3117

Aggregated Final Confusion Matrix (Entire Holdout)



== Average Stage Balanced Accuracies ==

Stage 1 (I vs Rest): 0.8997

Stage 2 (DC vs Rest): 0.5756

Stage 3 (Multiclass): 0.2762

Final (Soft-Gated): 0.3127

=====  
AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET  
=====

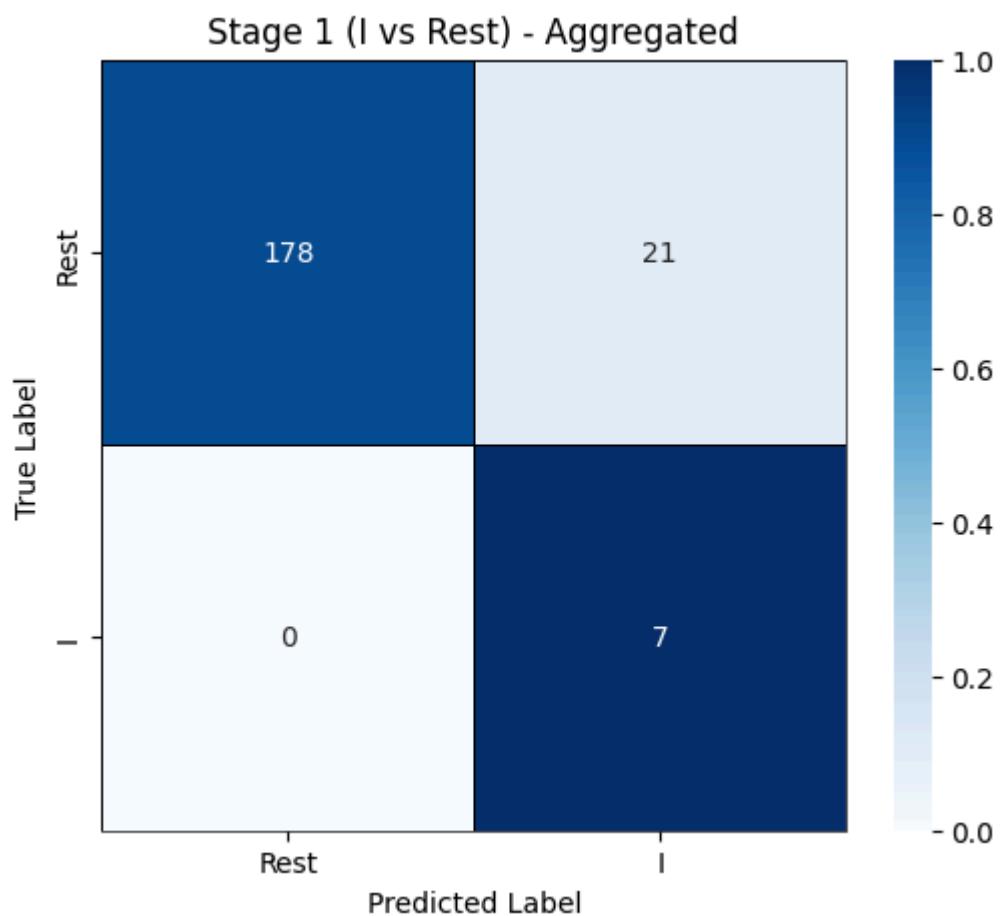
-----  
STAGE 1 (I vs Rest) - AGGREGATED  
-----

precision recall f1-score support

0	1.0000	0.8945	0.9443	199
1	0.2500	1.0000	0.4000	7

accuracy		0.8981	206	
macro avg	0.6250	0.9472	0.6721	206
weighted avg	0.9745	0.8981	0.9258	206

Balanced Accuracy: 0.9472



---

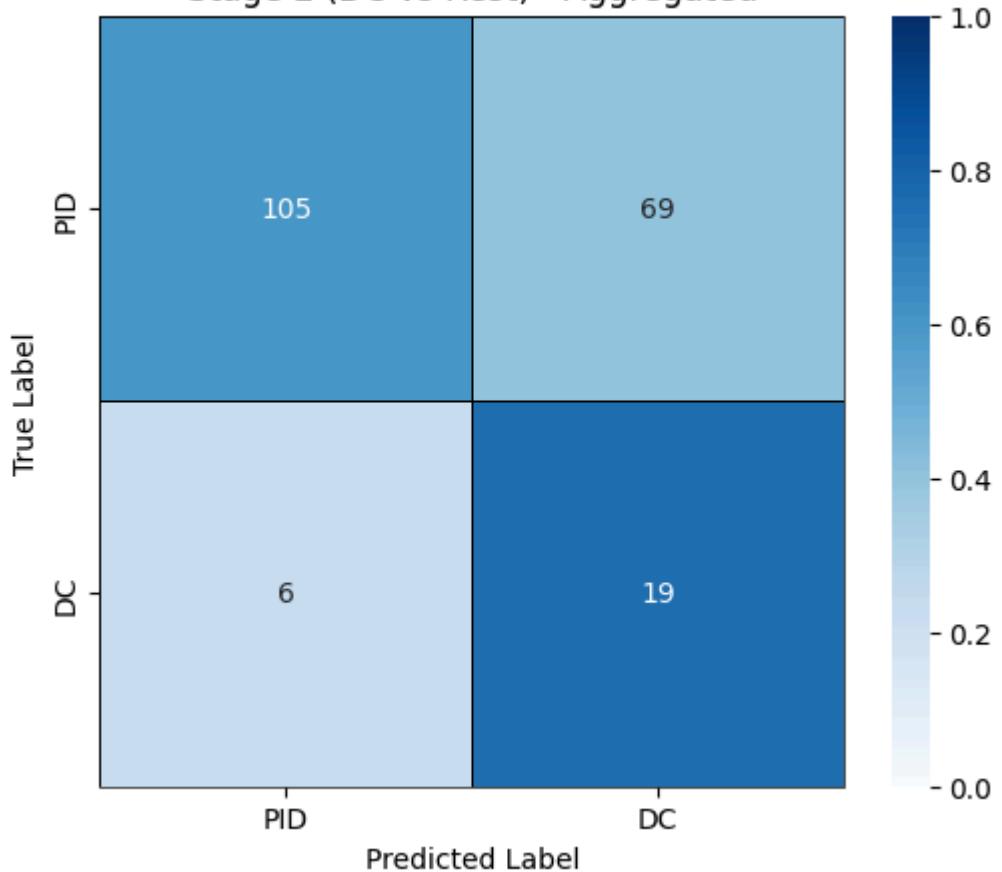
STAGE 2 (DC vs Rest) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.9459	0.6034	0.7368	174
1	0.2159	0.7600	0.3363	25
accuracy		0.6231		199
macro avg	0.5809	0.6817	0.5366	199
weighted avg	0.8542	0.6231	0.6865	199

Balanced Accuracy: 0.6817

Stage 2 (DC vs Rest) - Aggregated



---

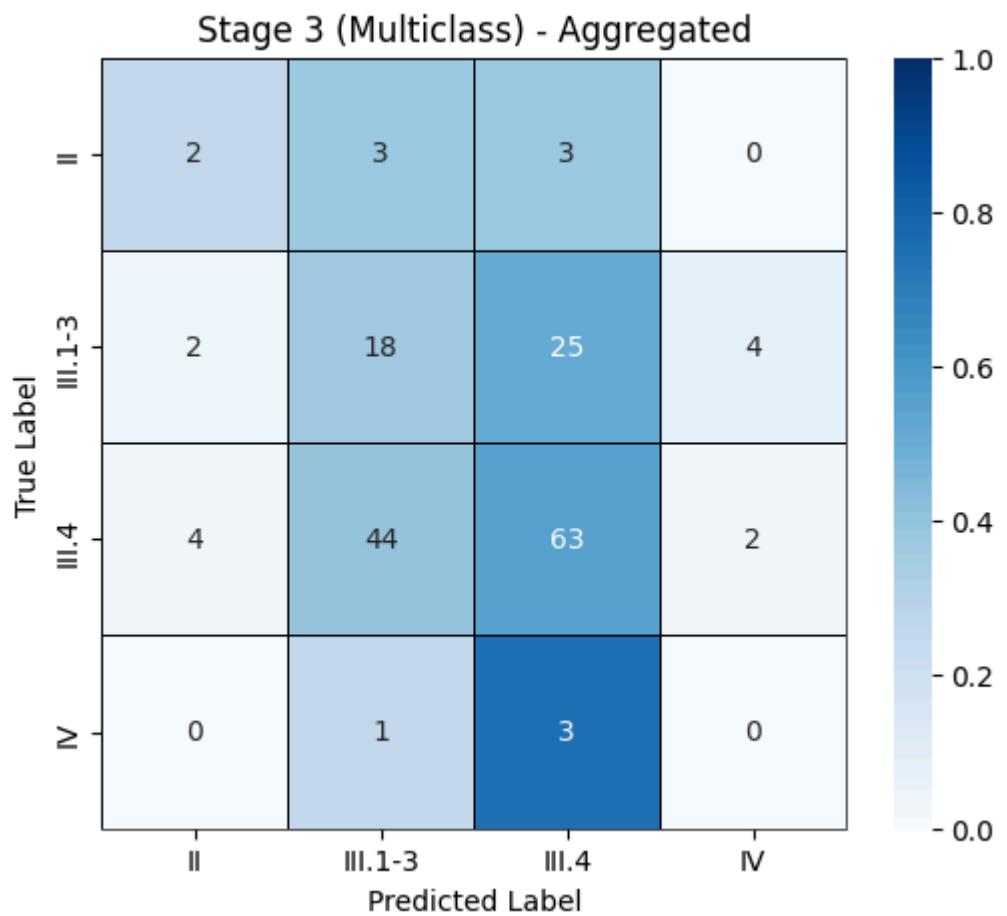
STAGE 3 (Multiclass) - AGGREGATED

---

	precision	recall	f1-score	support
0	0.2500	0.2500	0.2500	8
1	0.2727	0.3673	0.3130	49
2	0.6702	0.5575	0.6087	113
3	0.0000	0.0000	0.0000	4

accuracy                    0.4770    174  
macro avg    0.2982    0.2937    0.2929    174  
weighted avg    0.5235    0.4770    0.4950    174

Balanced Accuracy: 0.2937




---

**SOFT-GATED OVERALL - AGGREGATED**

---

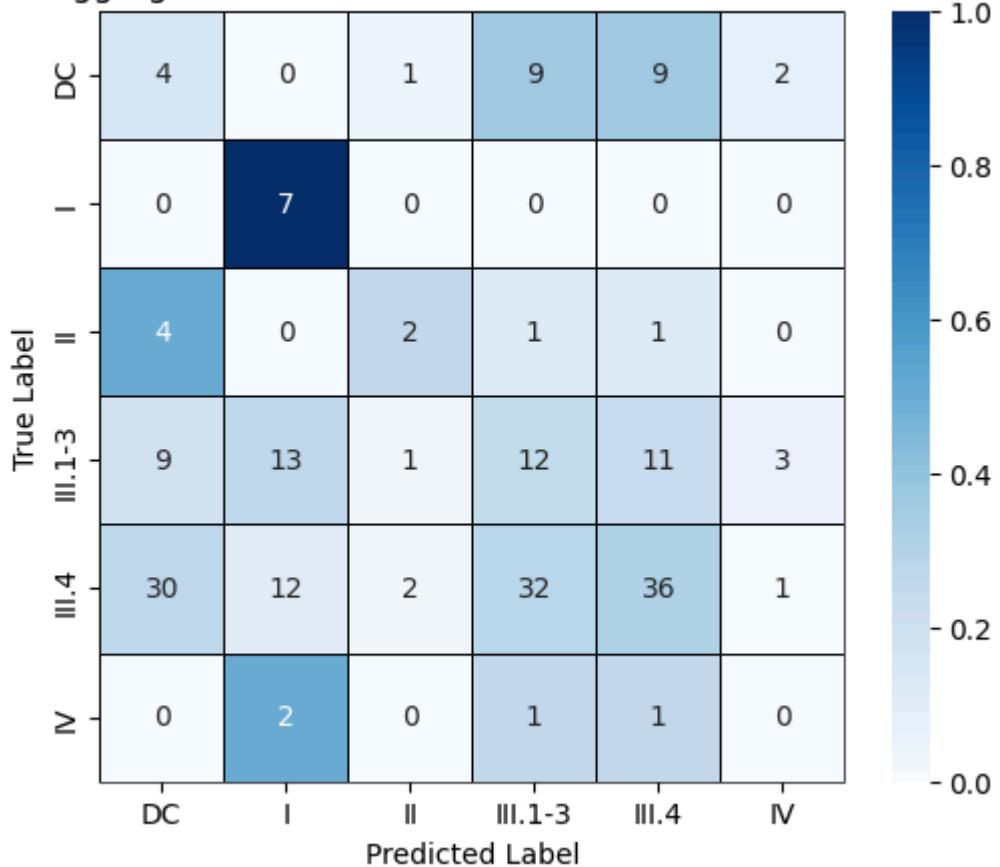
== Aggregated Classification Report ==  
precision recall f1-score support

DC	0.0851	0.1600	0.1111	25
I	0.2059	1.0000	0.3415	7
II	0.3333	0.2500	0.2857	8
III.1-3	0.2182	0.2449	0.2308	49
III.4	0.6207	0.3186	0.4211	113
IV	0.0000	0.0000	0.0000	4

accuracy	0.2961	206		
macro avg	0.2439	0.3289	0.2317	206
weighted avg	0.4226	0.2961	0.3220	206

Aggregated Balanced Accuracy: 0.3289

Aggregated Final Confusion Matrix (Entire Holdout)



== Average Stage Balanced Accuracies ==

Stage 1 (I vs Rest): 0.9472

Stage 2 (DC vs Rest): 0.4668

Stage 3 (Multiclass): 0.2951

Final (Soft-Gated): 0.3302

```
[0.9472361809045227, 0.6817241379310345, 0.29371726566732886, 0.3289136716633556]
[0.9447236180904522, 0.5533333333333333, 0.4326801517067004, 0.36427006200710377]
[0.9195979899497487, 0.6245977011494253, 0.29837908614773345, 0.3610538197579917]
[0.8908829863603733, 0.5534482758620689, 0.3721216362651255, 0.40519986755764253]
[0.8908829863603733, 0.5814942528735632, 0.2778693335741376, 0.311740412979351]
```

$0.919 \pm 0.011$

$0.599 \pm 0.022$

$0.335 \pm 0.026$

$0.354 \pm 0.014$

FLAT

RUN 1: Using Split A for validation, Split B for testing

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

4/4 ————— 0s 35ms/step

== Test Set Evaluation - Run 1 ==

precision recall f1-score support

DC	0.6667	0.1538	0.2500	13
I	0.3333	0.6667	0.4444	3
II	0.2000	0.5000	0.2857	4
IIIa	0.2000	0.5833	0.2979	24
IIIb	0.5000	0.1228	0.1972	57
IV	0.0000	0.0000	0.0000	2

accuracy		0.2621		103
macro avg	0.3167	0.3378	0.2459	103
weighted avg	0.4249	0.2621	0.2341	103

Balanced Accuracy: 0.3377755285650022

Confusion Matrix:

```
[[2 0 0 7 4 0]
 [0 2 1 0 0 0]
 [0 0 2 2 0 0]
 [0 3 4 14 3 0]
 [1 1 3 45 7 0]
 [0 0 0 2 0 0]]
```



=====  
RUN 2: Using Split B for validation, Split A for testing (SWAP)  
=====

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

**1/4** ————— 0s 71ms/step  
**4/4** ————— 0s 22ms/step

== Test Set Evaluation - Run 2 ==

precision recall f1-score support

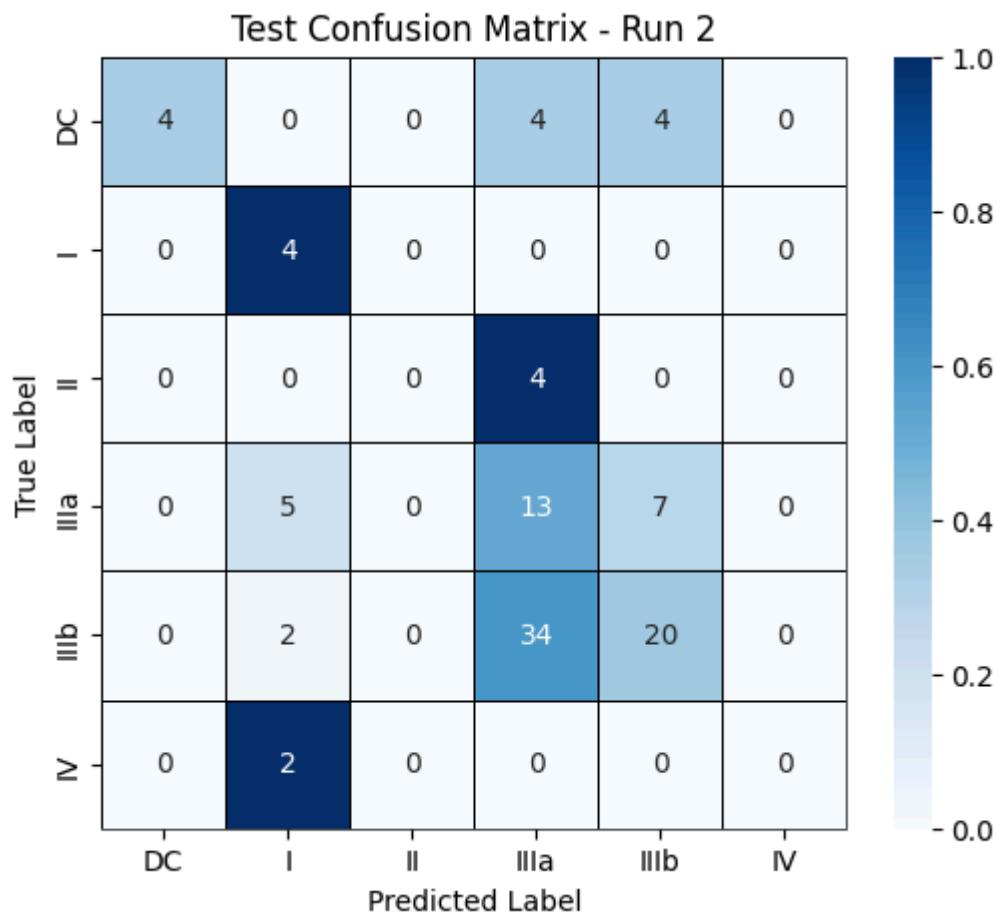
DC	1.0000	0.3333	0.5000	12
I	0.3077	1.0000	0.4706	4
II	0.0000	0.0000	0.0000	4
IIIa	0.2364	0.5200	0.3250	25
IIIb	0.6452	0.3571	0.4598	56
IV	0.0000	0.0000	0.0000	2

accuracy		0.3981		103
macro avg	0.3649	0.3684	0.2926	103
weighted avg	0.5366	0.3981	0.4054	103

Balanced Accuracy: 0.36841269841269836

Confusion Matrix:

```
[[ 4  0  0  4  4  0]
 [ 0  4  0  0  0  0]
 [ 0  0  0  4  0  0]
 [ 0  5  0  13  7  0]
 [ 0  2  0  34  20  0]
 [ 0  2  0  0  0  0]]
```



=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

== Aggregated Classification Report ==

precision recall f1-score support

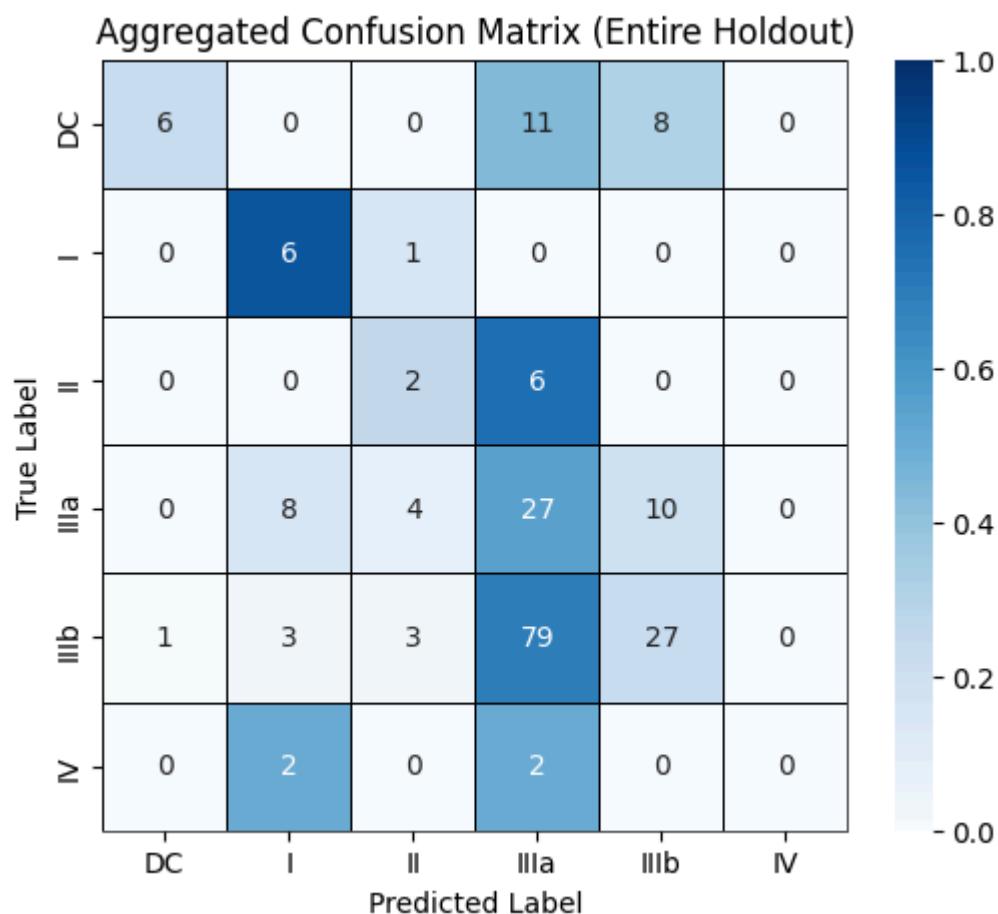
DC	0.8571	0.2400	0.3750	25
I	0.3158	0.8571	0.4615	7
II	0.2000	0.2500	0.2222	8
IIIa	0.2160	0.5510	0.3103	49
IIIb	0.6000	0.2389	0.3418	113
IV	0.0000	0.0000	0.0000	4

accuracy	0.3301	206		
macro avg	0.3648	0.3562	0.2851	206
weighted avg	0.5030	0.3301	0.3311	206

Aggregated Balanced Accuracy: 0.3562

Aggregated Confusion Matrix:

```
[[ 6  0  11  8  0]
 [ 0  6  1  0  0  0]
 [ 0  0  2  6  0  0]
 [ 0  8  4  27 10  0]
 [ 1  3  3  79 27  0]
 [ 0  2  0  2  0  0]]
```



-- Average Balanced Accuracy Across Runs --

Run 1: 0.3378

Run 2: 0.3684

Average: 0.3531

=====  
RUN 1: Using Split A for validation, Split B for testing  
=====

These features will be dropped:

```
['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127',  
'C101', 'C92', 'C40']
```

**4/4** ————— **0s** 26ms/step

== Test Set Evaluation - Run 1 ==

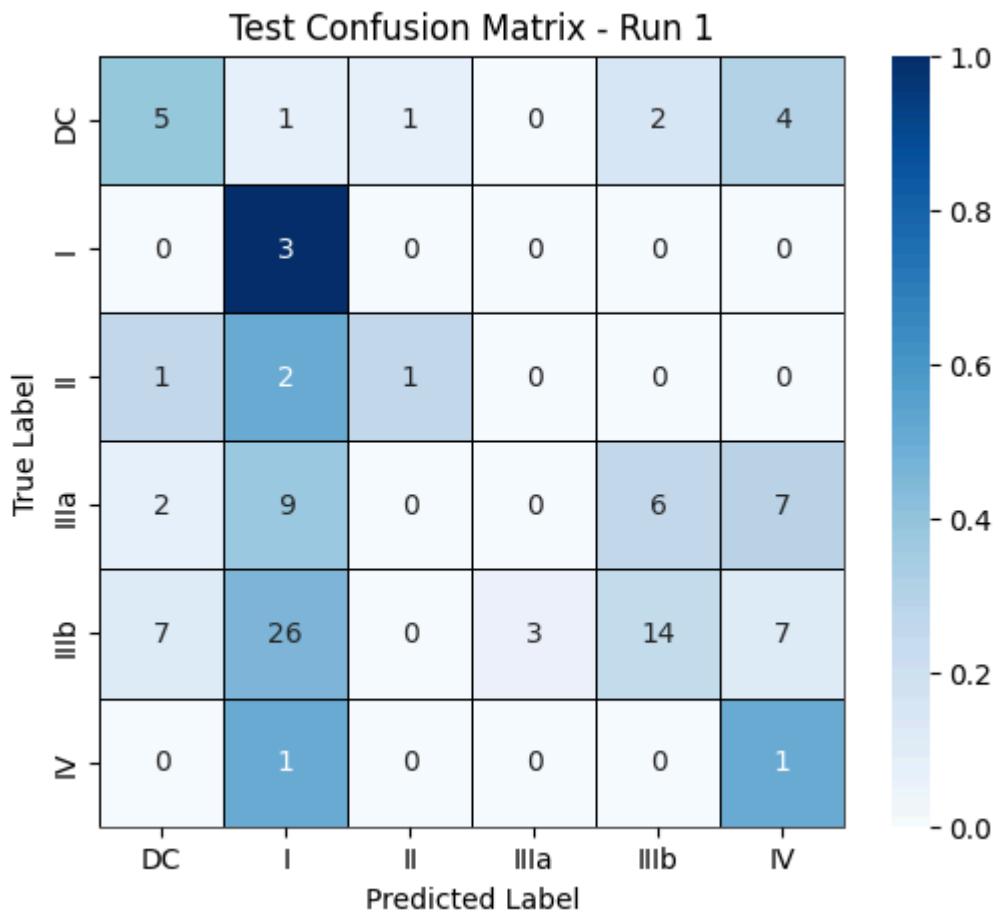
precision recall f1-score support

	DC	0.3333	0.3846	0.3571	13
I	0.0714	1.0000	0.1333	3	
II	0.5000	0.2500	0.3333	4	
IIIa	0.0000	0.0000	0.0000	24	
IIIb	0.6364	0.2456	0.3544	57	
IV	0.0526	0.5000	0.0952	2	
	accuracy		0.2330		103
	macro avg	0.2656	0.3967	0.2122	103
	weighted avg	0.4168	0.2330	0.2599	103

Balanced Accuracy: 0.3967049032838506

Confusion Matrix:

```
[[ 5 1 1 0 2 4]  
 [0 3 0 0 0 0]  
 [1 2 1 0 0 0]  
 [2 9 0 0 6 7]  
 [7 2 6 0 3 14 7]  
 [0 1 0 0 0 1]]
```



=====

RUN 2: Using Split B for validation, Split A for testing (SWAP)

=====

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

**4/4** ————— **0s 21ms/step**

== Test Set Evaluation - Run 2 ==

precision recall f1-score support

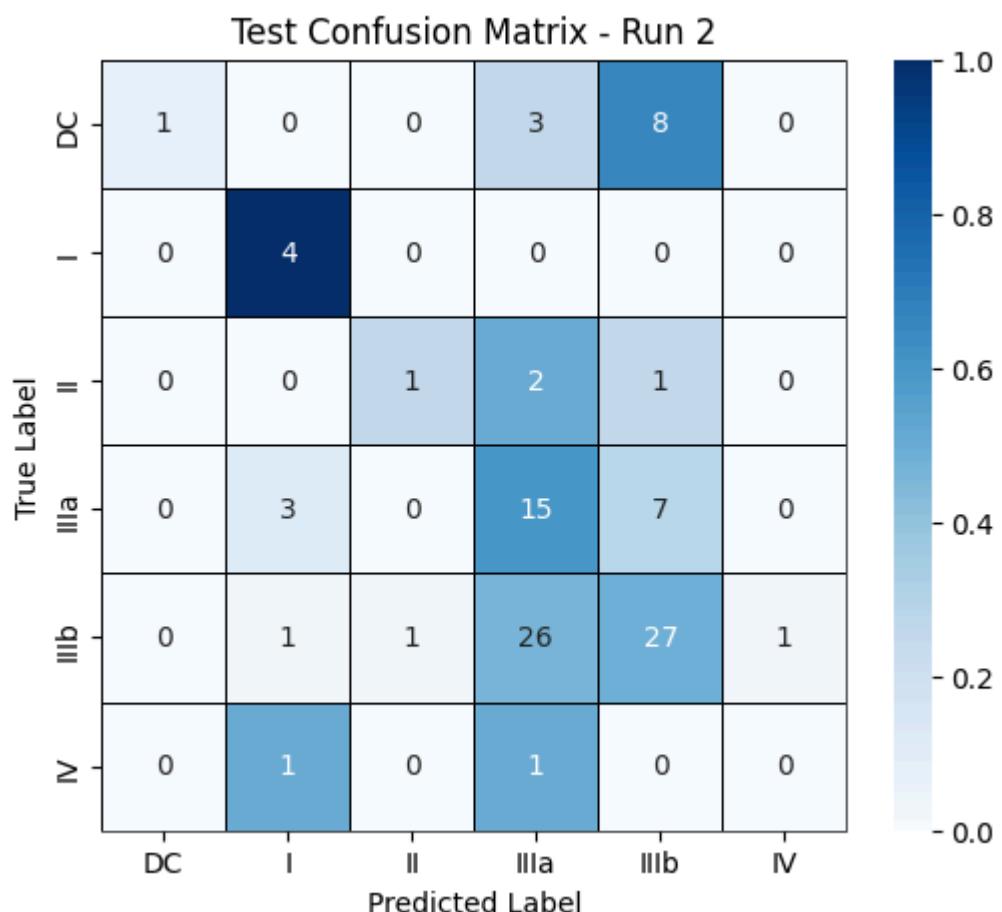
DC	1.0000	0.0833	0.1538	12
I	0.4444	1.0000	0.6154	4
II	0.5000	0.2500	0.3333	4
IIIa	0.3191	0.6000	0.4167	25
IIIb	0.6279	0.4821	0.5455	56
IV	0.0000	0.0000	0.0000	2

accuracy	0.4660	103		
macro avg	0.4819	0.4026	0.3441	103
weighted avg	0.5720	0.4660	0.4525	103

Balanced Accuracy: 0.40257936507936504

Confusion Matrix:

```
[[1 0 0 3 8 0]
 [0 4 0 0 0 0]
 [0 0 1 2 1 0]
 [0 3 0 15 7 0]
 [0 1 1 26 27 1]
 [0 1 0 1 0 0]]
```



=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

== Aggregated Classification Report ==

precision recall f1-score support

DC	0.3750	0.2400	0.2927	25
I	0.1373	1.0000	0.2414	7
II	0.5000	0.2500	0.3333	8
IIIa	0.3000	0.3061	0.3030	49
IIIb	0.6308	0.3628	0.4607	113
IV	0.0500	0.2500	0.0833	4

accuracy 0.3495 206

macro avg	0.3322	0.4015	0.2857	206
weighted avg	0.4879	0.3495	0.3831	206

Aggregated Balanced Accuracy: 0.4015

Aggregated Confusion Matrix:

```
[6 1 1 3 10 4]
```

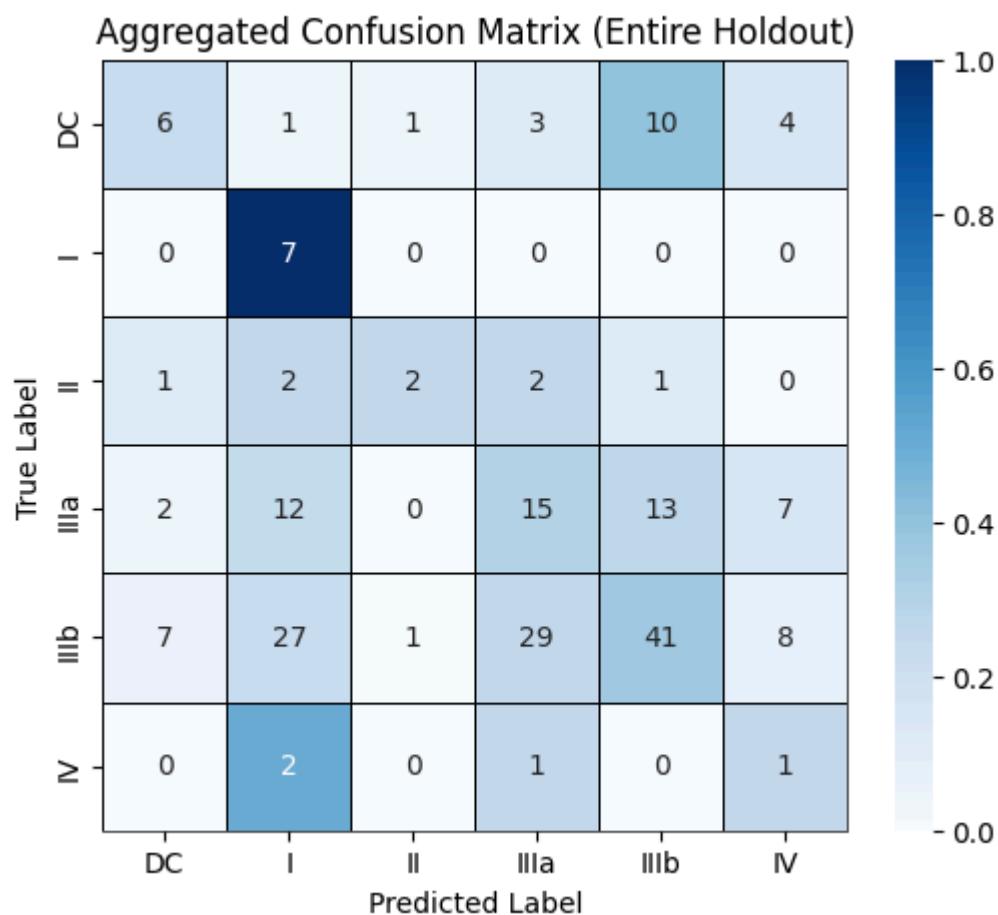
```
[0 7 0 0 0 0]
```

```
[1 2 2 2 1 0]
```

```
[2 12 0 15 13 7]
```

```
[7 27 1 29 41 8]
```

```
[0 2 0 1 0 1]]
```



-- Average Balanced Accuracy Across Runs --

Run 1: 0.3967

Run 2: 0.4026

Average: 0.3996

=====

RUN 1: Using Split A for validation, Split B for testing

=====

These features will be dropped:

```
[C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127',  
'C101', 'C92', 'C40']
```

**4/4** ————— **0s** 26ms/step

== Test Set Evaluation - Run 1 ==

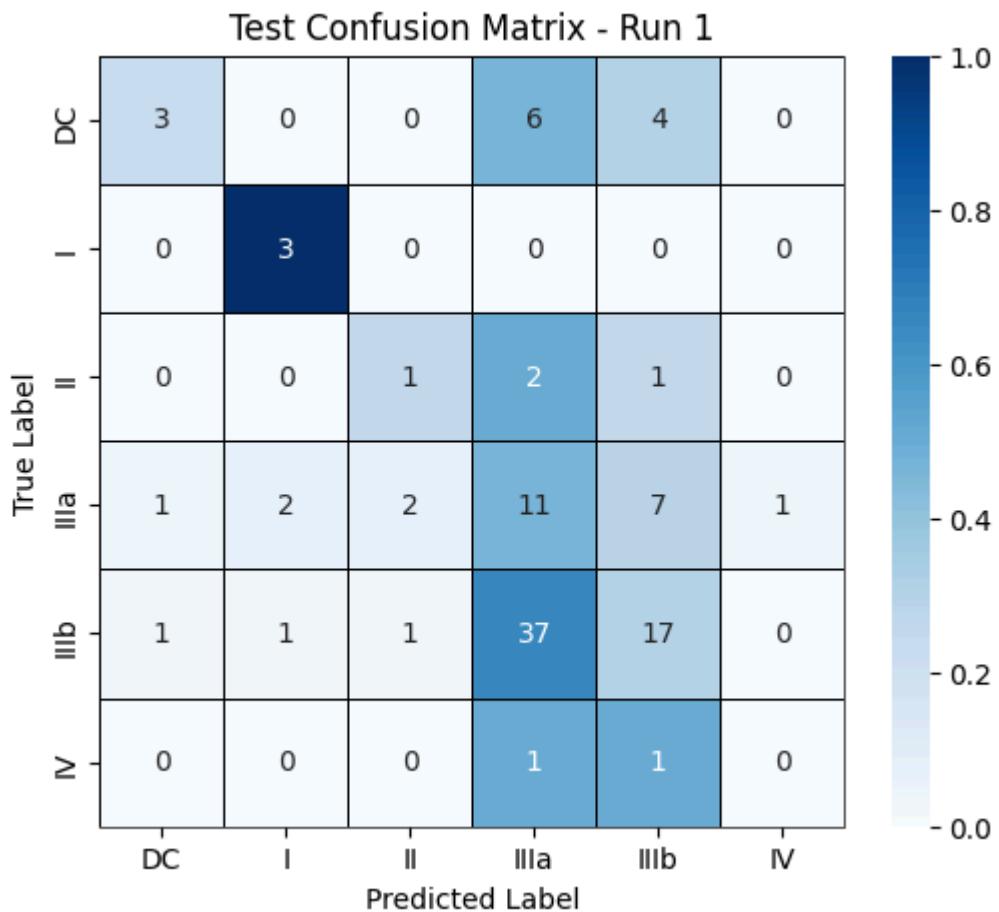
precision recall f1-score support

	DC	0.6000	0.2308	0.3333	13
I	0.5000	1.0000	0.6667	3	
II	0.2500	0.2500	0.2500	4	
IIIa	0.1930	0.4583	0.2716	24	
IIIb	0.5667	0.2982	0.3908	57	
IV	0.0000	0.0000	0.0000	2	
	accuracy		0.3398		103
	macro avg	0.3516	0.3729	0.3187	103
	weighted avg	0.4586	0.3398	0.3508	103

Balanced Accuracy: 0.37289136302294196

Confusion Matrix:

```
[[ 3  0  0  6  4  0]  
 [ 0  3  0  0  0  0]  
 [ 0  0  1  2  1  0]  
 [ 1  2  2  11  7  1]  
 [ 1  1  1  37  17  0]  
 [ 0  0  0  1  1  0]]
```



=====  
RUN 2: Using Split B for validation, Split A for testing (SWAP)  
=====

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

**4/4** ————— **0s 34ms/step**

== Test Set Evaluation - Run 2 ==

precision recall f1-score support

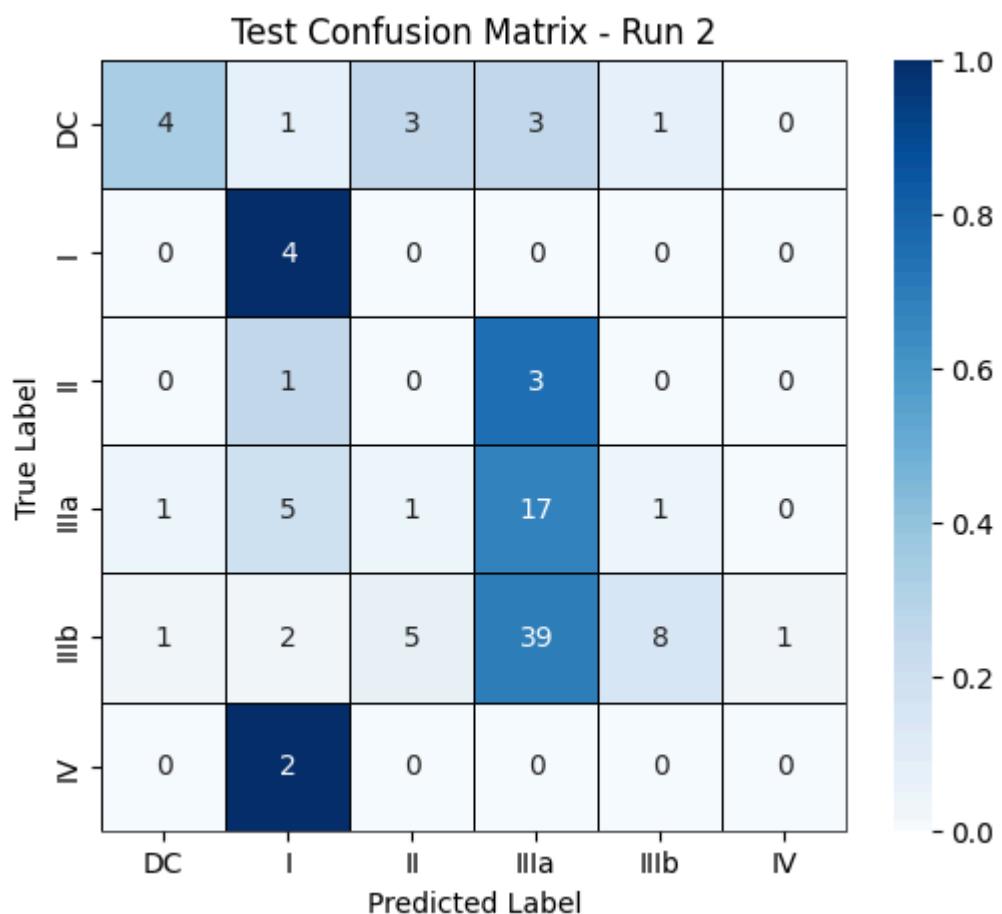
DC	0.6667	0.3333	0.4444	12
I	0.2667	1.0000	0.4211	4
II	0.0000	0.0000	0.0000	4
IIIa	0.2742	0.6800	0.3908	25
IIIb	0.8000	0.1429	0.2424	56
IV	0.0000	0.0000	0.0000	2

accuracy                    0.3204    103  
macro avg    0.3346    0.3594    0.2498    103  
weighted avg    0.5895    0.3204    0.2948    103

Balanced Accuracy: 0.35936507936507933

Confusion Matrix:

```
[[4 1 3 3 1 0]
 [0 4 0 0 0 0]
 [0 1 0 3 0 0]
 [1 5 1 17 1 0]
 [1 2 5 39 8 1]
 [0 2 0 0 0 0]]
```



=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

== Aggregated Classification Report ==

precision recall f1-score support

DC	0.6364	0.2800	0.3889	25
I	0.3333	1.0000	0.5000	7
II	0.0769	0.1250	0.0952	8
IIIa	0.2353	0.5714	0.3333	49
IIIb	0.6250	0.2212	0.3268	113
IV	0.0000	0.0000	0.0000	4

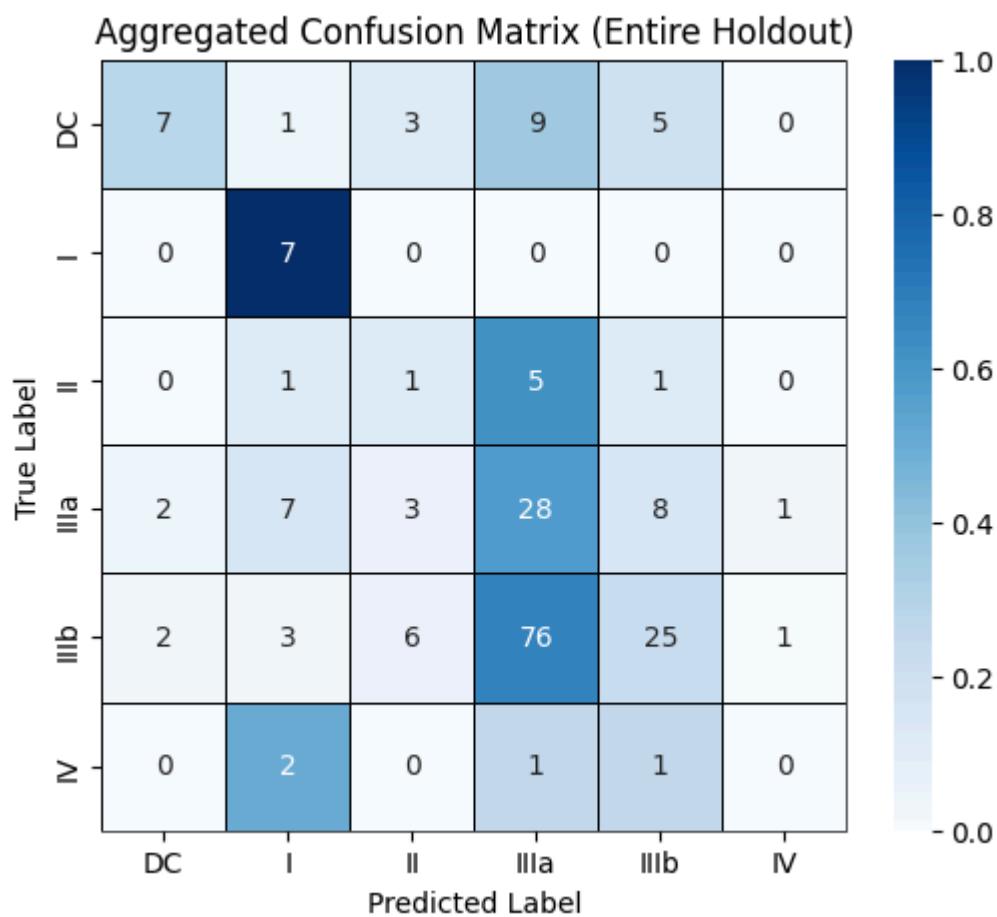
accuracy                    0.3301            206

macro avg	0.3178	0.3663	0.2740	206
weighted avg	0.4904	0.3301	0.3264	206

Aggregated Balanced Accuracy: 0.3663

Aggregated Confusion Matrix:

```
[[ 7  1  3  9  5  0]
 [ 0  7  0  0  0  0]
 [ 0  1  1  5  1  0]
 [ 2  7  3  28  8  1]
 [ 2  3  6  76  25  1]
 [ 0  2  0  1  1  0]]
```



== Average Balanced Accuracy Across Runs ==

Run 1: 0.3729

Run 2: 0.3594

Average: 0.3661

=====

RUN 1: Using Split A for validation, Split B for testing

=====

These features will be dropped:

```
[C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127',  
'C101', 'C92', 'C40']
```

**4/4** ————— **0s** 31ms/step

== Test Set Evaluation - Run 1 ==

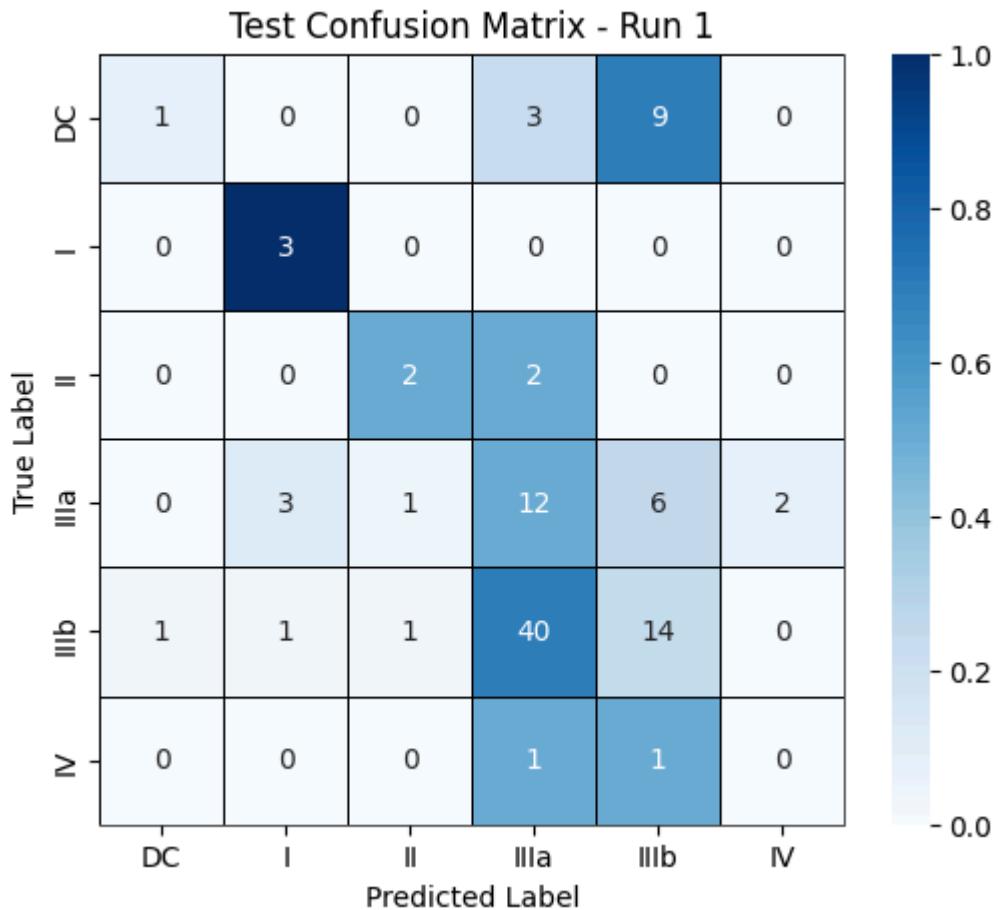
precision recall f1-score support

	DC	0.5000	0.0769	0.1333	13
I	0.4286	1.0000	0.6000	3	
II	0.5000	0.5000	0.5000	4	
IIIa	0.2069	0.5000	0.2927	24	
IIIb	0.4667	0.2456	0.3218	57	
IV	0.0000	0.0000	0.0000	2	
	accuracy		0.3107		103
	macro avg	0.3504	0.3871	0.3080	103
	weighted avg	0.4015	0.3107	0.3000	103

Balanced Accuracy: 0.38708951866846597

Confusion Matrix:

```
[[ 1  0  0  3  9  0]  
 [ 0  3  0  0  0  0]  
 [ 0  0  2  2  0  0]  
 [ 0  3  1 12  6  2]  
 [ 1  1  1 40 14  0]  
 [ 0  0  0  1  1  0]]
```



=====

RUN 2: Using Split B for validation, Split A for testing (SWAP)

=====

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

**4/4** ————— **0s** 20ms/step

== Test Set Evaluation - Run 2 ==

precision recall f1-score support

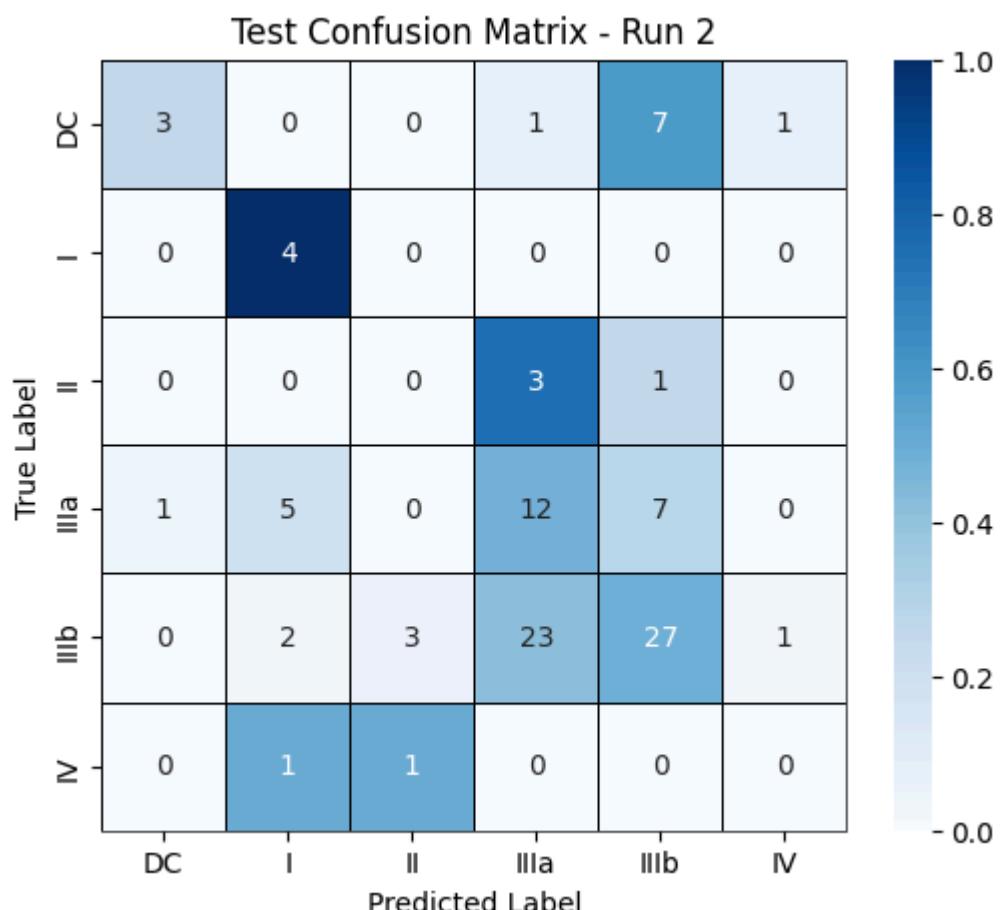
DC	0.7500	0.2500	0.3750	12
I	0.3333	1.0000	0.5000	4
II	0.0000	0.0000	0.0000	4
IIIa	0.3077	0.4800	0.3750	25
IIIb	0.6429	0.4821	0.5510	56
IV	0.0000	0.0000	0.0000	2

accuracy	0.4466	103		
macro avg	0.3390	0.3687	0.3002	103
weighted avg	0.5245	0.4466	0.4537	103

Balanced Accuracy: 0.3686904761904762

Confusion Matrix:

```
[[3 0 0 1 7 1]
 [0 4 0 0 0 0]
 [0 0 0 3 1 0]
 [1 5 0 12 7 0]
 [0 2 3 23 27 1]
 [0 1 1 0 0 0]]
```



=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

== Aggregated Classification Report ==

precision recall f1-score support

DC	0.6667	0.1600	0.2581	25
I	0.3684	1.0000	0.5385	7
II	0.2500	0.2500	0.2500	8
IIIa	0.2474	0.4898	0.3288	49
IIIb	0.5694	0.3628	0.4432	113
IV	0.0000	0.0000	0.0000	4

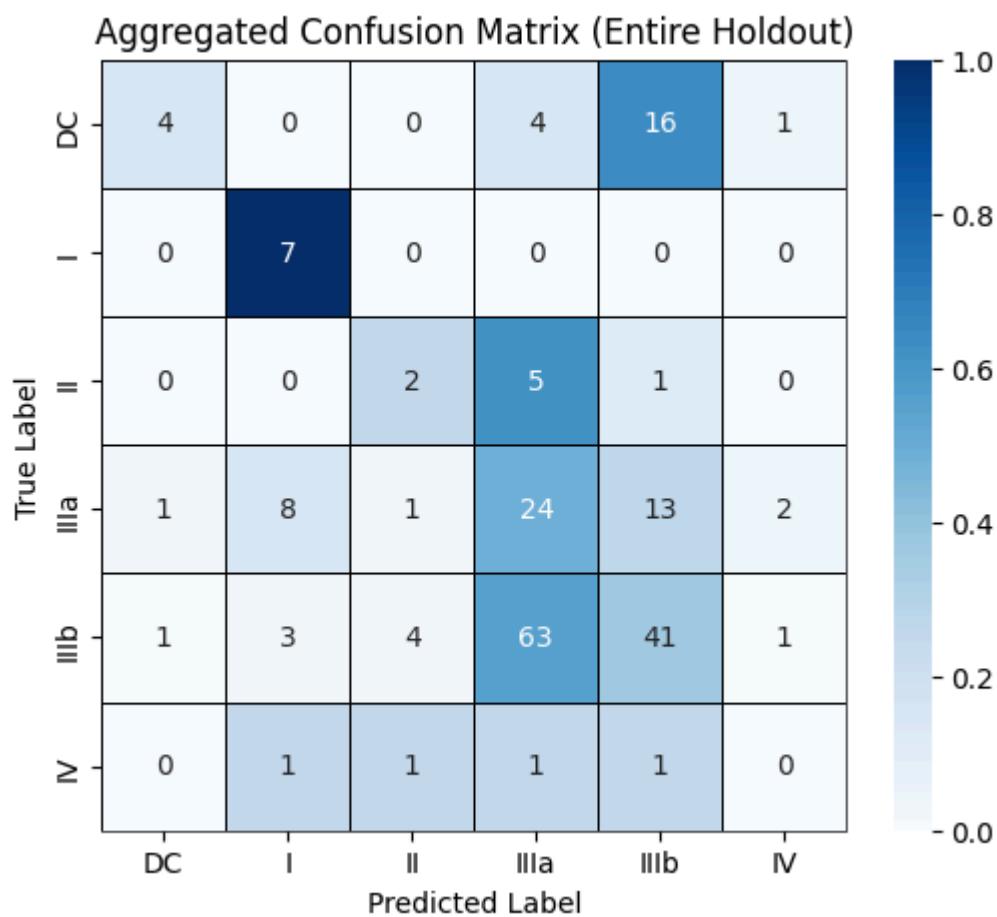
accuracy                    0.3786            206

macro avg	0.3503	0.3771	0.3031	206
weighted avg	0.4744	0.3786	0.3807	206

Aggregated Balanced Accuracy: 0.3771

Aggregated Confusion Matrix:

```
[[ 4  0  0  4 16  1]
 [ 0  7  0  0  0  0]
 [ 0  0  2  5  1  0]
 [ 1  8  1 24 13  2]
 [ 1  3  4 63 41  1]
 [ 0  1  1  1  1  0]]
```



== Average Balanced Accuracy Across Runs ==

Run 1: 0.3871

Run 2: 0.3687

Average: 0.3779

RUN 1: Using Split A for validation, Split B for testing

These features will be dropped:

```
[C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127',  
'C101', 'C92', 'C40']
```

**4/4** ————— **0s** 29ms/step

== Test Set Evaluation - Run 1 ==

precision recall f1-score support

	DC	0.3500	0.5385	0.4242	13
I	0.3750	1.0000	0.5455	3	
II	0.5000	0.2500	0.3333	4	
IIIa	0.1714	0.2500	0.2034	24	
IIIb	0.6571	0.4035	0.5000	57	
IV	0.0000	0.0000	0.0000	2	
	accuracy		0.3883		103
	macro avg	0.3423	0.4070	0.3344	103
	weighted avg	0.4781	0.3883	0.4065	103

Balanced Accuracy: 0.4069950517318938

Confusion Matrix:

```
[[ 7  0  0  2  4  0]  
 [ 0  3  0  0  0  0]  
 [ 1  0  1  1  1  0]  
 [ 6  3  0  6  6  3]  
 [ 6  2  1  25 23  0]  
 [ 0  0  0  1  1  0]]
```



=====

RUN 2: Using Split B for validation, Split A for testing (SWAP)

=====

These features will be dropped:

['C113', 'C126', 'C100', 'C66', 'C11', 'C82', 'C135', 'C51', 'C86', 'C85', 'C96', 'C64', 'C87', 'C139', 'C127', 'C101', 'C92', 'C40']

**4/4** ————— **0s** 30ms/step

== Test Set Evaluation - Run 2 ==

precision recall f1-score support

DC	0.6667	0.1667	0.2667	12
I	0.1739	1.0000	0.2963	4
II	0.0000	0.0000	0.0000	4
IIIa	0.2879	0.7600	0.4176	25
IIIb	0.2000	0.0179	0.0328	56
IV	0.0000	0.0000	0.0000	2

accuracy	0.2524	103		
macro avg	0.2214	0.3241	0.1689	103
weighted avg	0.2630	0.2524	0.1618	103

Balanced Accuracy: 0.3240873015873016

Confusion Matrix:

```
[[2 0 0 4 3 3]
 [0 4 0 0 0 0]
 [0 1 0 2 1 0]
 [0 6 0 19 0 0]
 [1 10 2 41 1 1]
 [0 2 0 0 0 0]]
```



=====

AGGREGATED RESULTS ACROSS ENTIRE HOLDOUT SET

=====

== Aggregated Classification Report ==

precision recall f1-score support

DC	0.3913	0.3600	0.3750	25
I	0.2258	1.0000	0.3684	7
II	0.2500	0.1250	0.1667	8
IIIa	0.2475	0.5102	0.3333	49
IIIb	0.6000	0.2124	0.3137	113
IV	0.0000	0.0000	0.0000	4

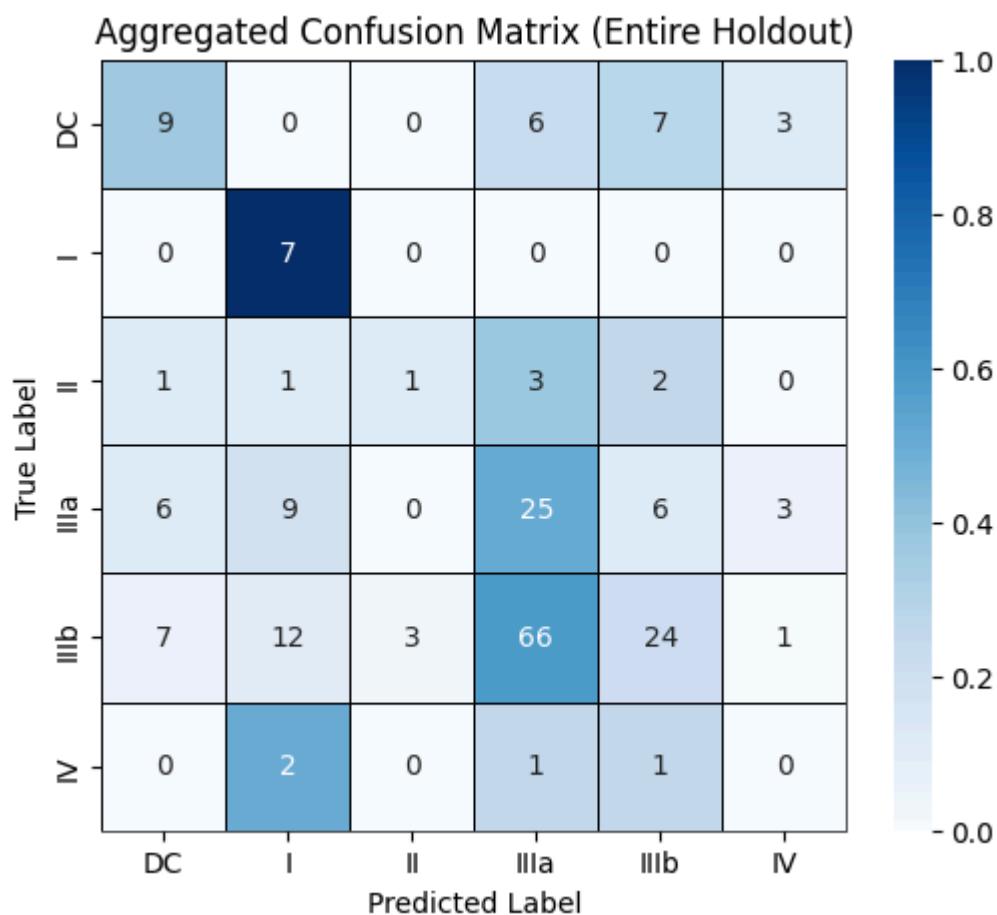
accuracy 0.3204 206

macro avg	0.2858	0.3679	0.2595	206
weighted avg	0.4529	0.3204	0.3159	206

Aggregated Balanced Accuracy: 0.3679

Aggregated Confusion Matrix:

```
[[ 9  0  0  6  7  3]
 [ 0  7  0  0  0  0]
 [ 1  1  1  3  2  0]
 [ 6  9  0 25  6  3]
 [ 7 12  3 66 24  1]
 [ 0  2  0  1  1  0]]
```



== Average Balanced Accuracy Across Runs ==

Run 1: 0.4070

Run 2: 0.3241

Average: 0.3655

[0.3562, 0.4015, 0.3663, 0.3771, 0.3679]

$0.373 \pm 0.009$